

RabbitMQ Overview

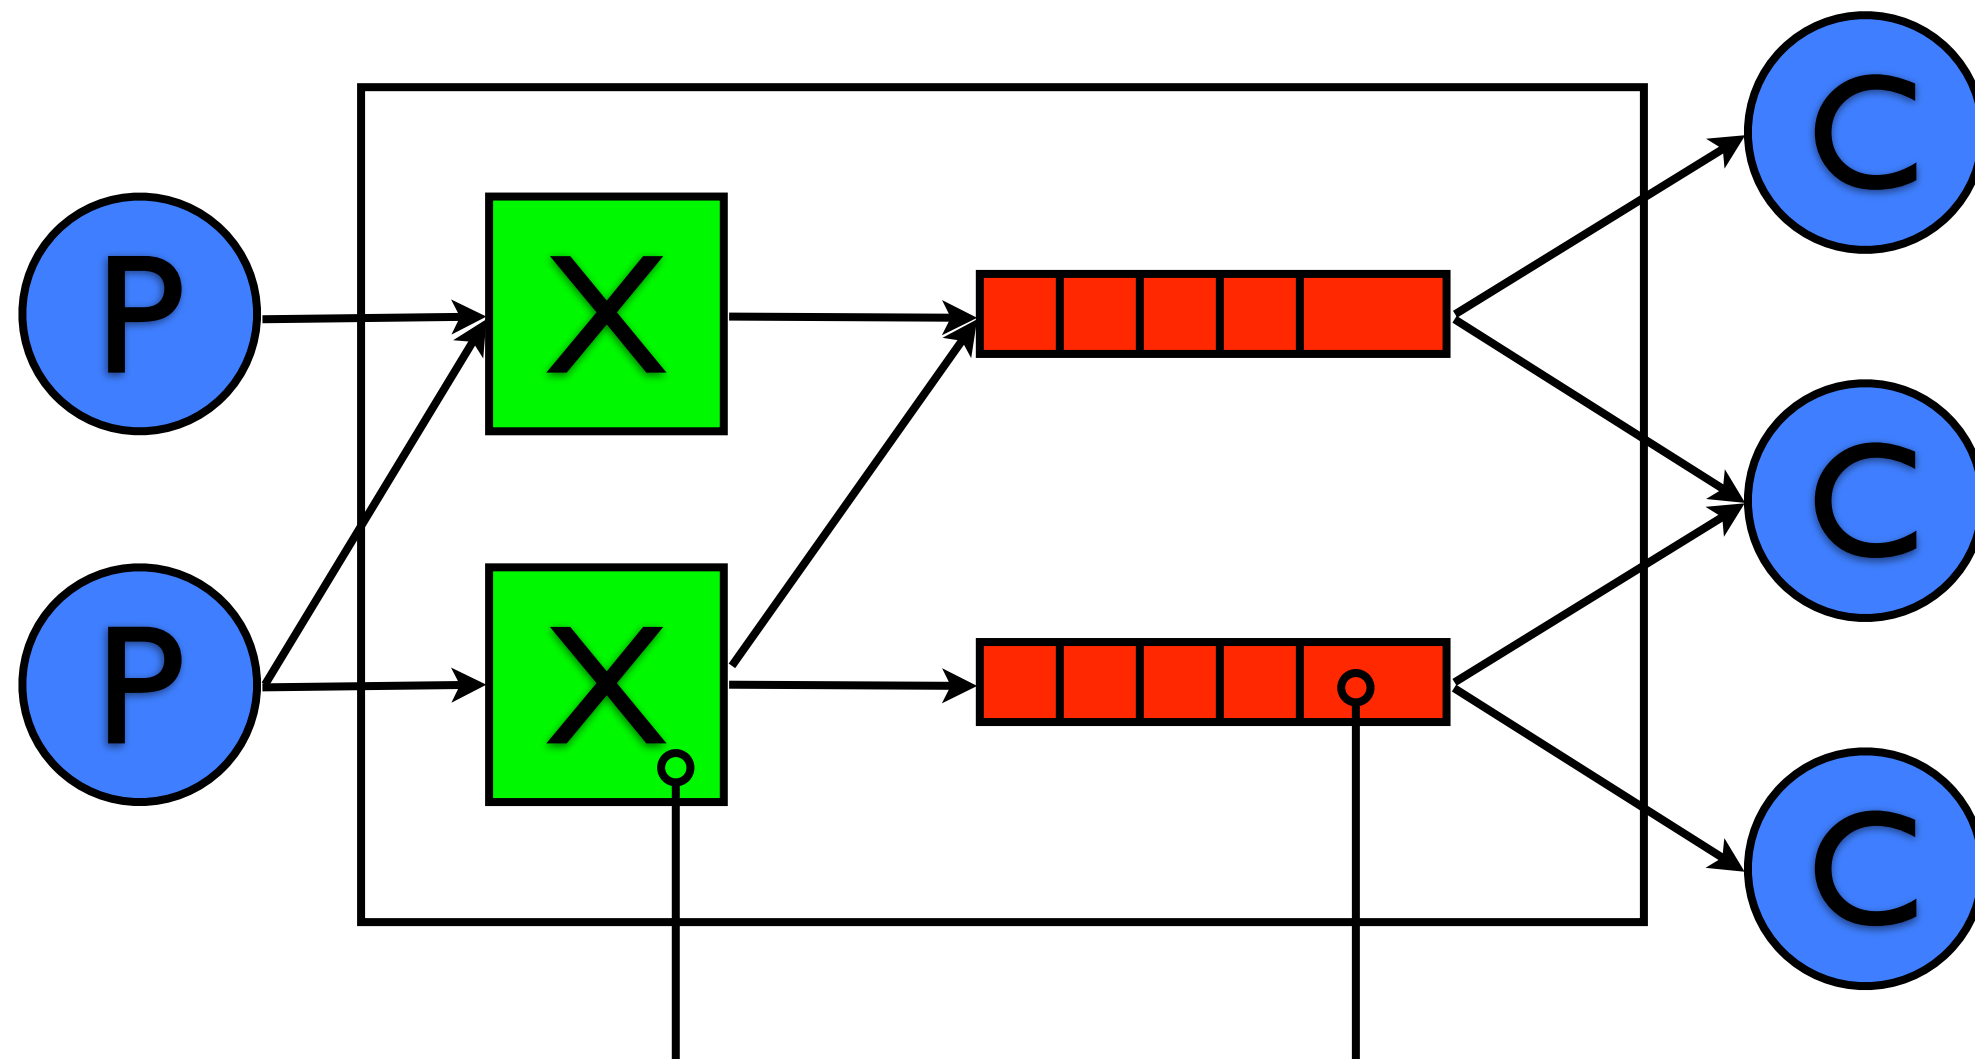
Tony Garnock-Jones <tonyg@lshift.net>

Agenda

- AMQP in 3 minutes
- RabbitMQ architecture
- Availability, Clustering, Federation
- Durability, Persistence, Memory usage
- Security
- Operational Tools
- Ongoing work

AMQP

AMQP Basics (0-9-1)

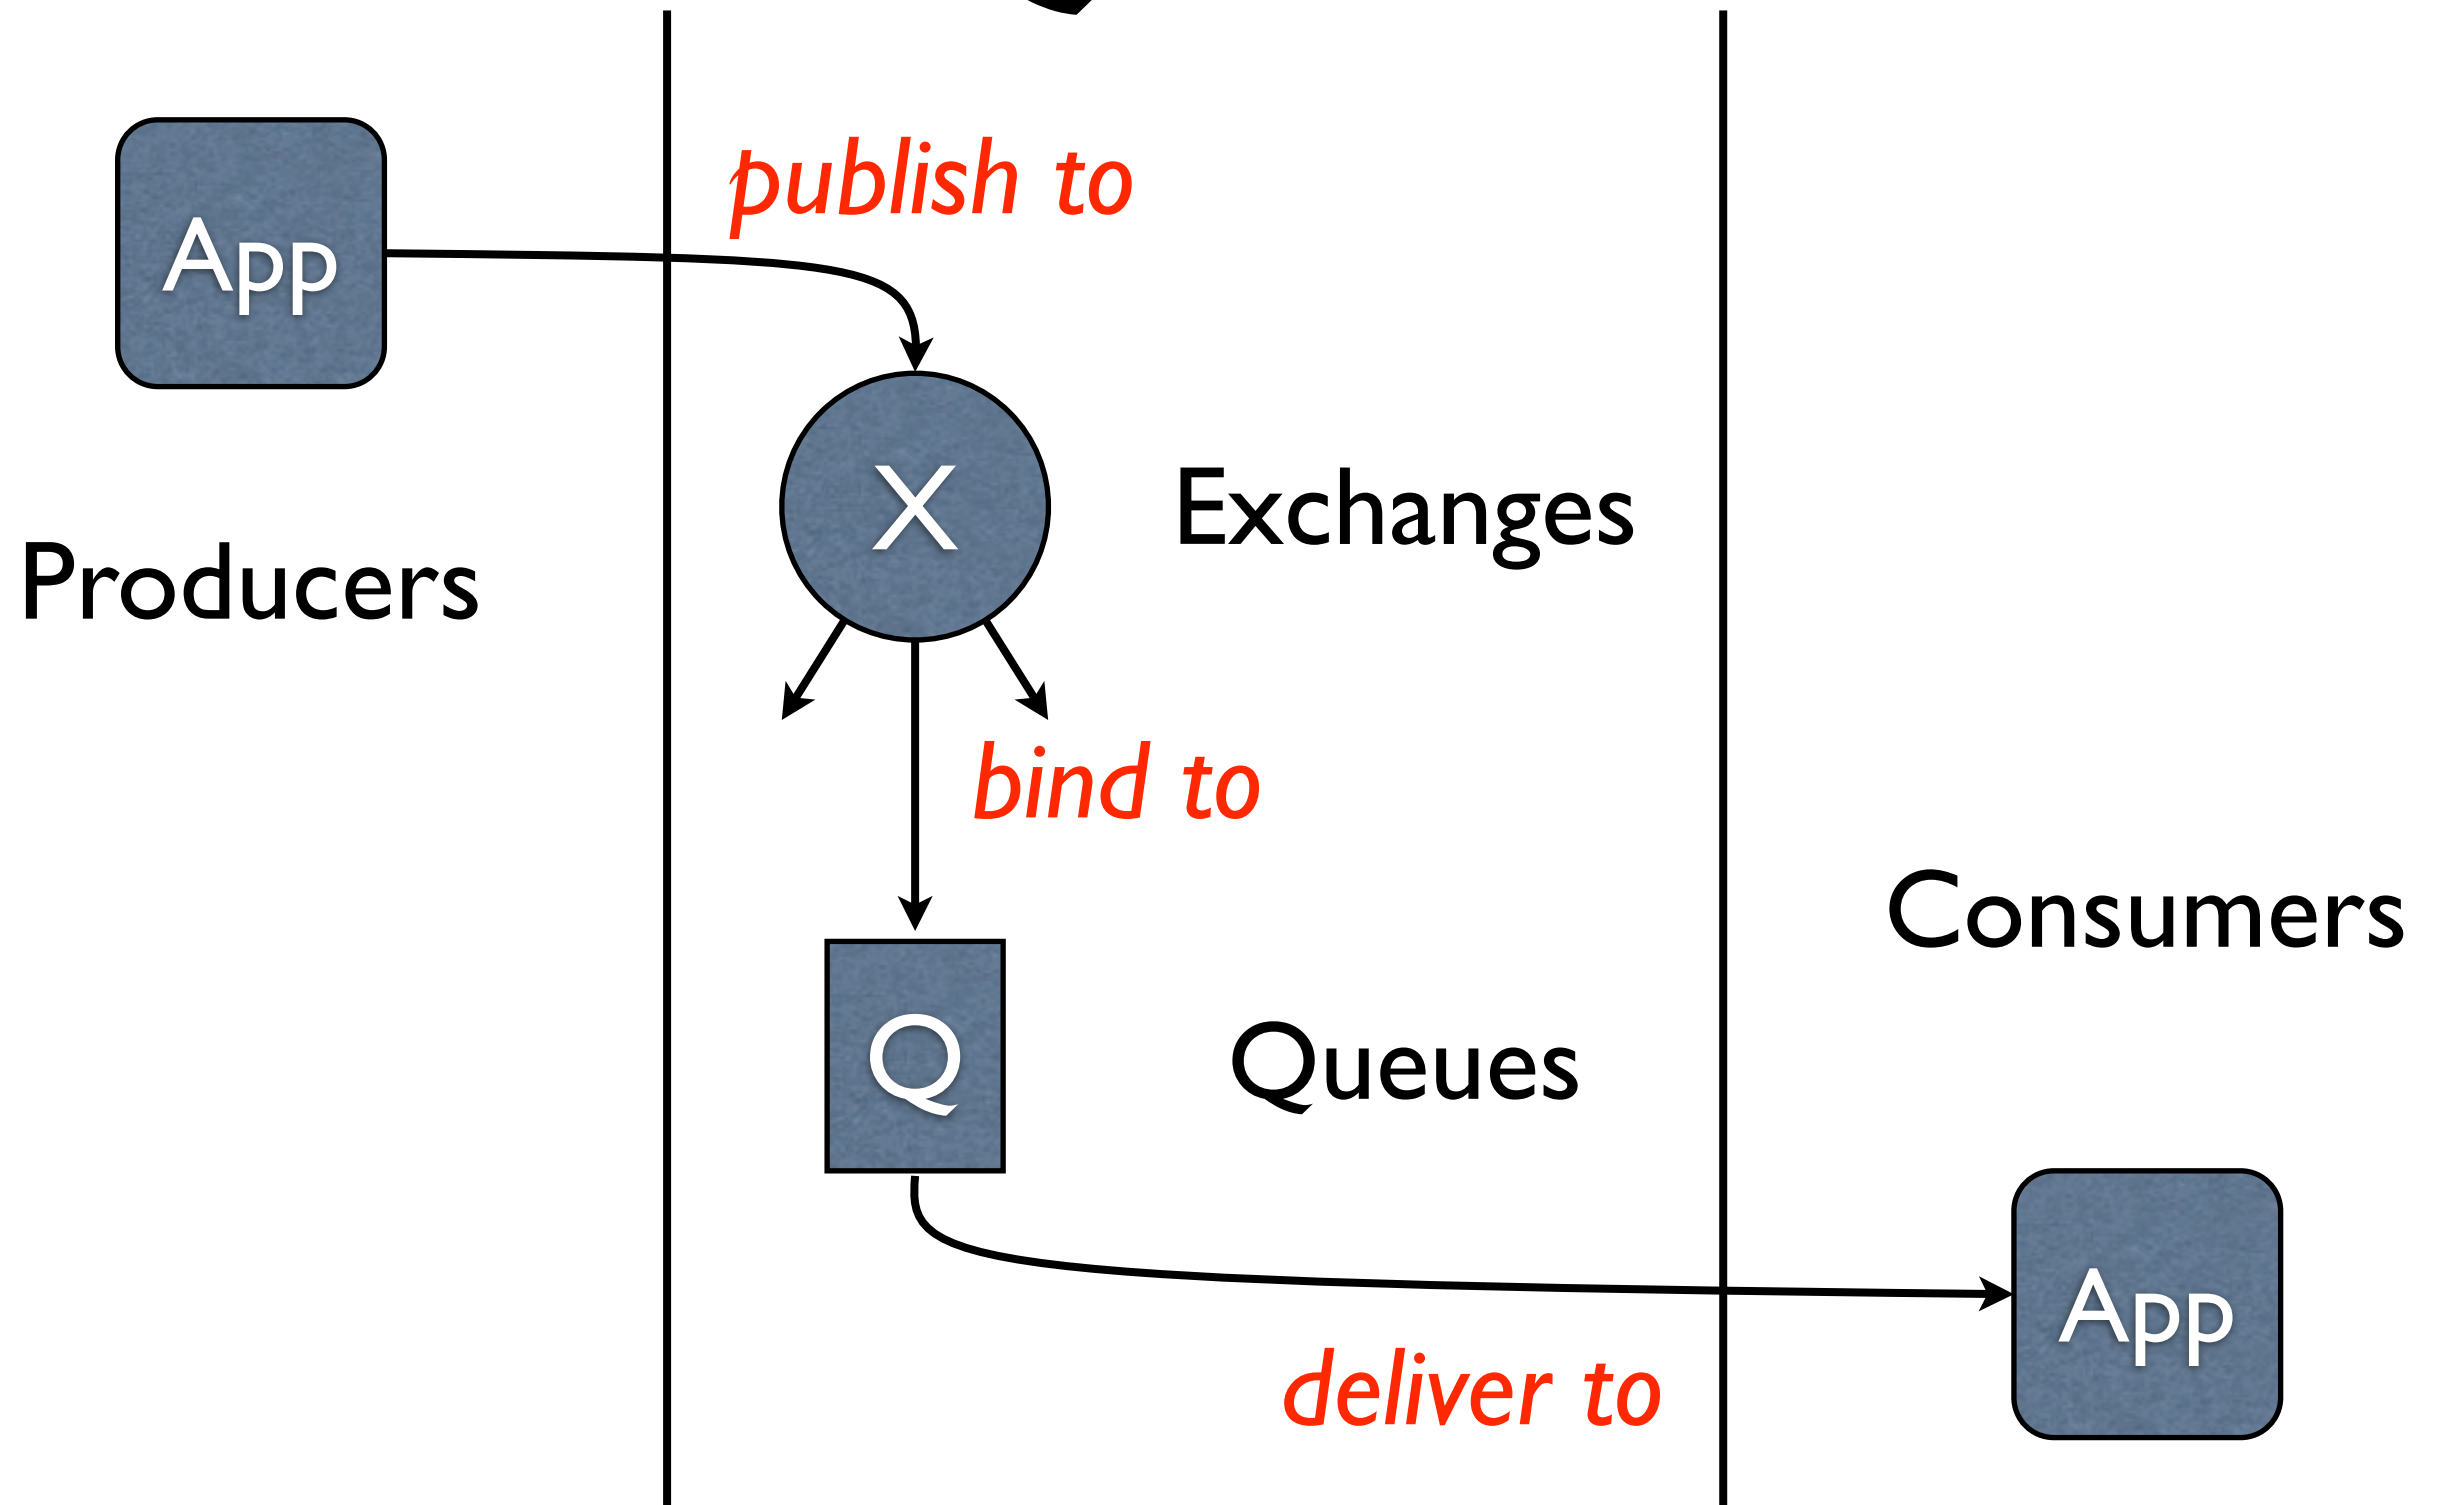


- Exchanges perform relaying, copying, and filtering

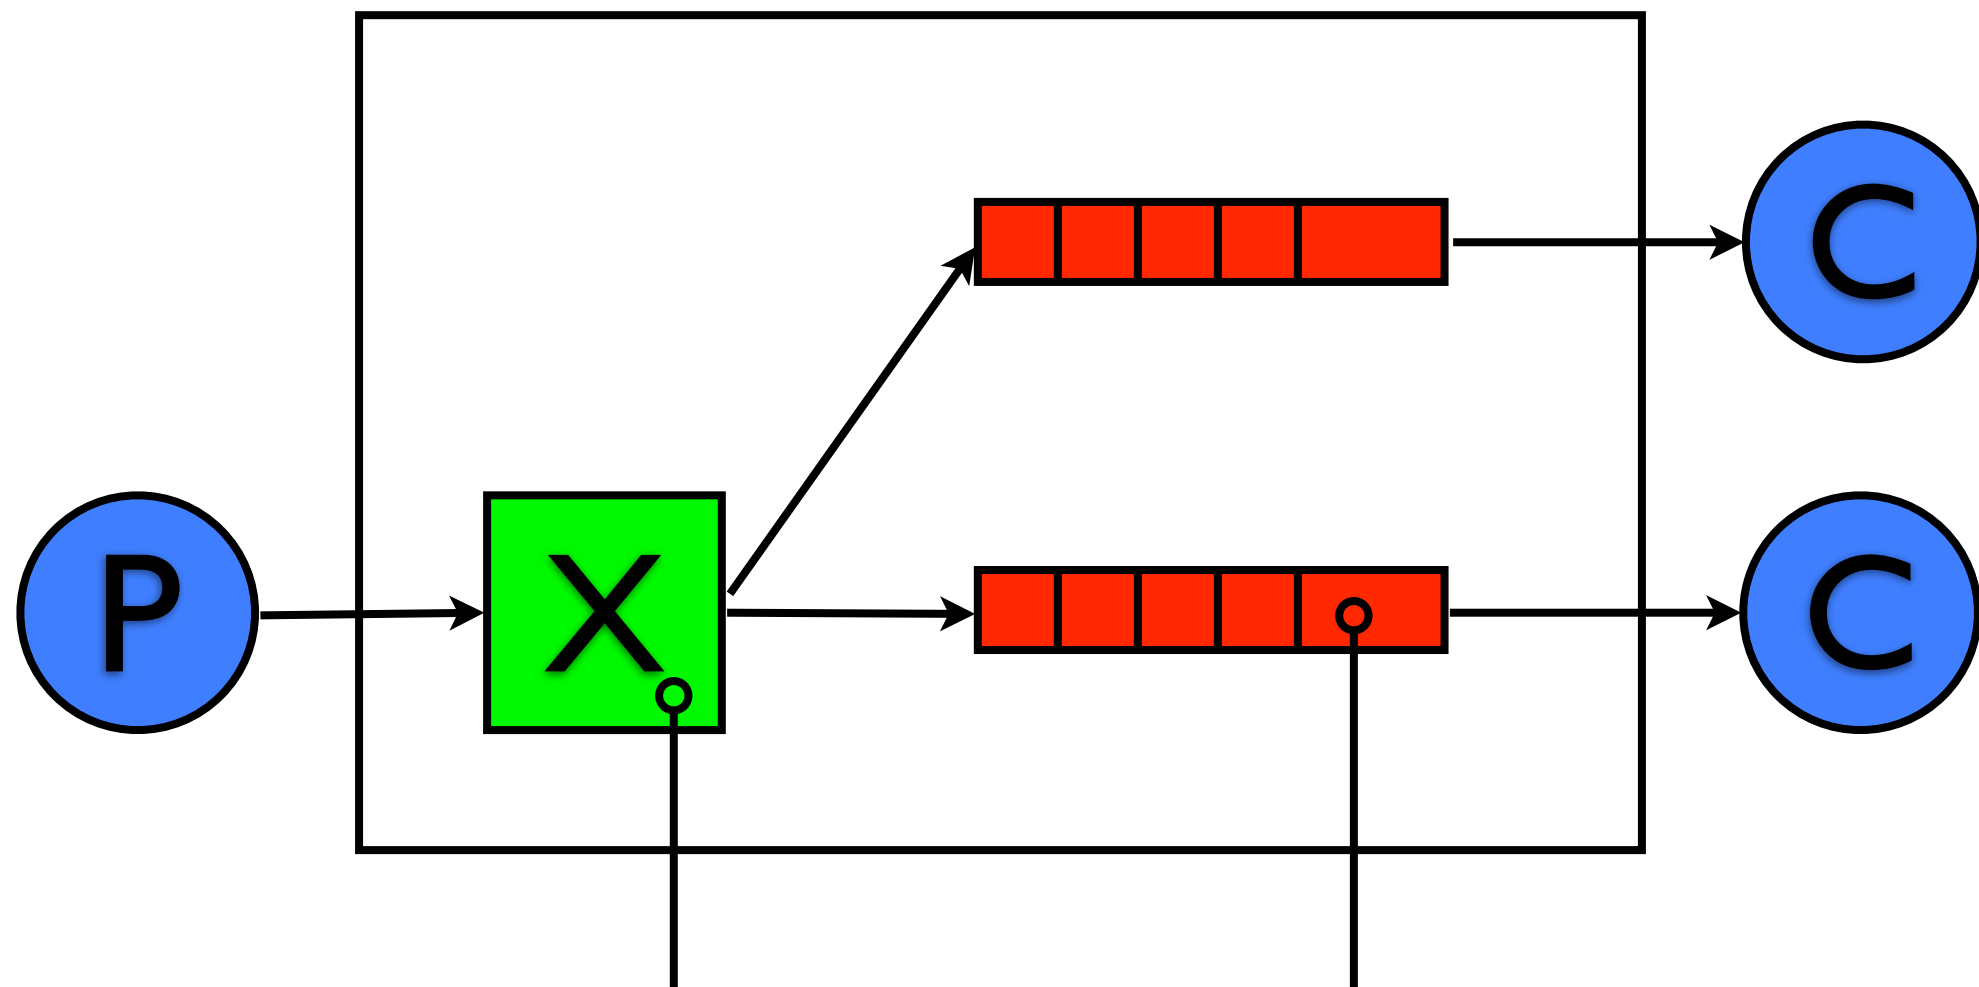
- Queues perform buffering and round-robin delivery

AMQP Basics

(0-9-1)

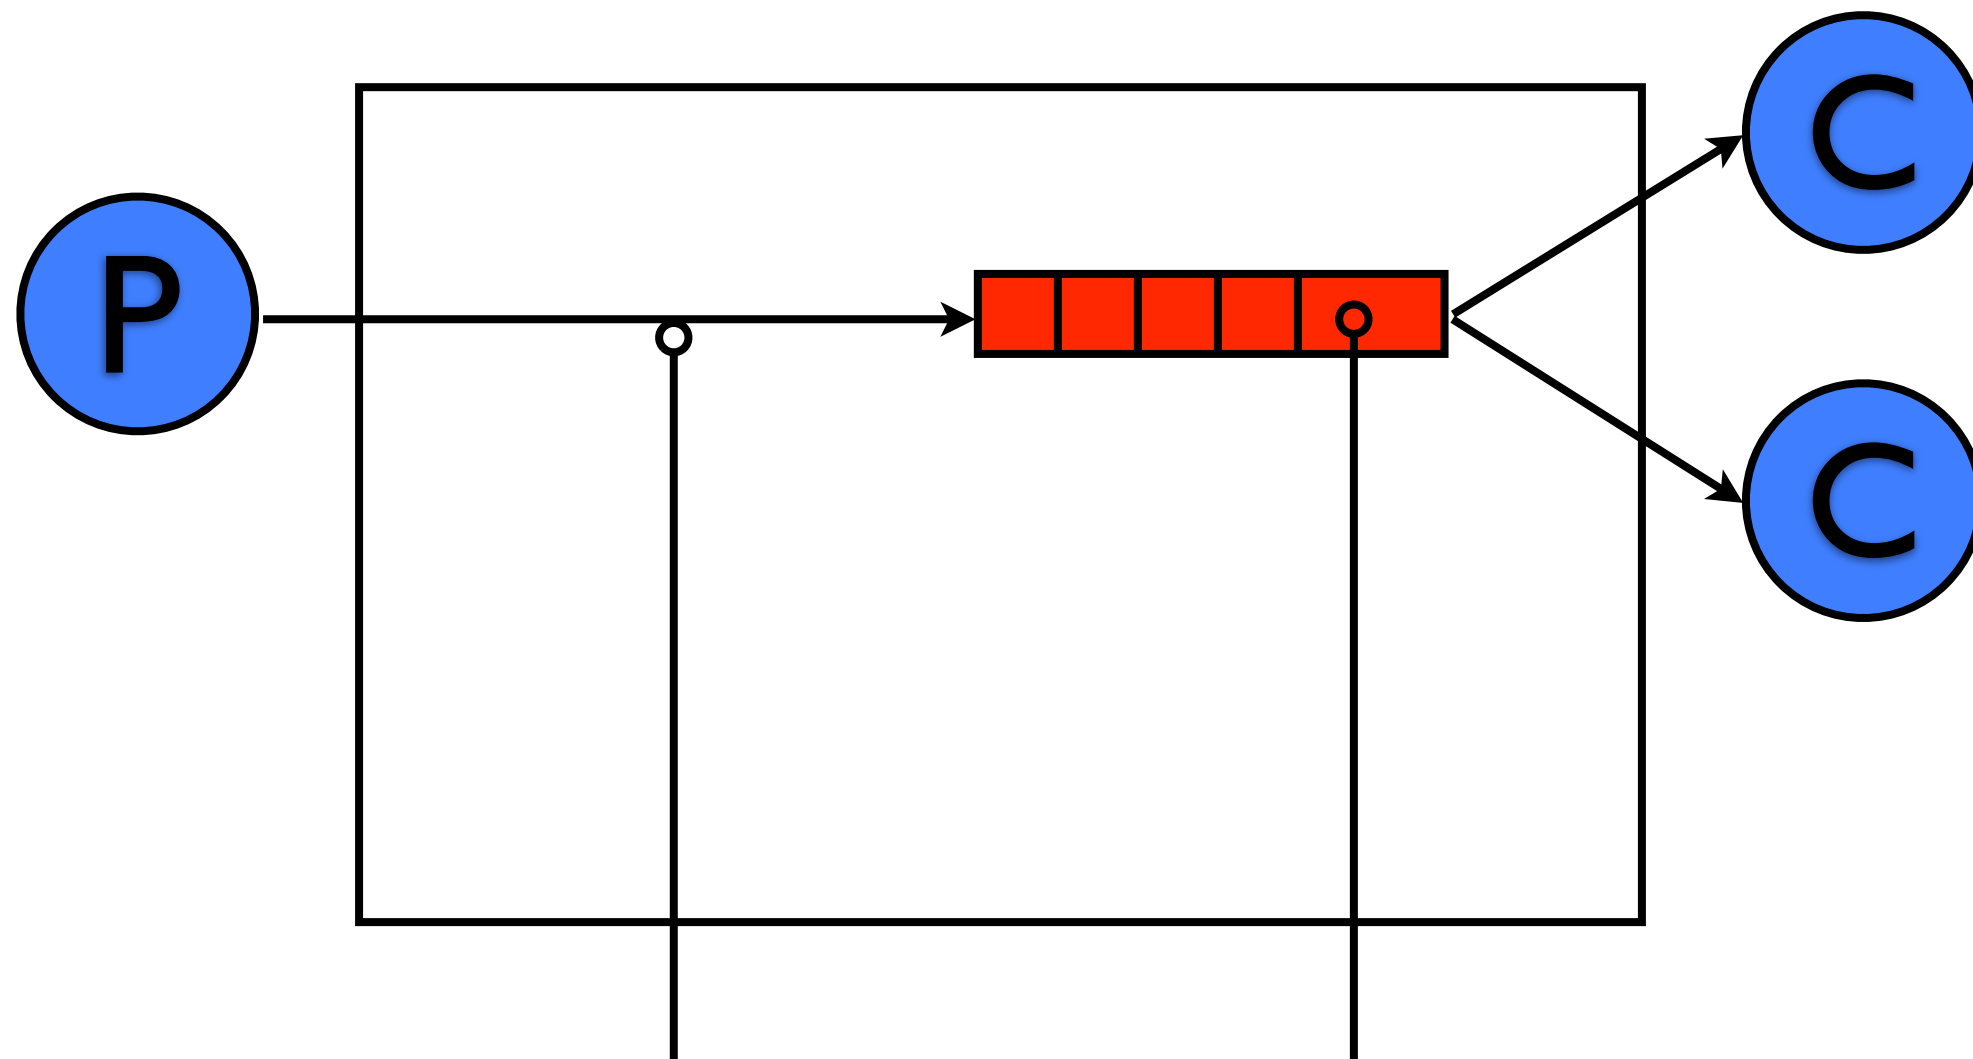


PubSub



- AMQP exchange used as the destination copies to multiple queues
- Per-consumer private queues receive topical messages

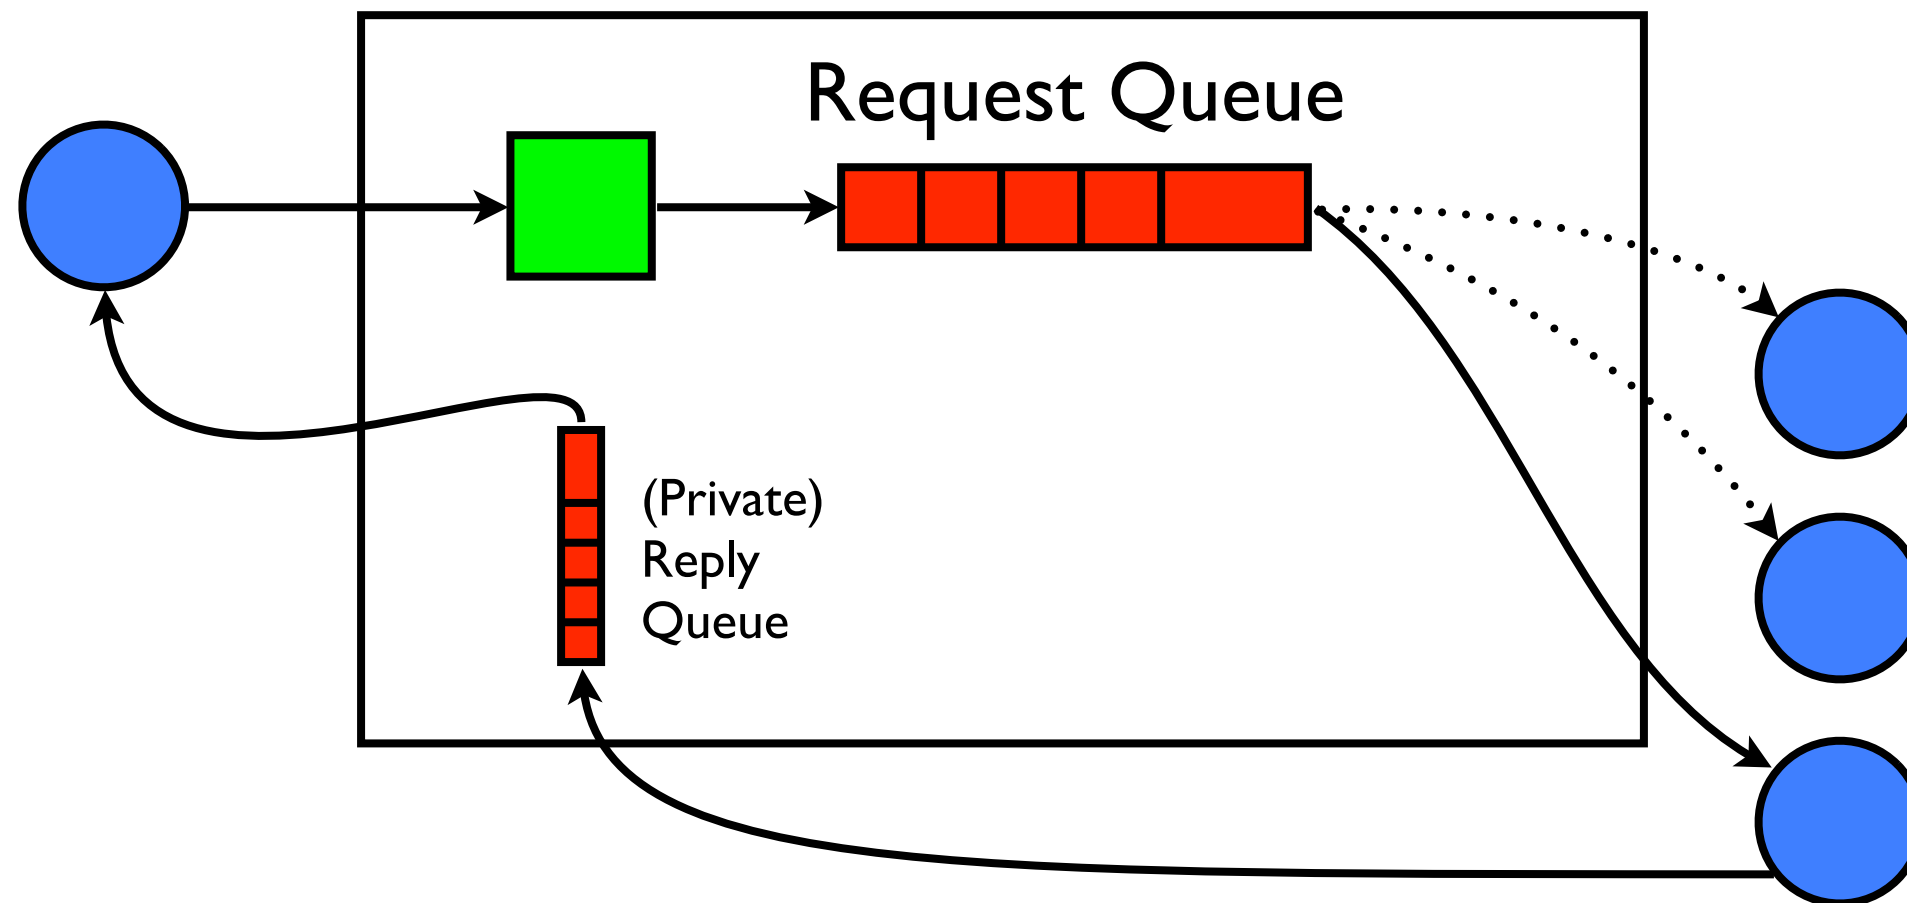
Queueing



- Using the default exchange (“”) routes directly to queues

- Both shared and private queues can be addressed like this

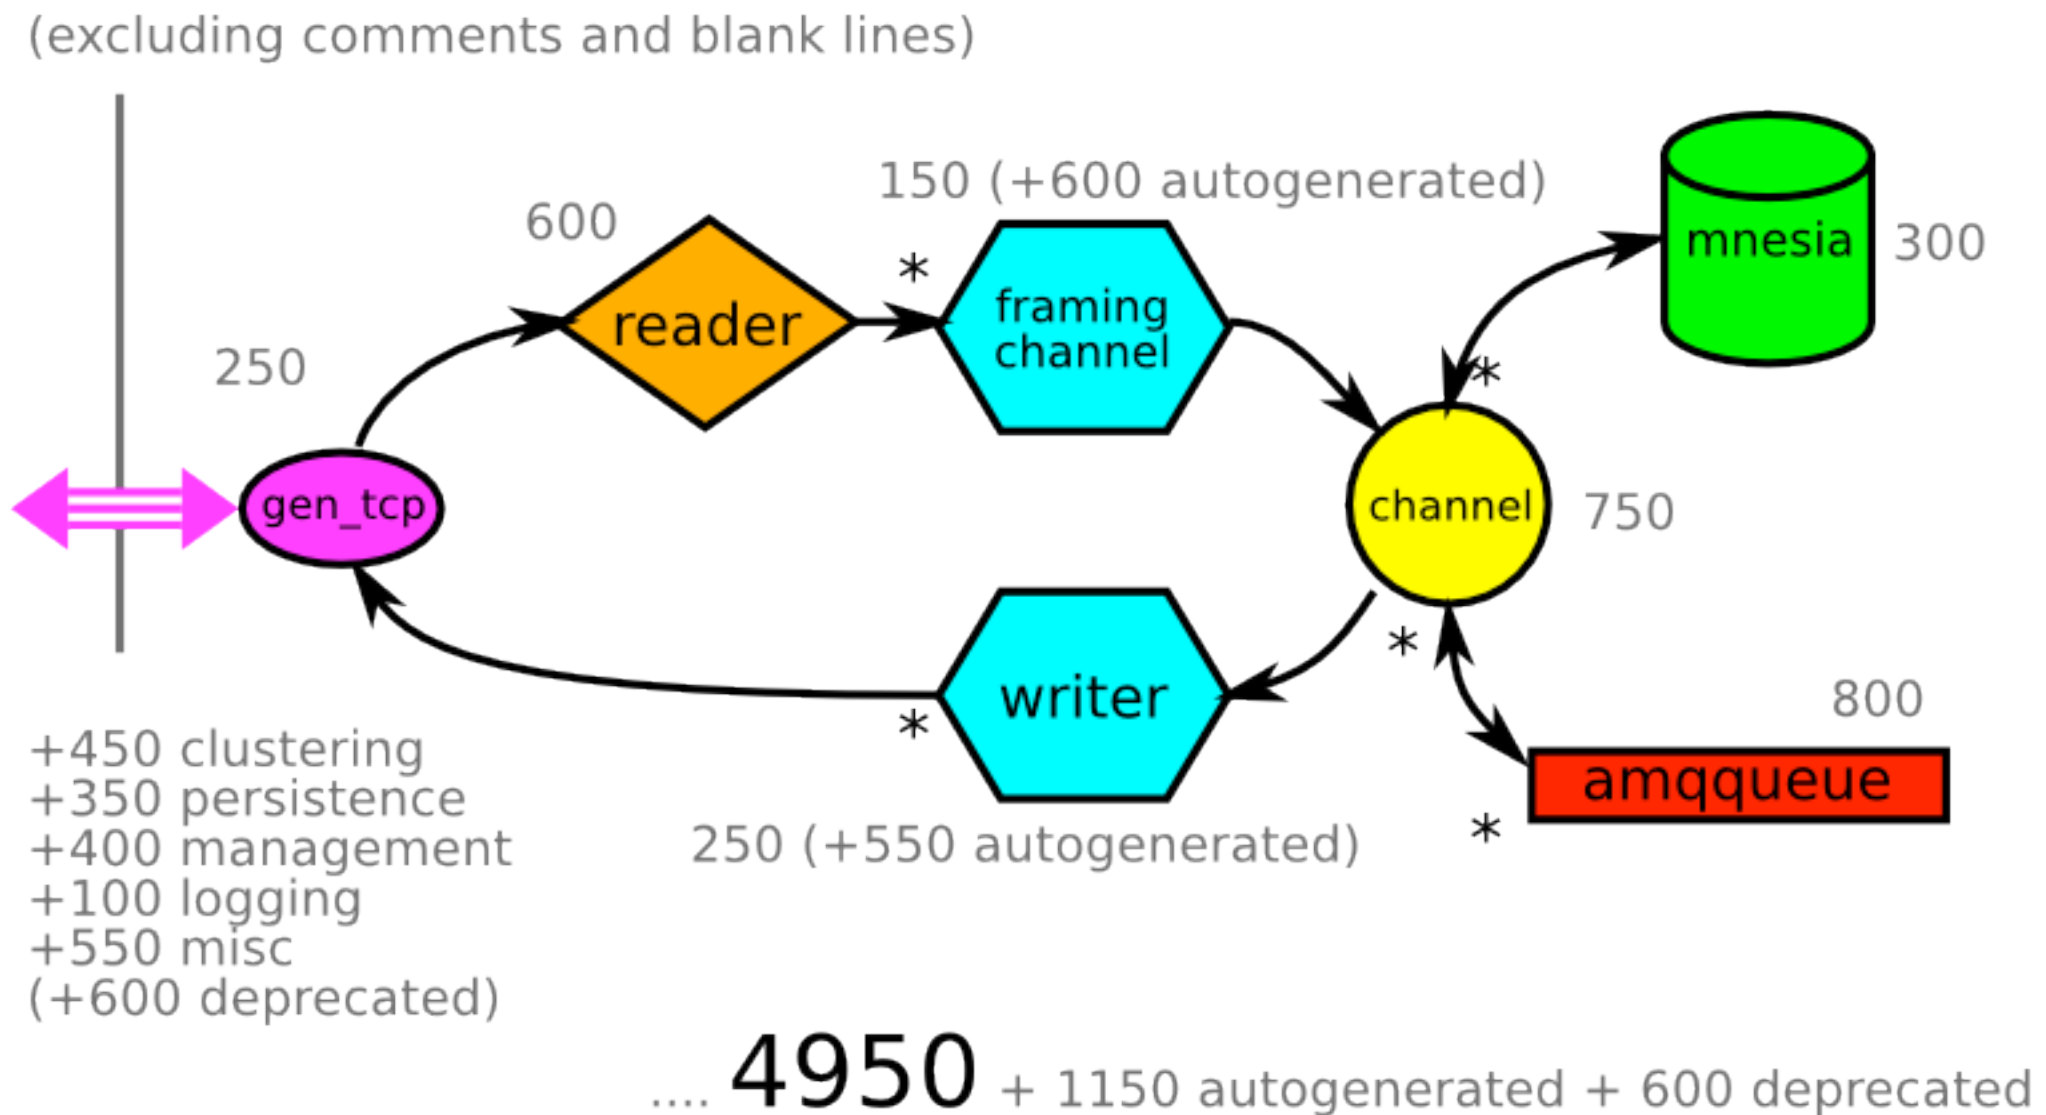
Simple load-balancing



- Shared queue mediates access to service instances
- Load-balancing, live upgrades, fault-tolerance

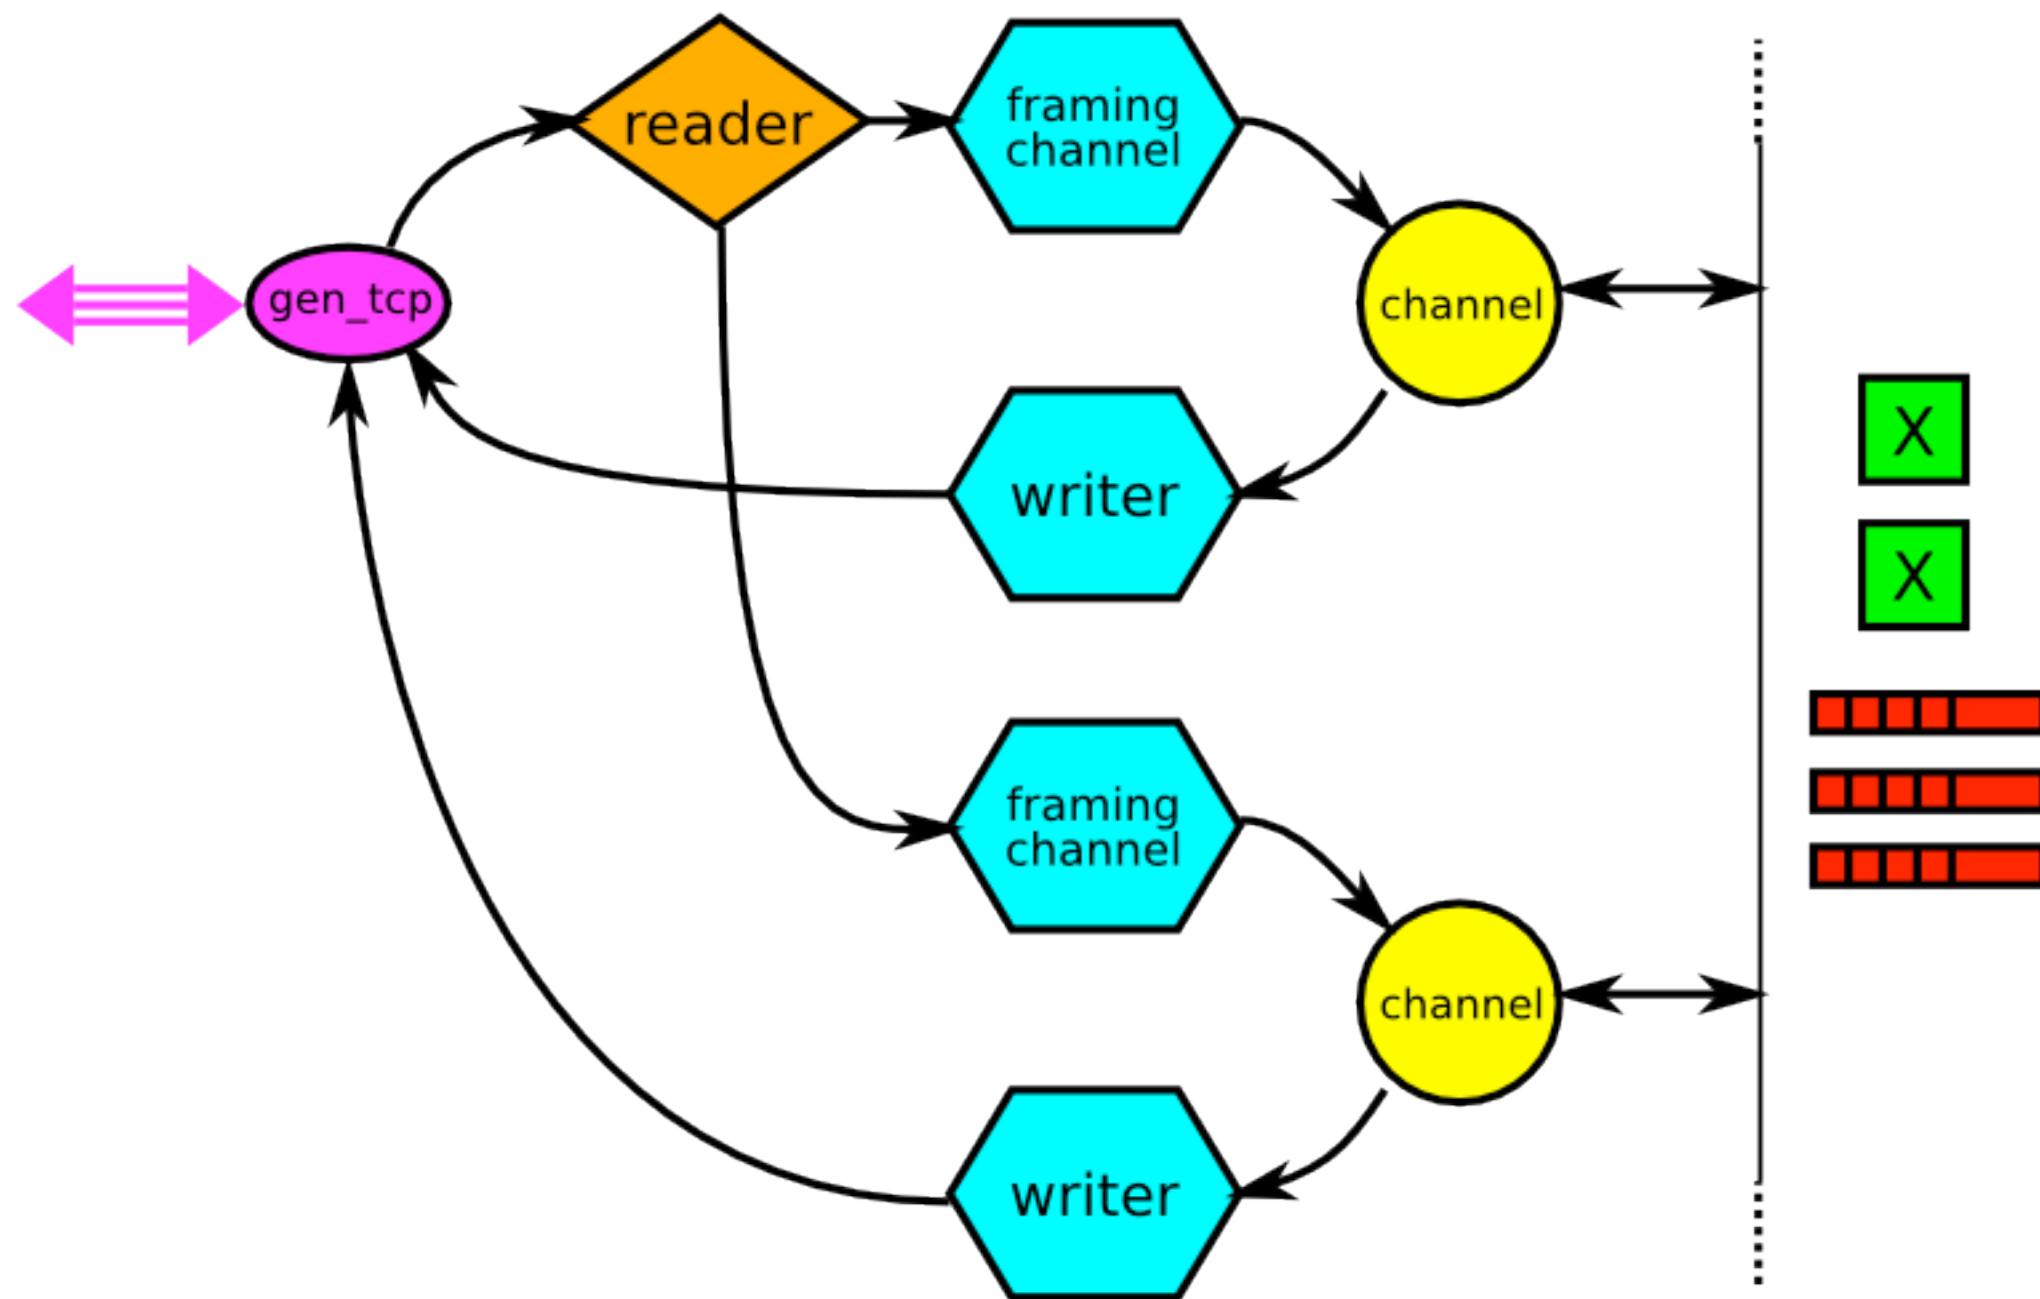
RabbitMQ

Code Overview

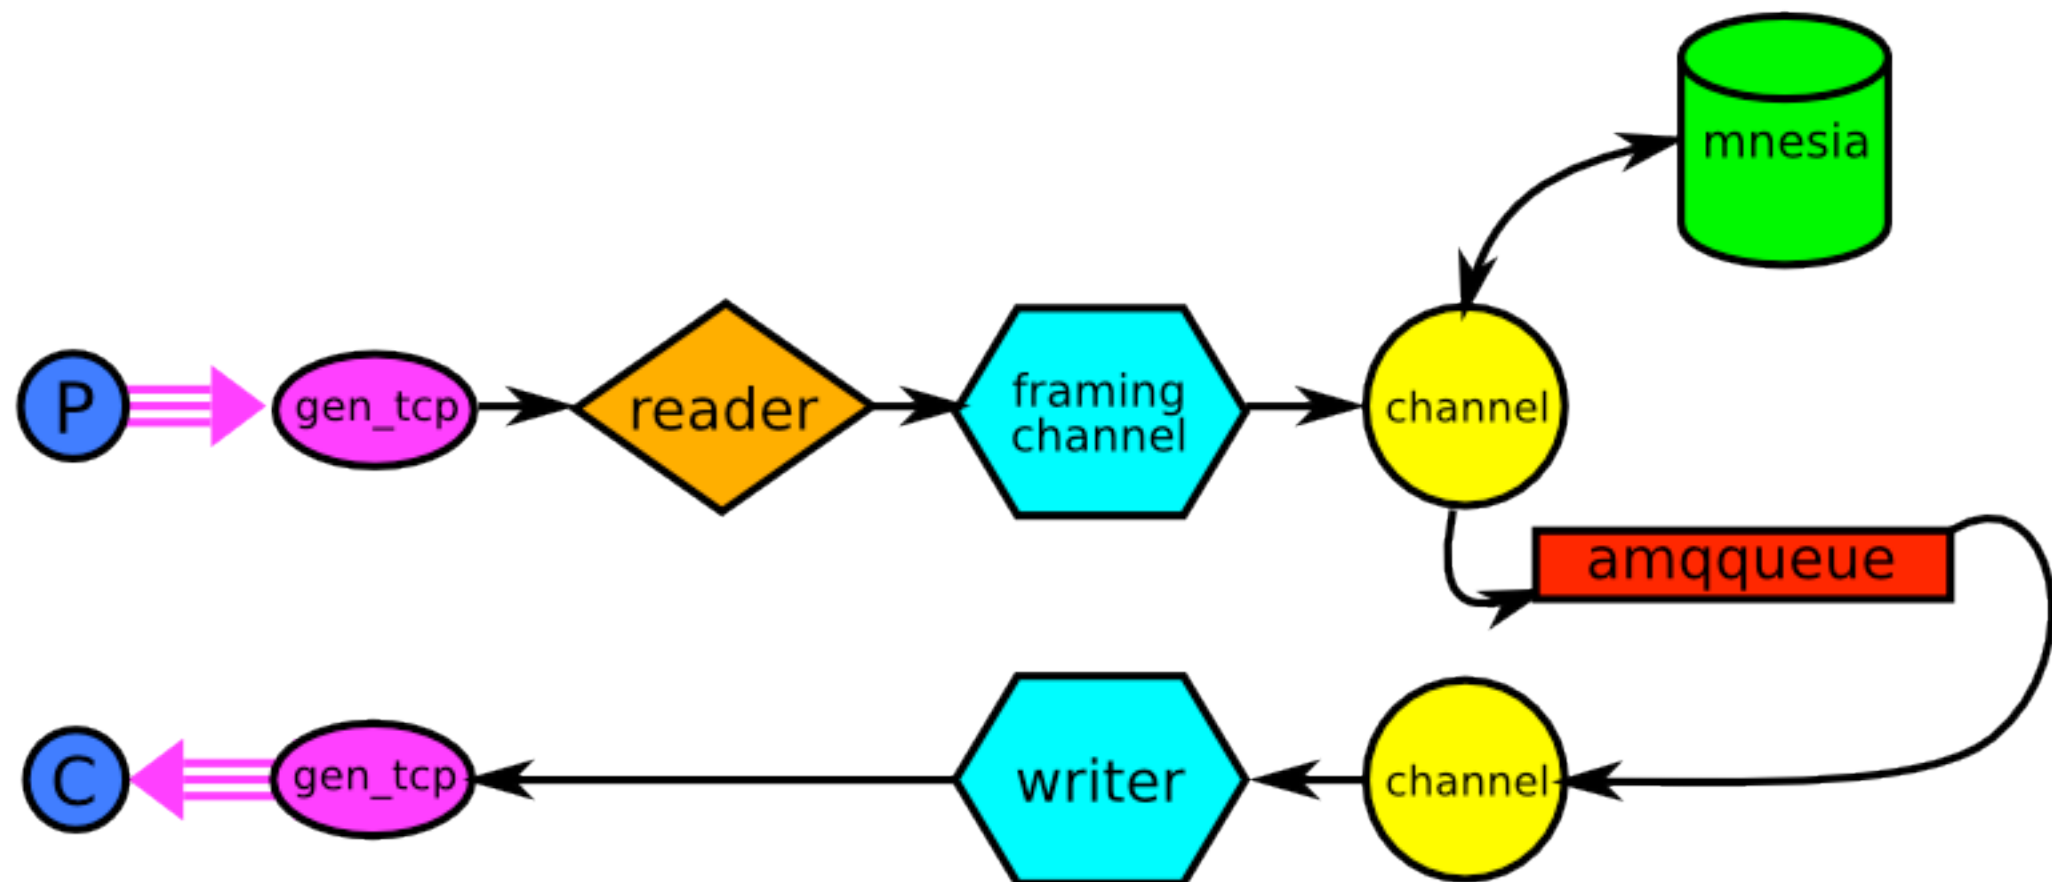
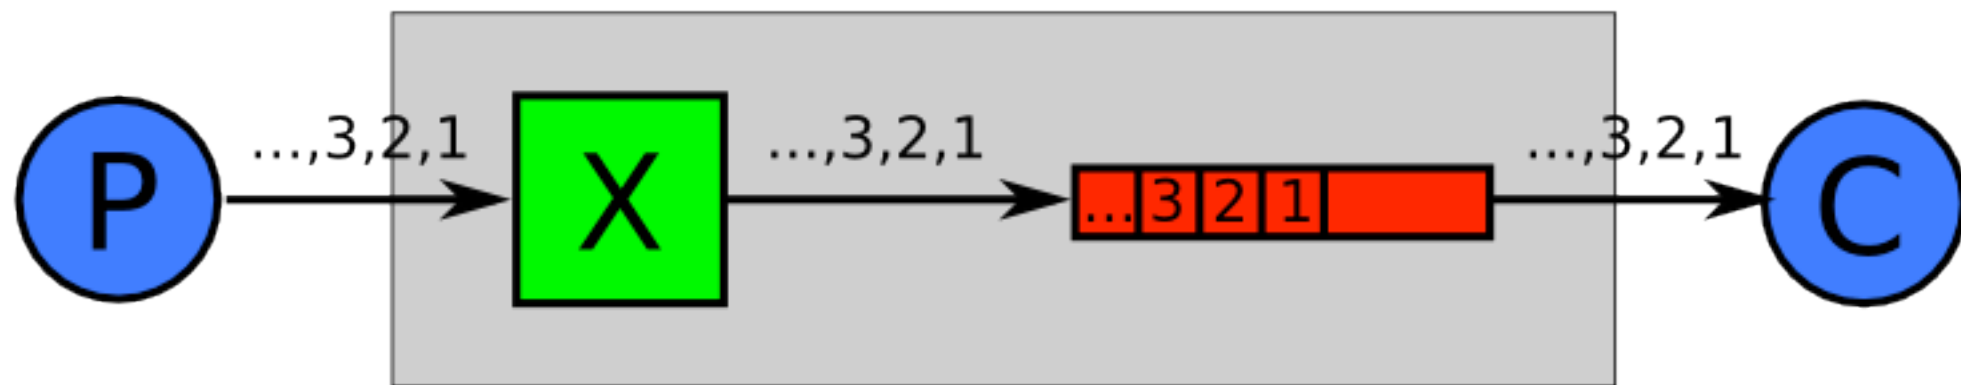


(out of date numbers! it's ~7500 LoC now, mid 2009)

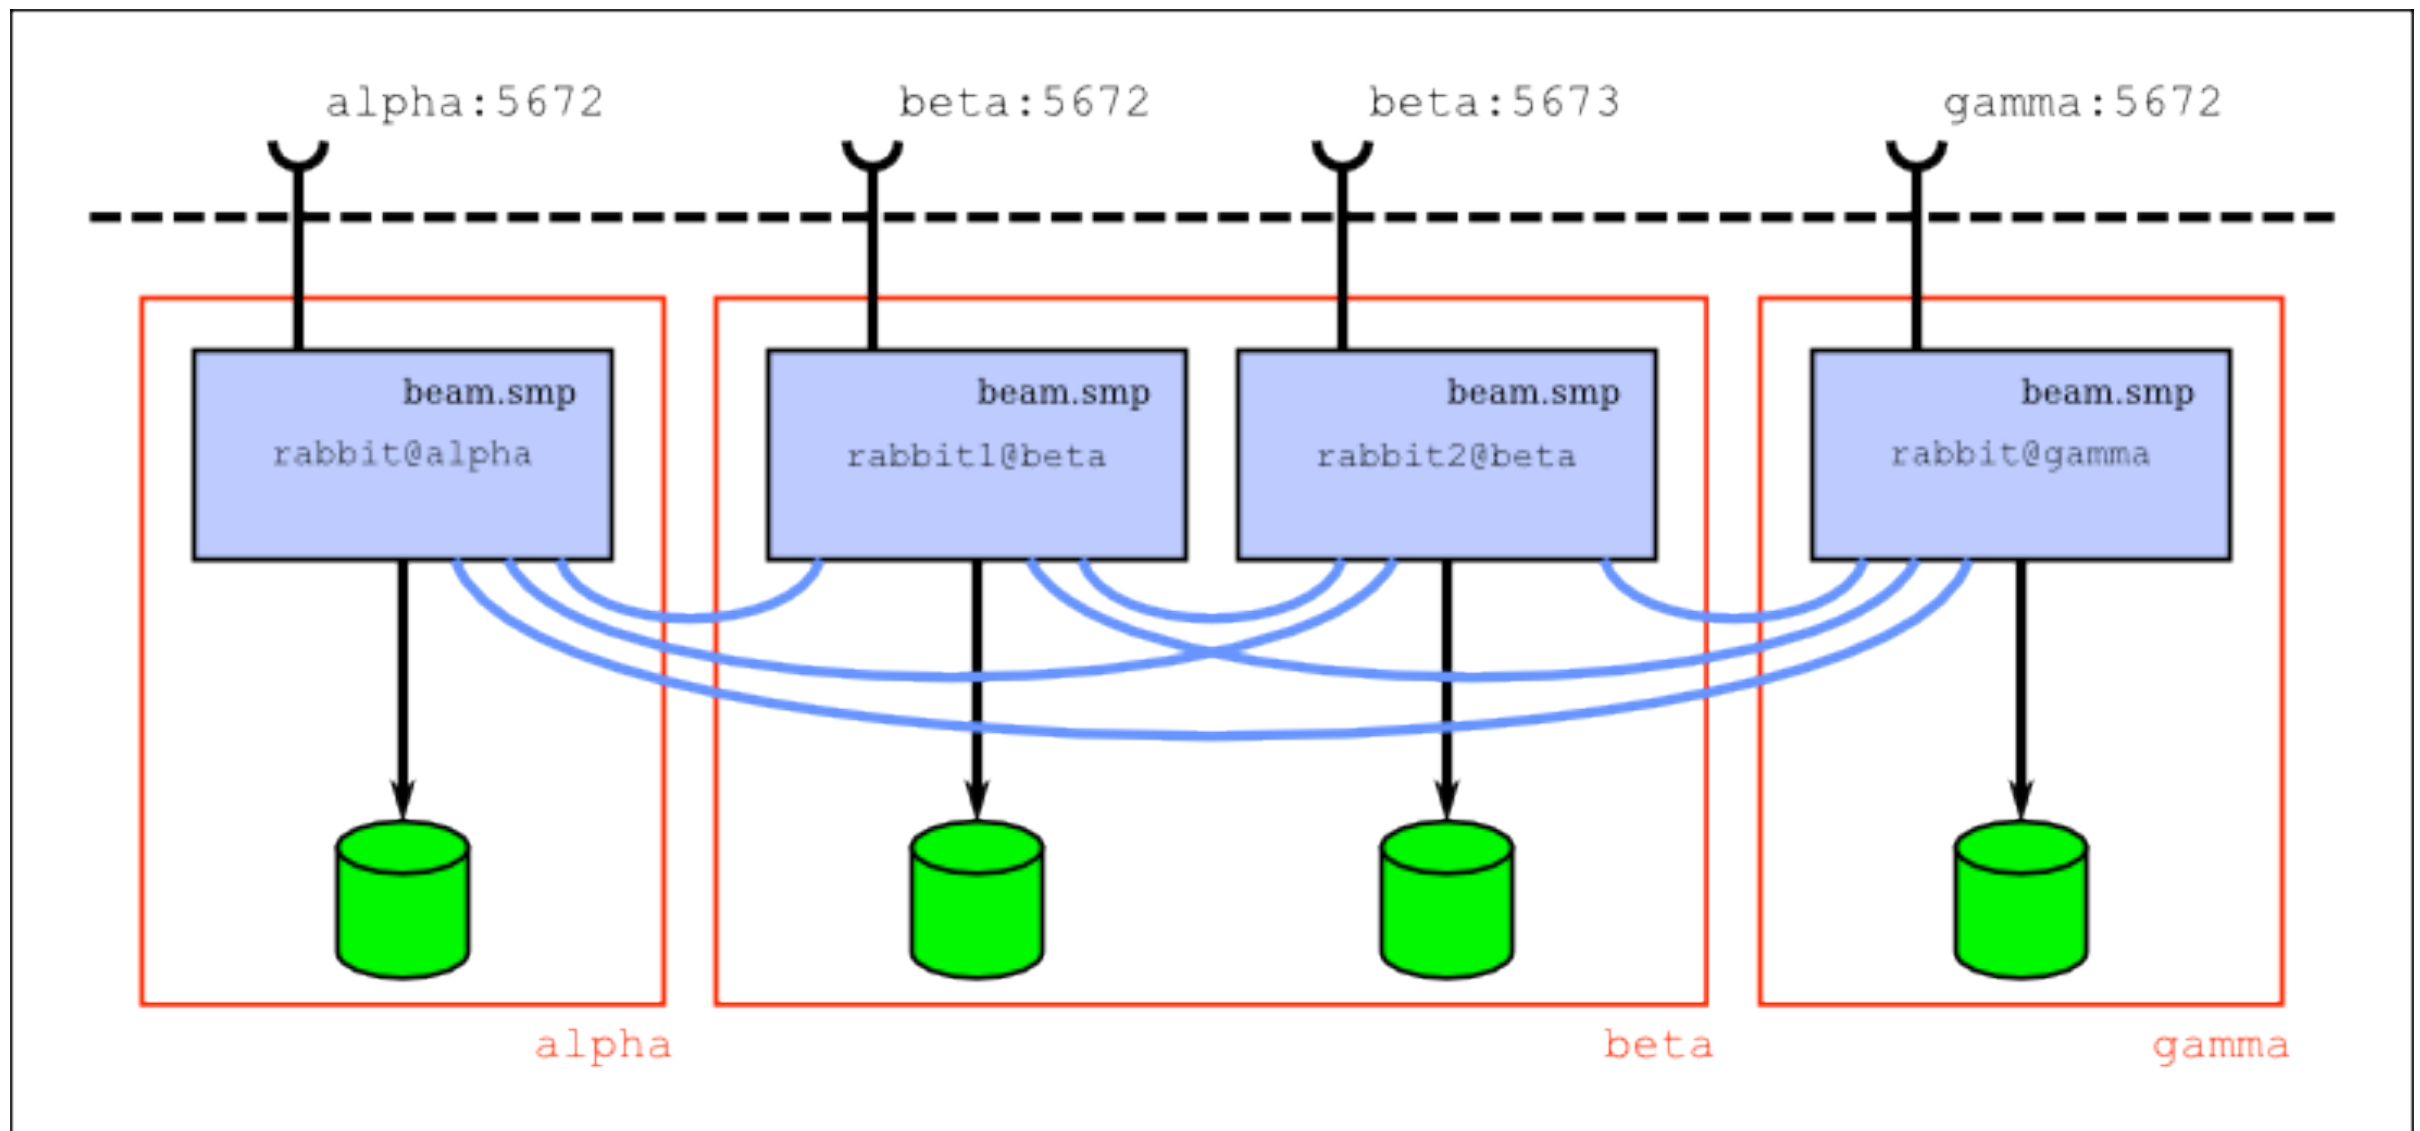
Concurrency



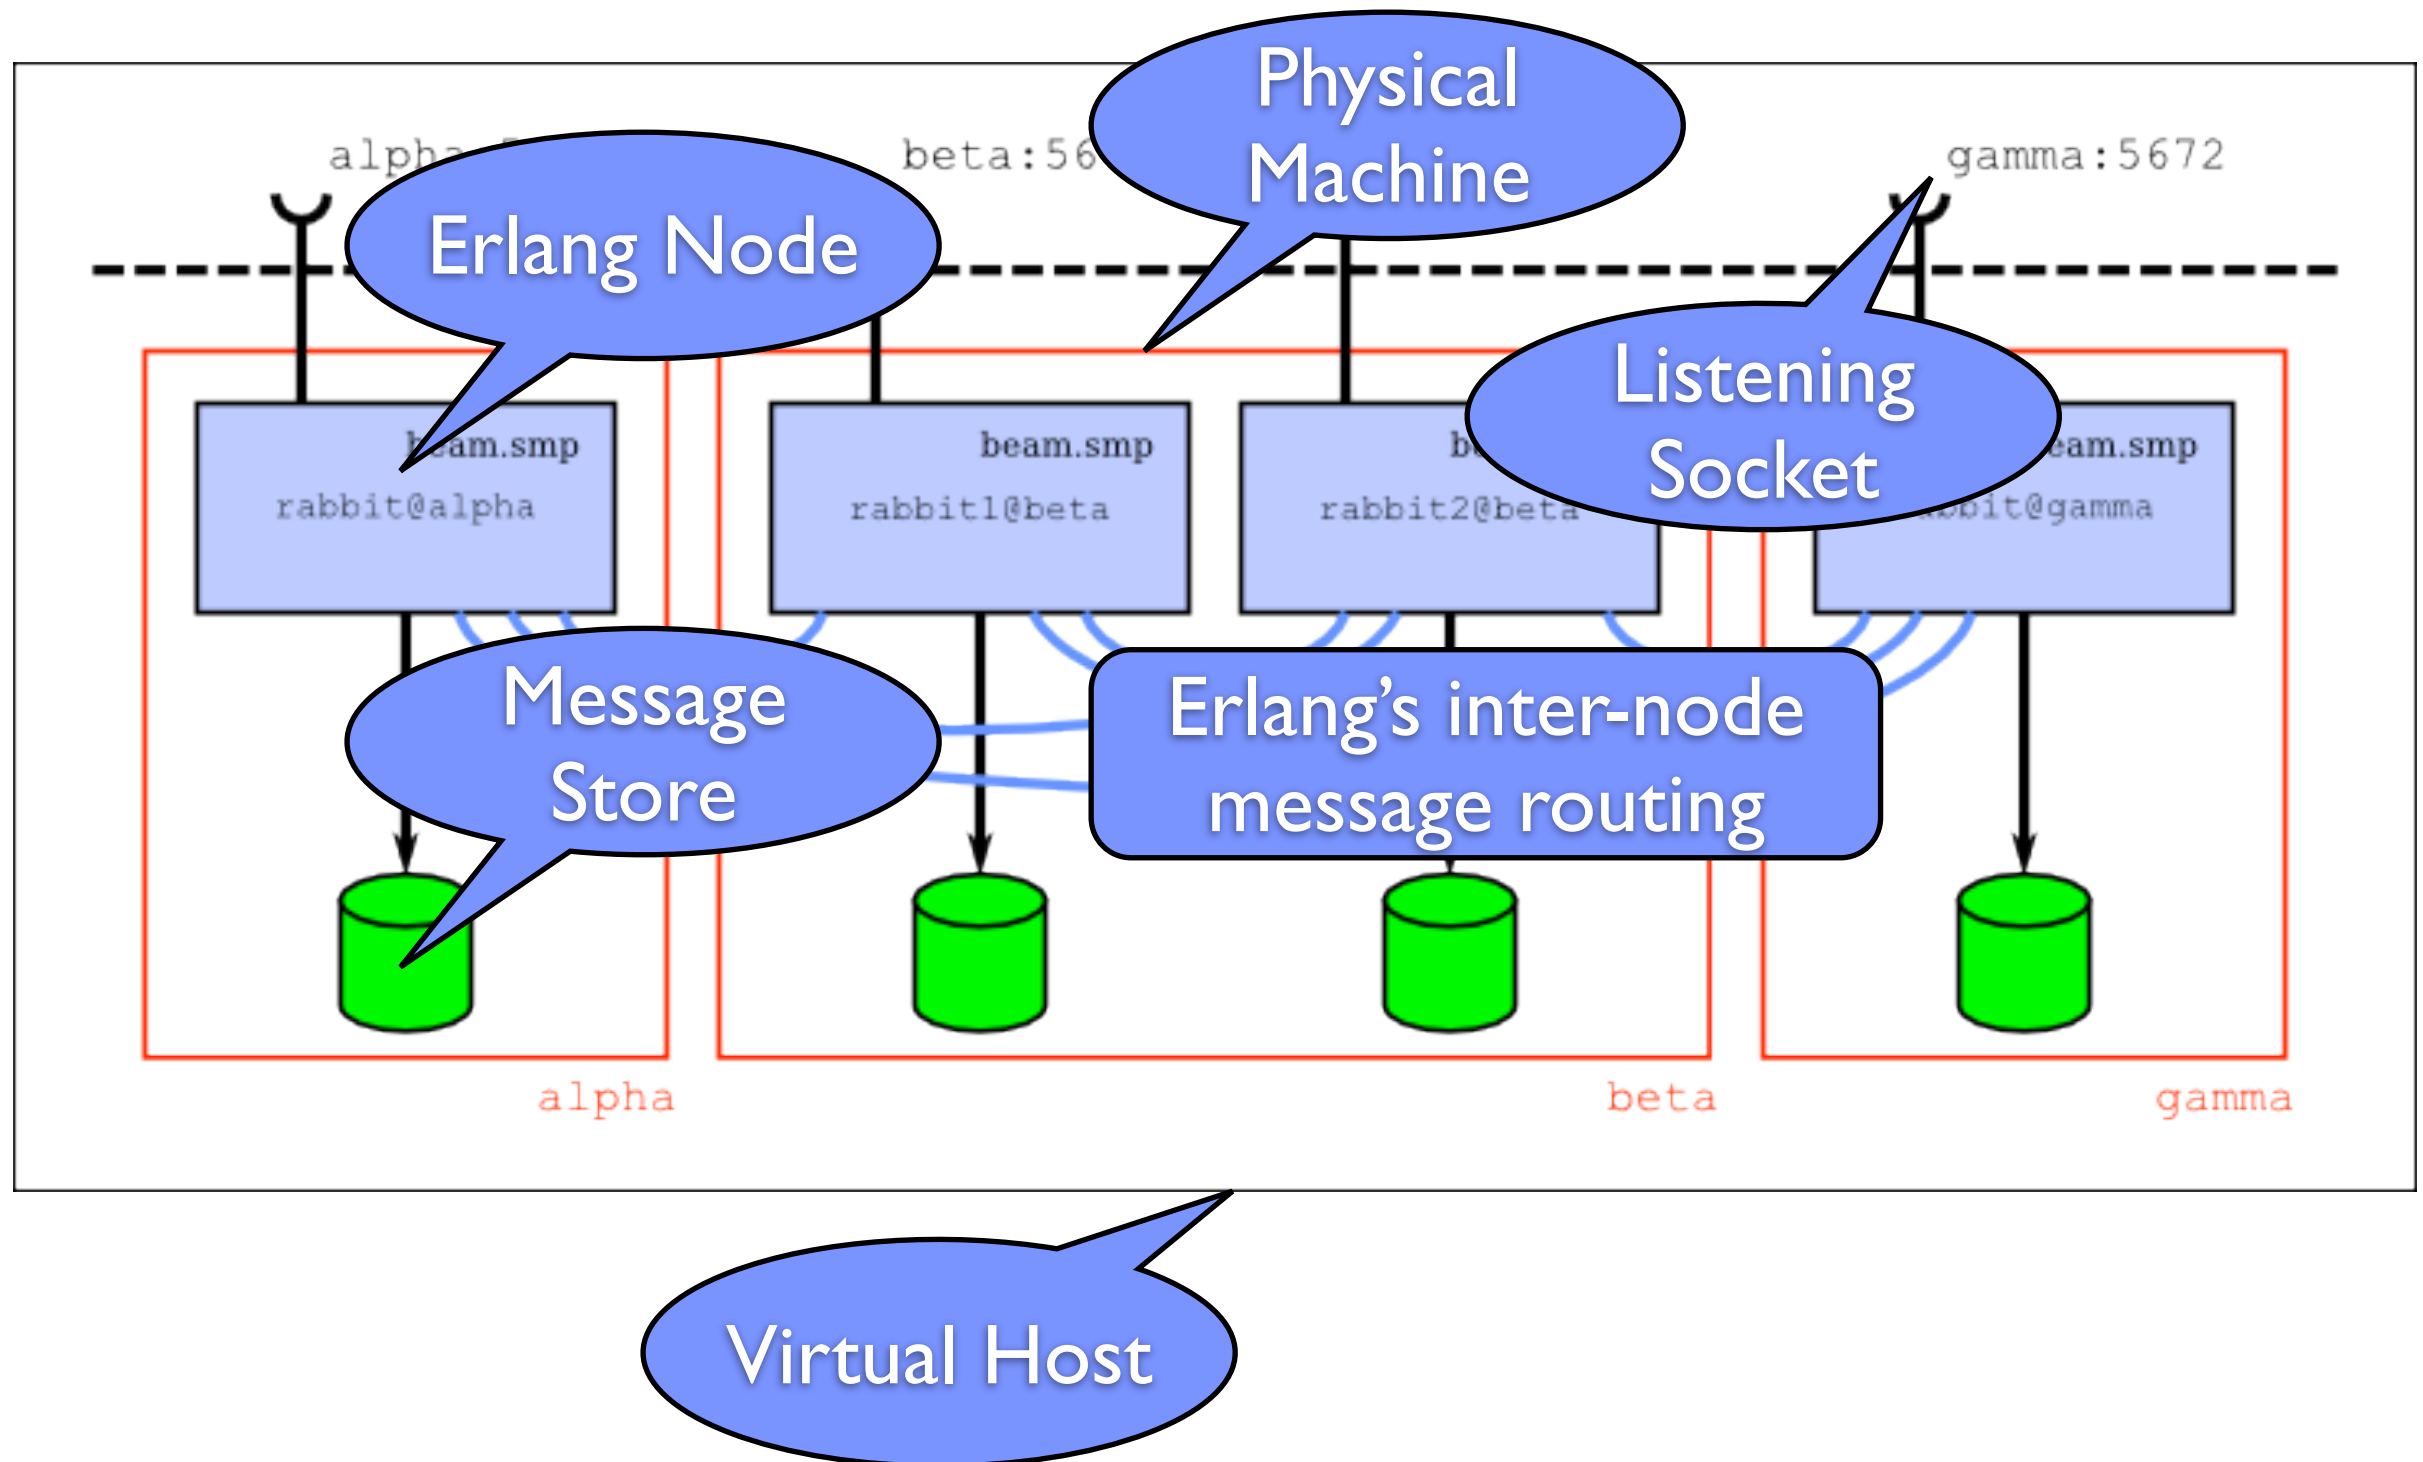
Order Preservation



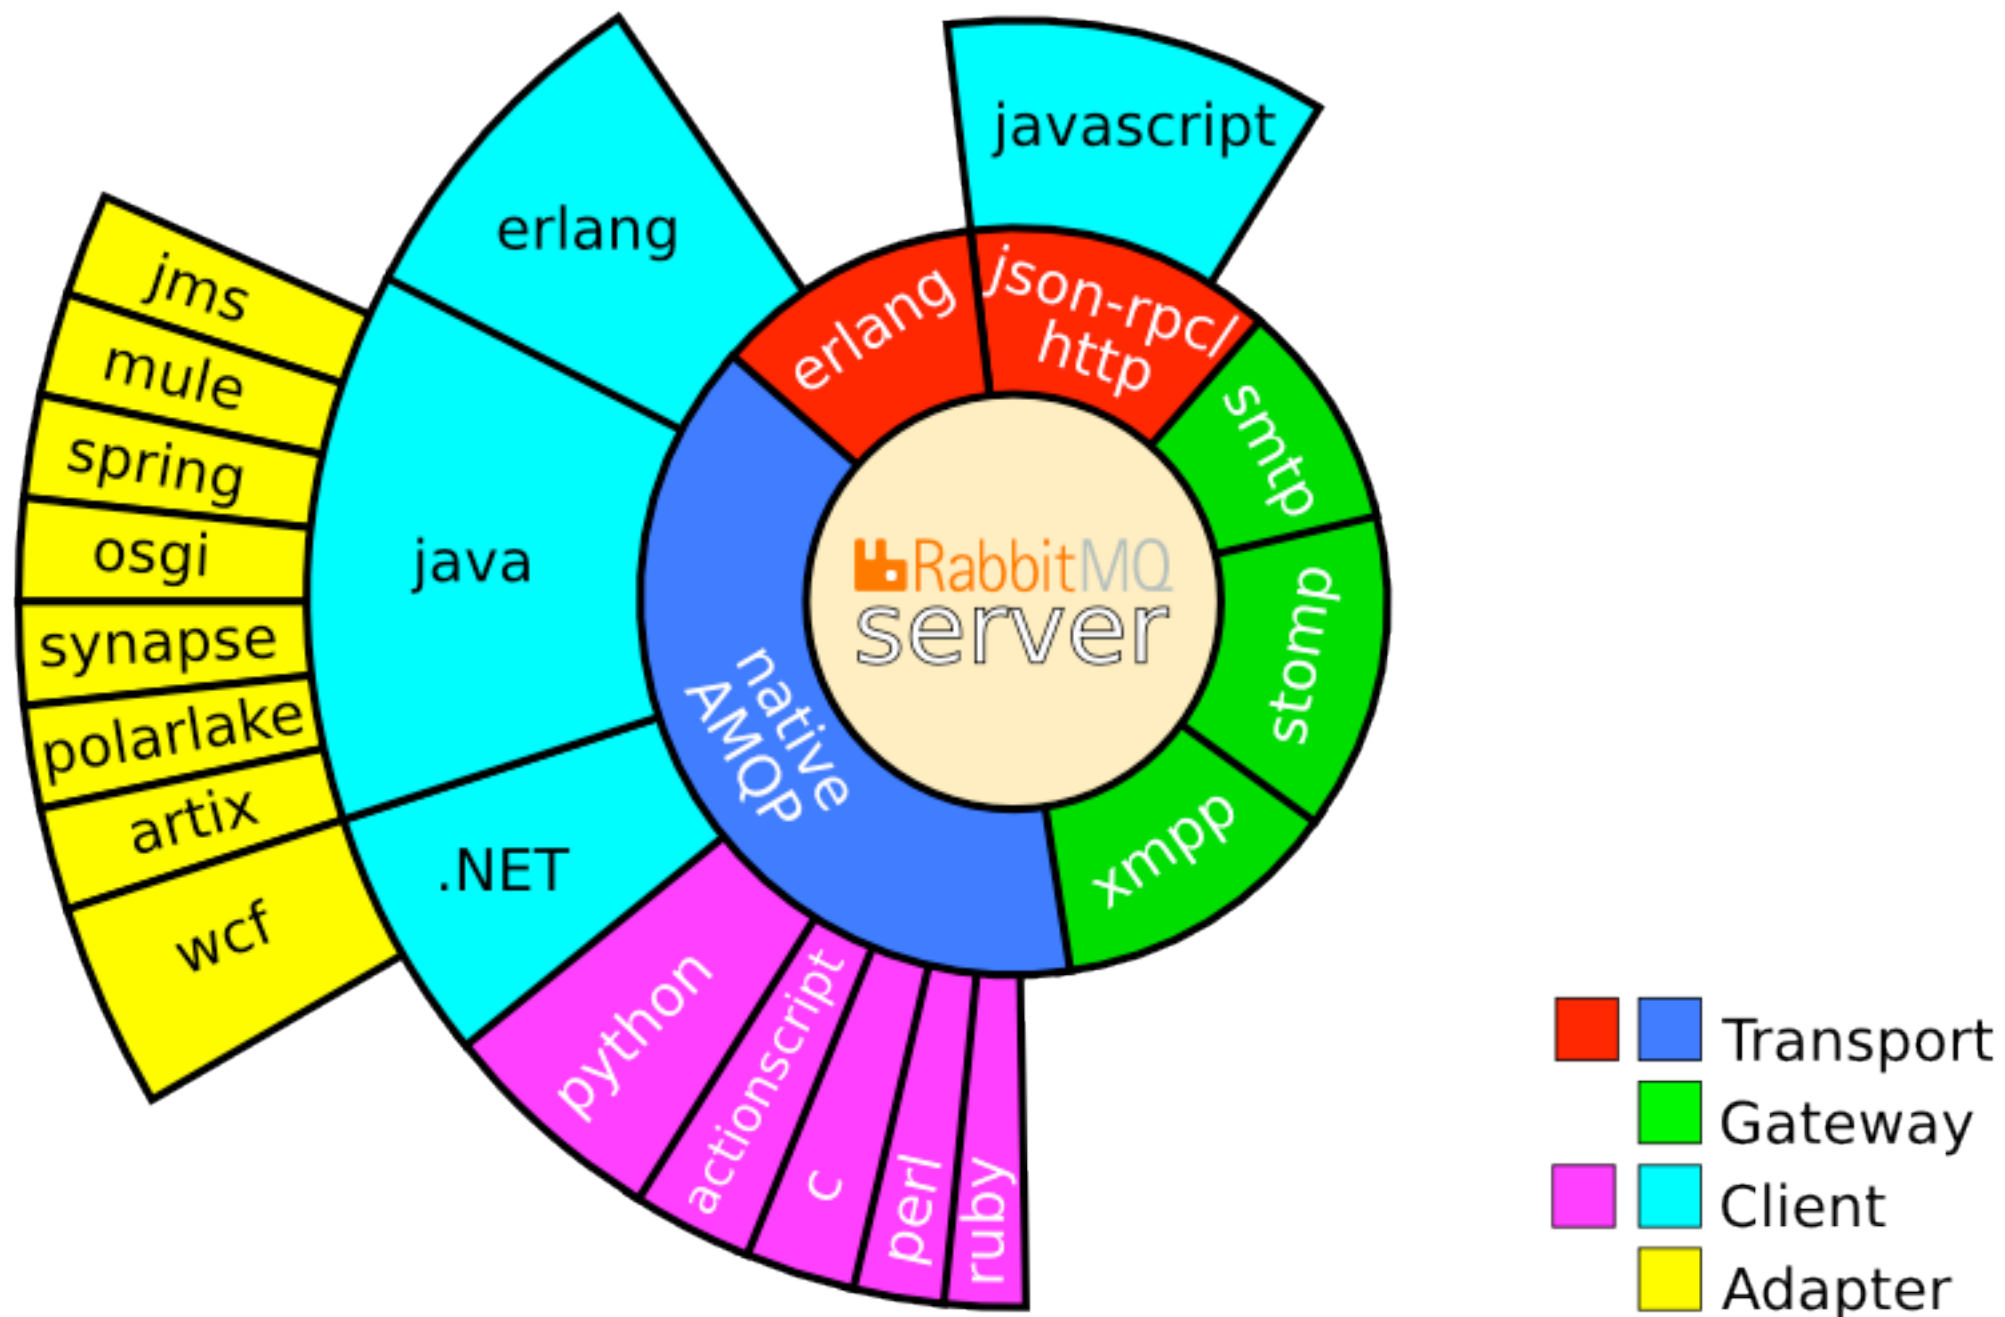
Clustering



Clustering



Connectivity



Availability, Clustering, Federation

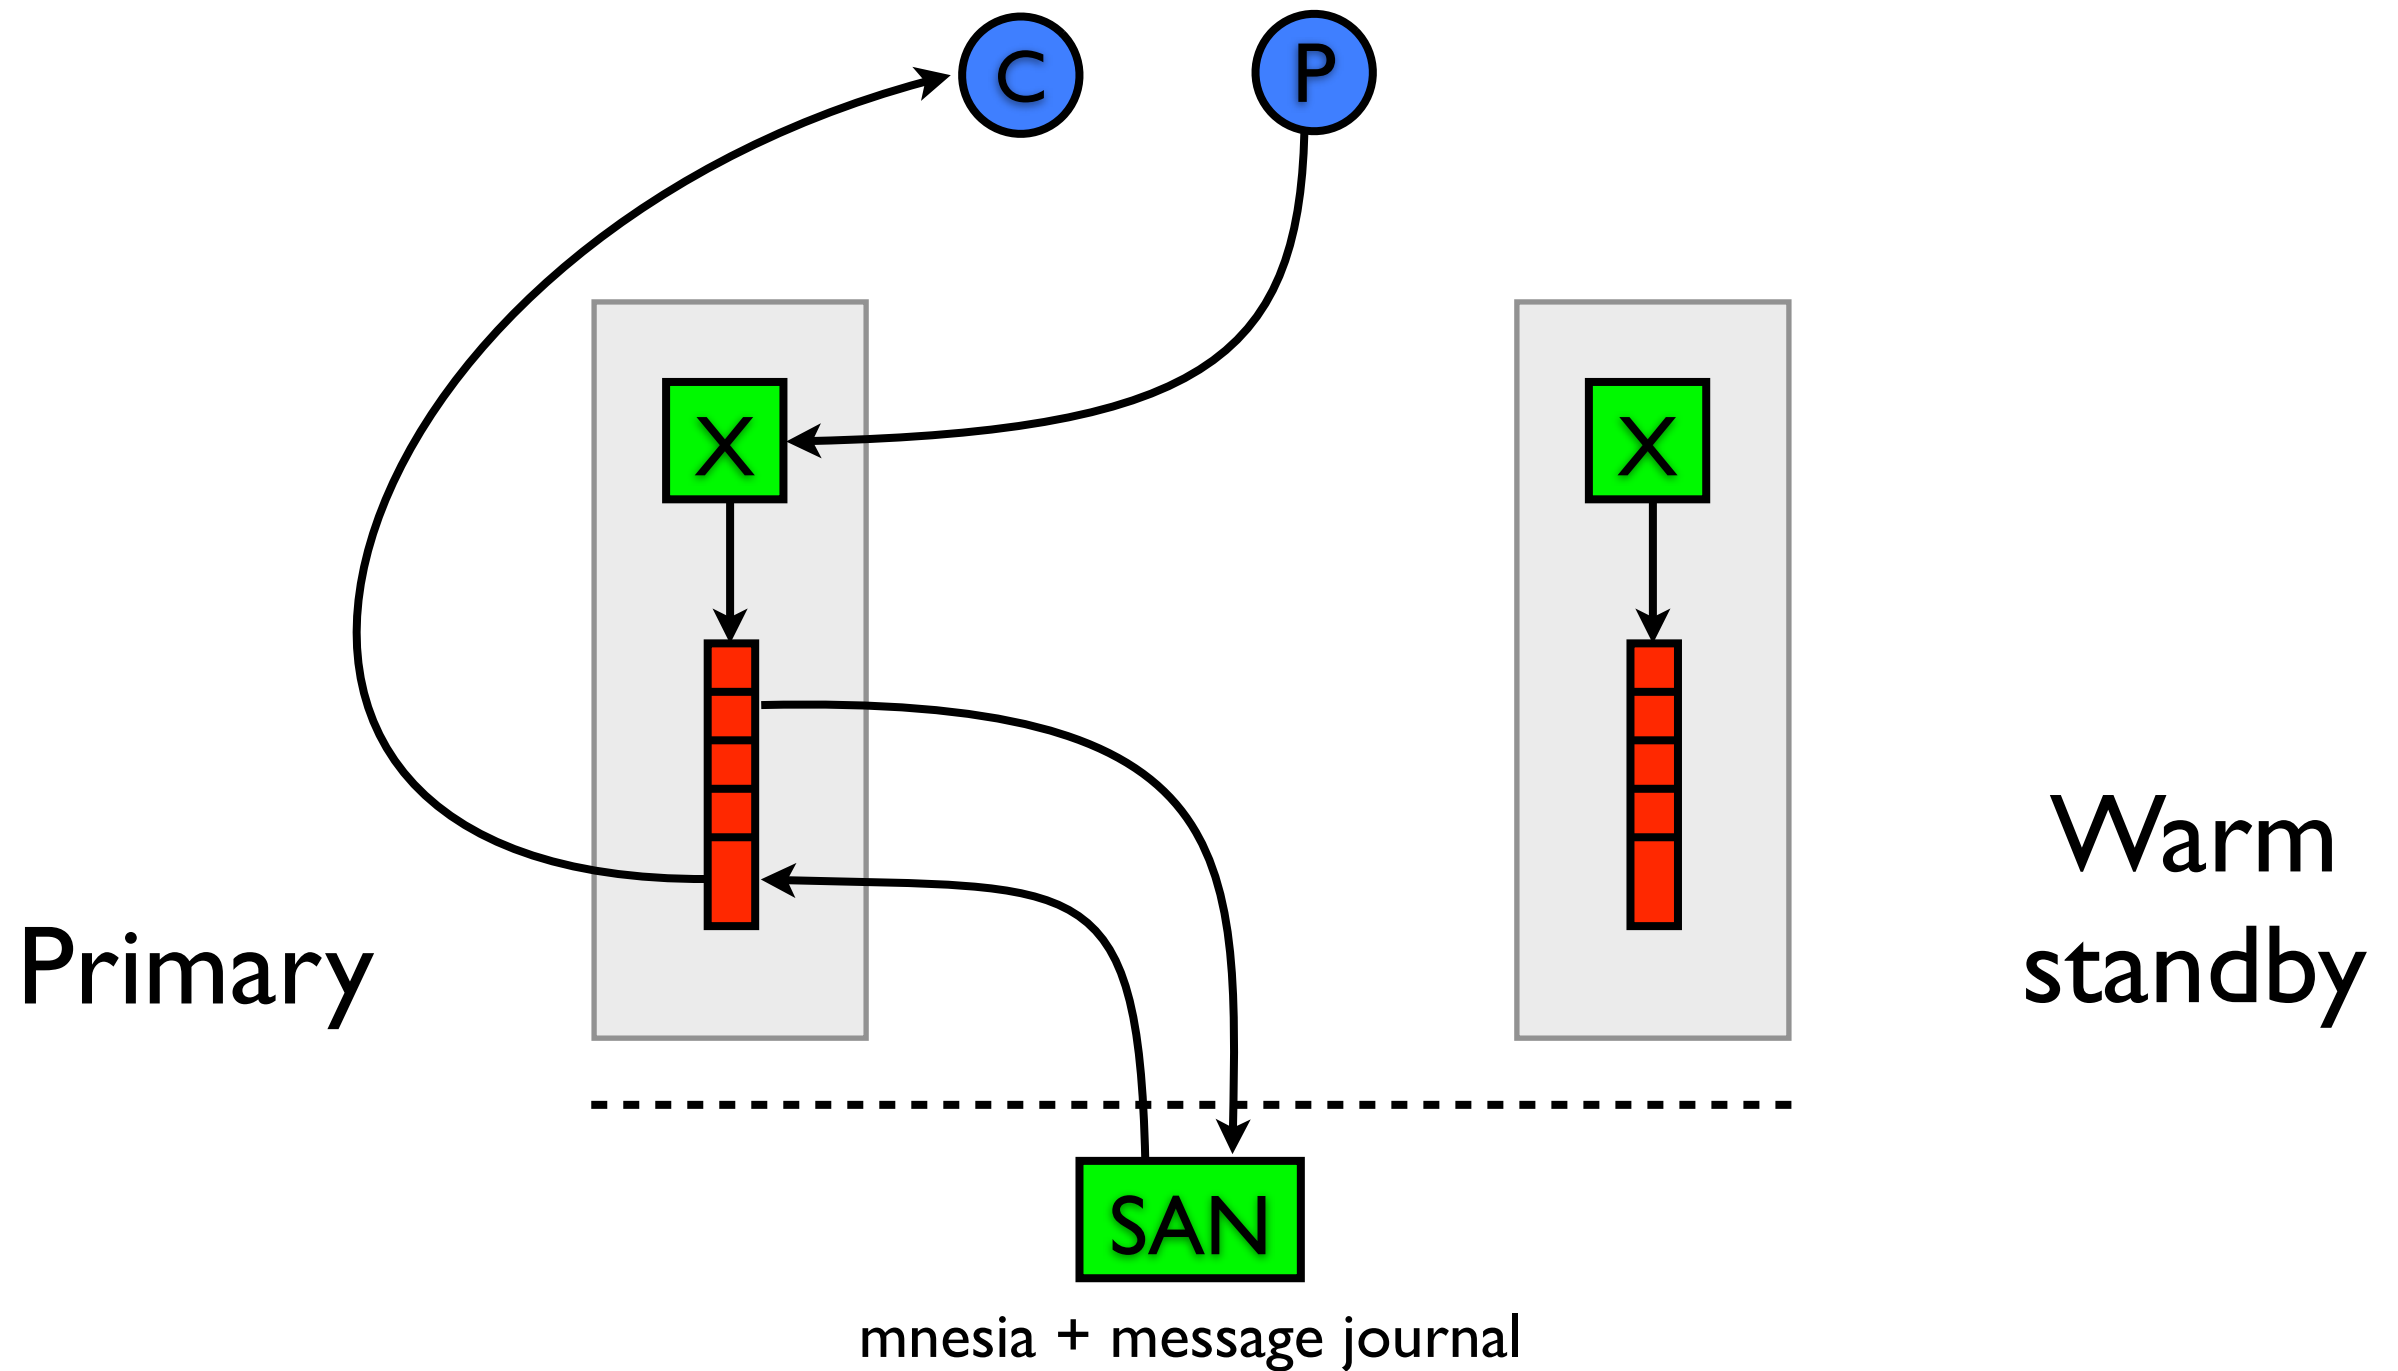
Failover

- Split “availability” into
 - Service availability: a broker’s ready when you need one
 - Data availability: your persisted messages survive failures
- Short outage during failover; non-ack’d messages will need to be retransmitted
- Need better? Use redundant data paths

Failover

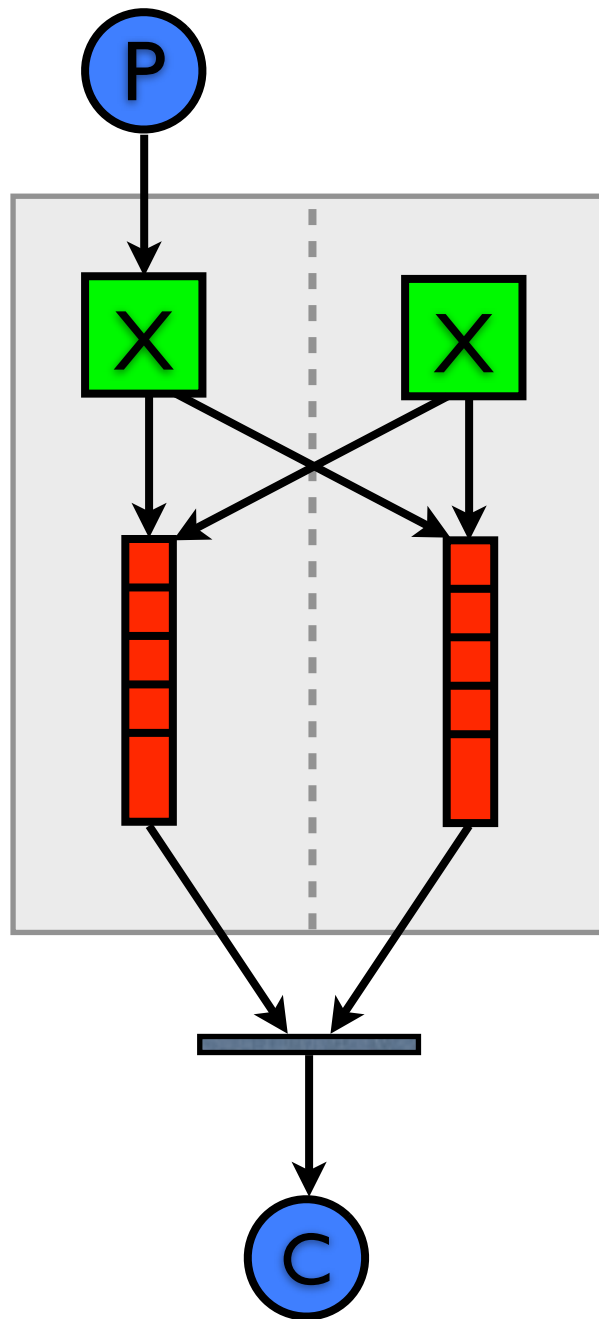
- Off-the-shelf components:
 - Networked fsync()able filesystem
 - Failure monitor: Linux-HA, ping + virtual ethernet, ...

Failover

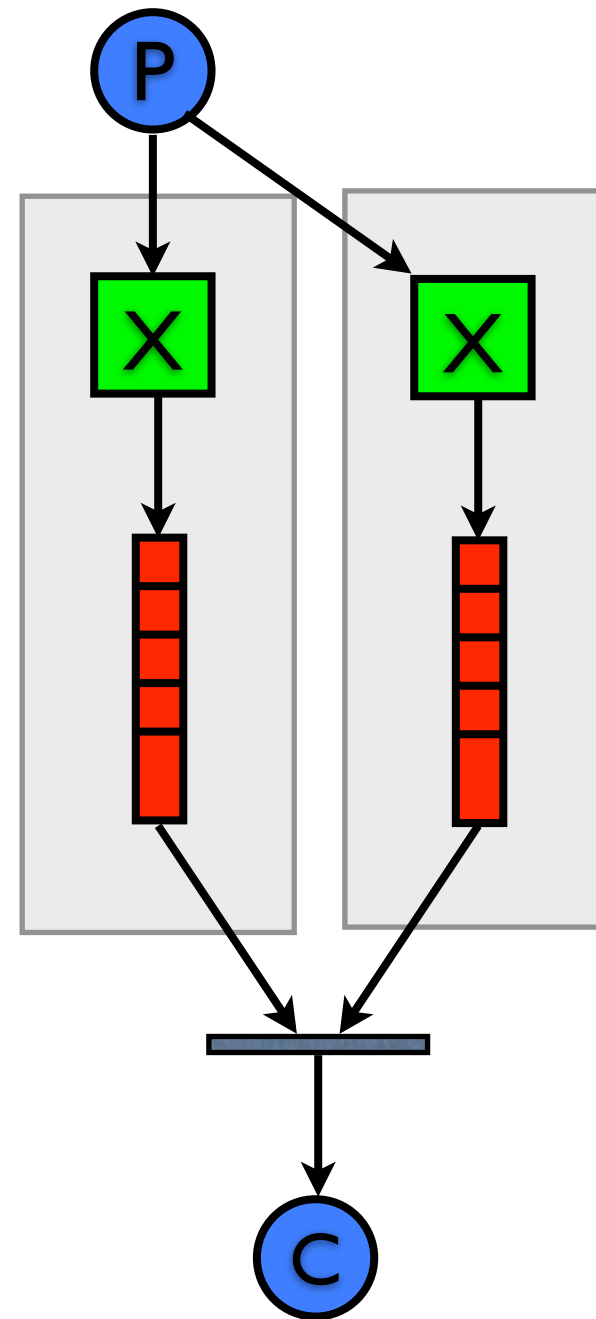


Redundancy for HA

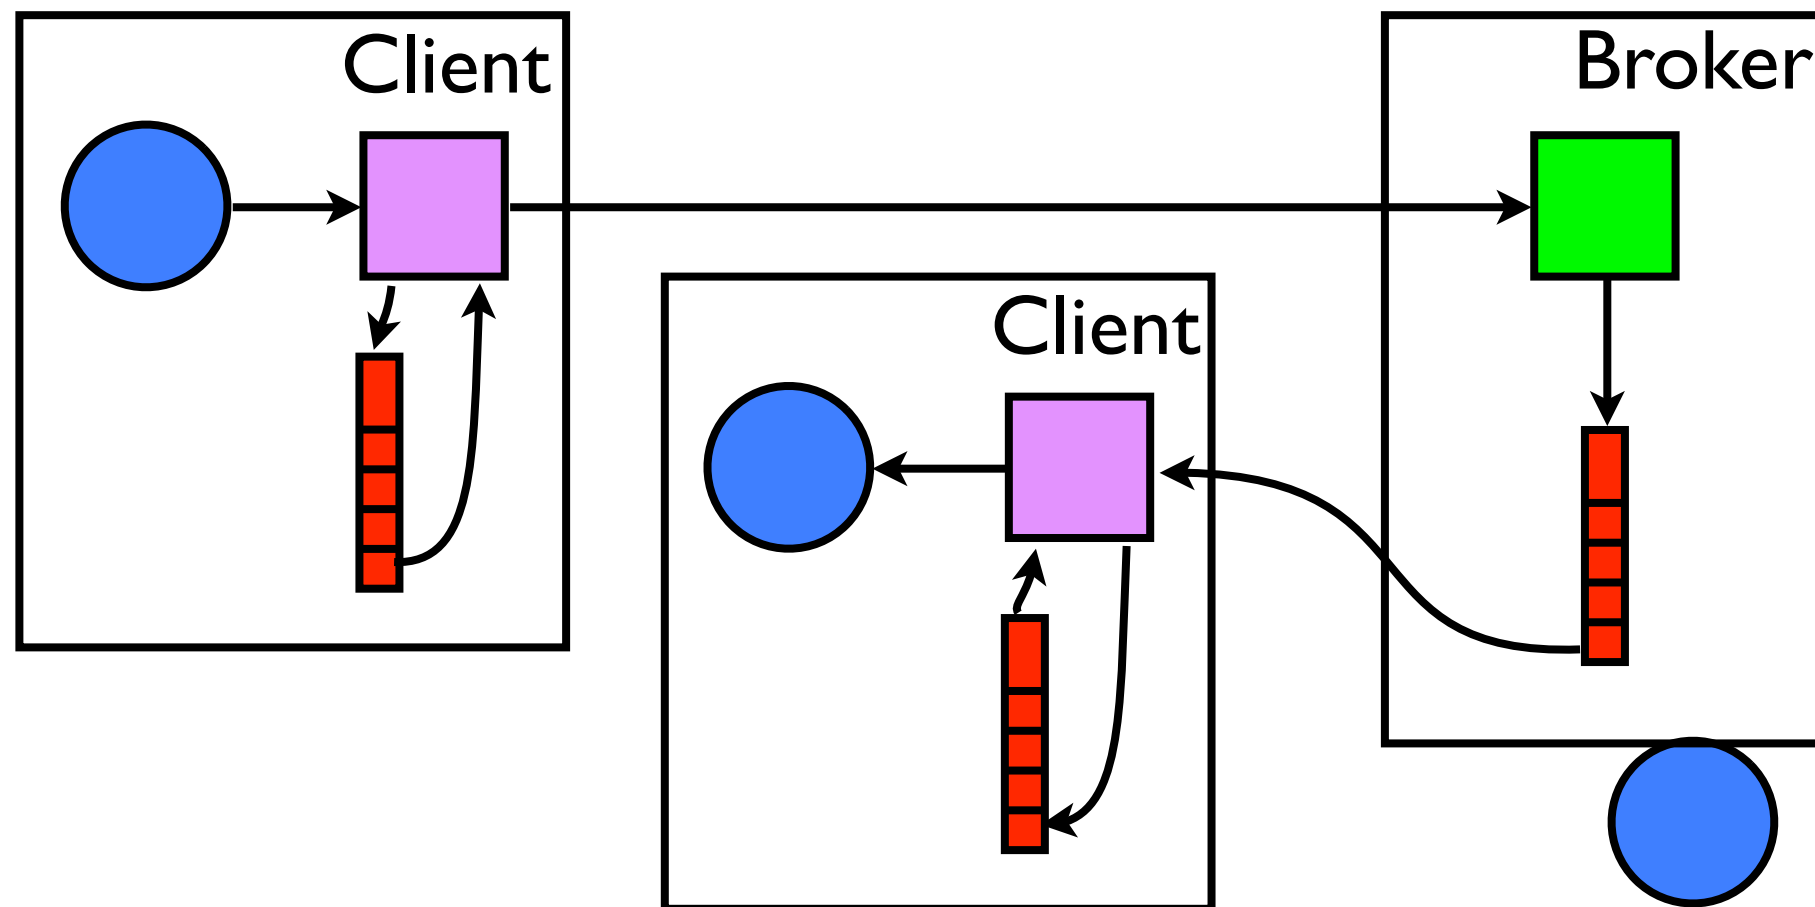
One
broker,
two nodes
each



Two
brokers,
one node
each



Exactly-Once Delivery



- “guaranteed delivery”, even with intermittent links

- “auto-deduplication”

Paul Baran's Networks

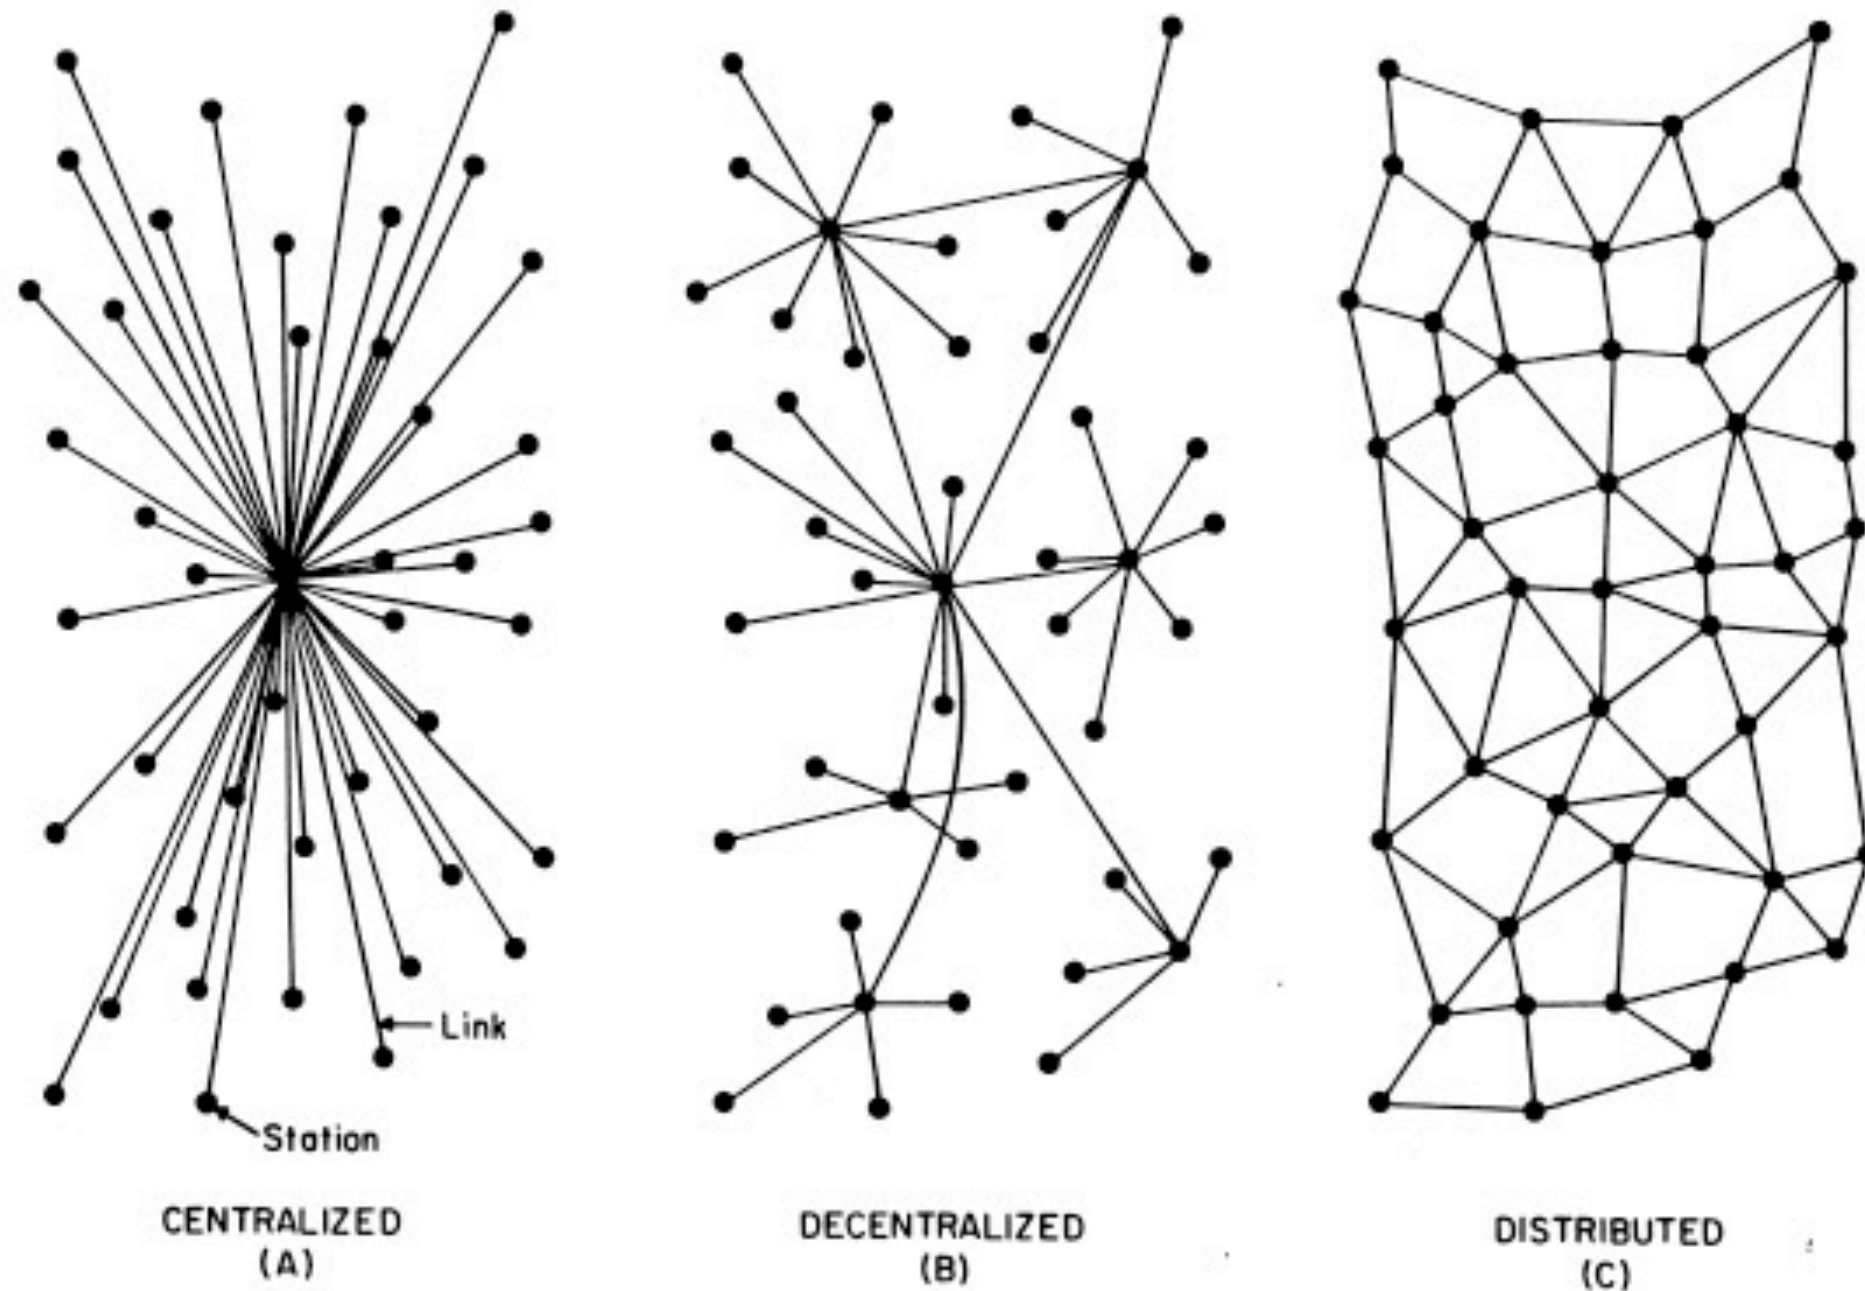
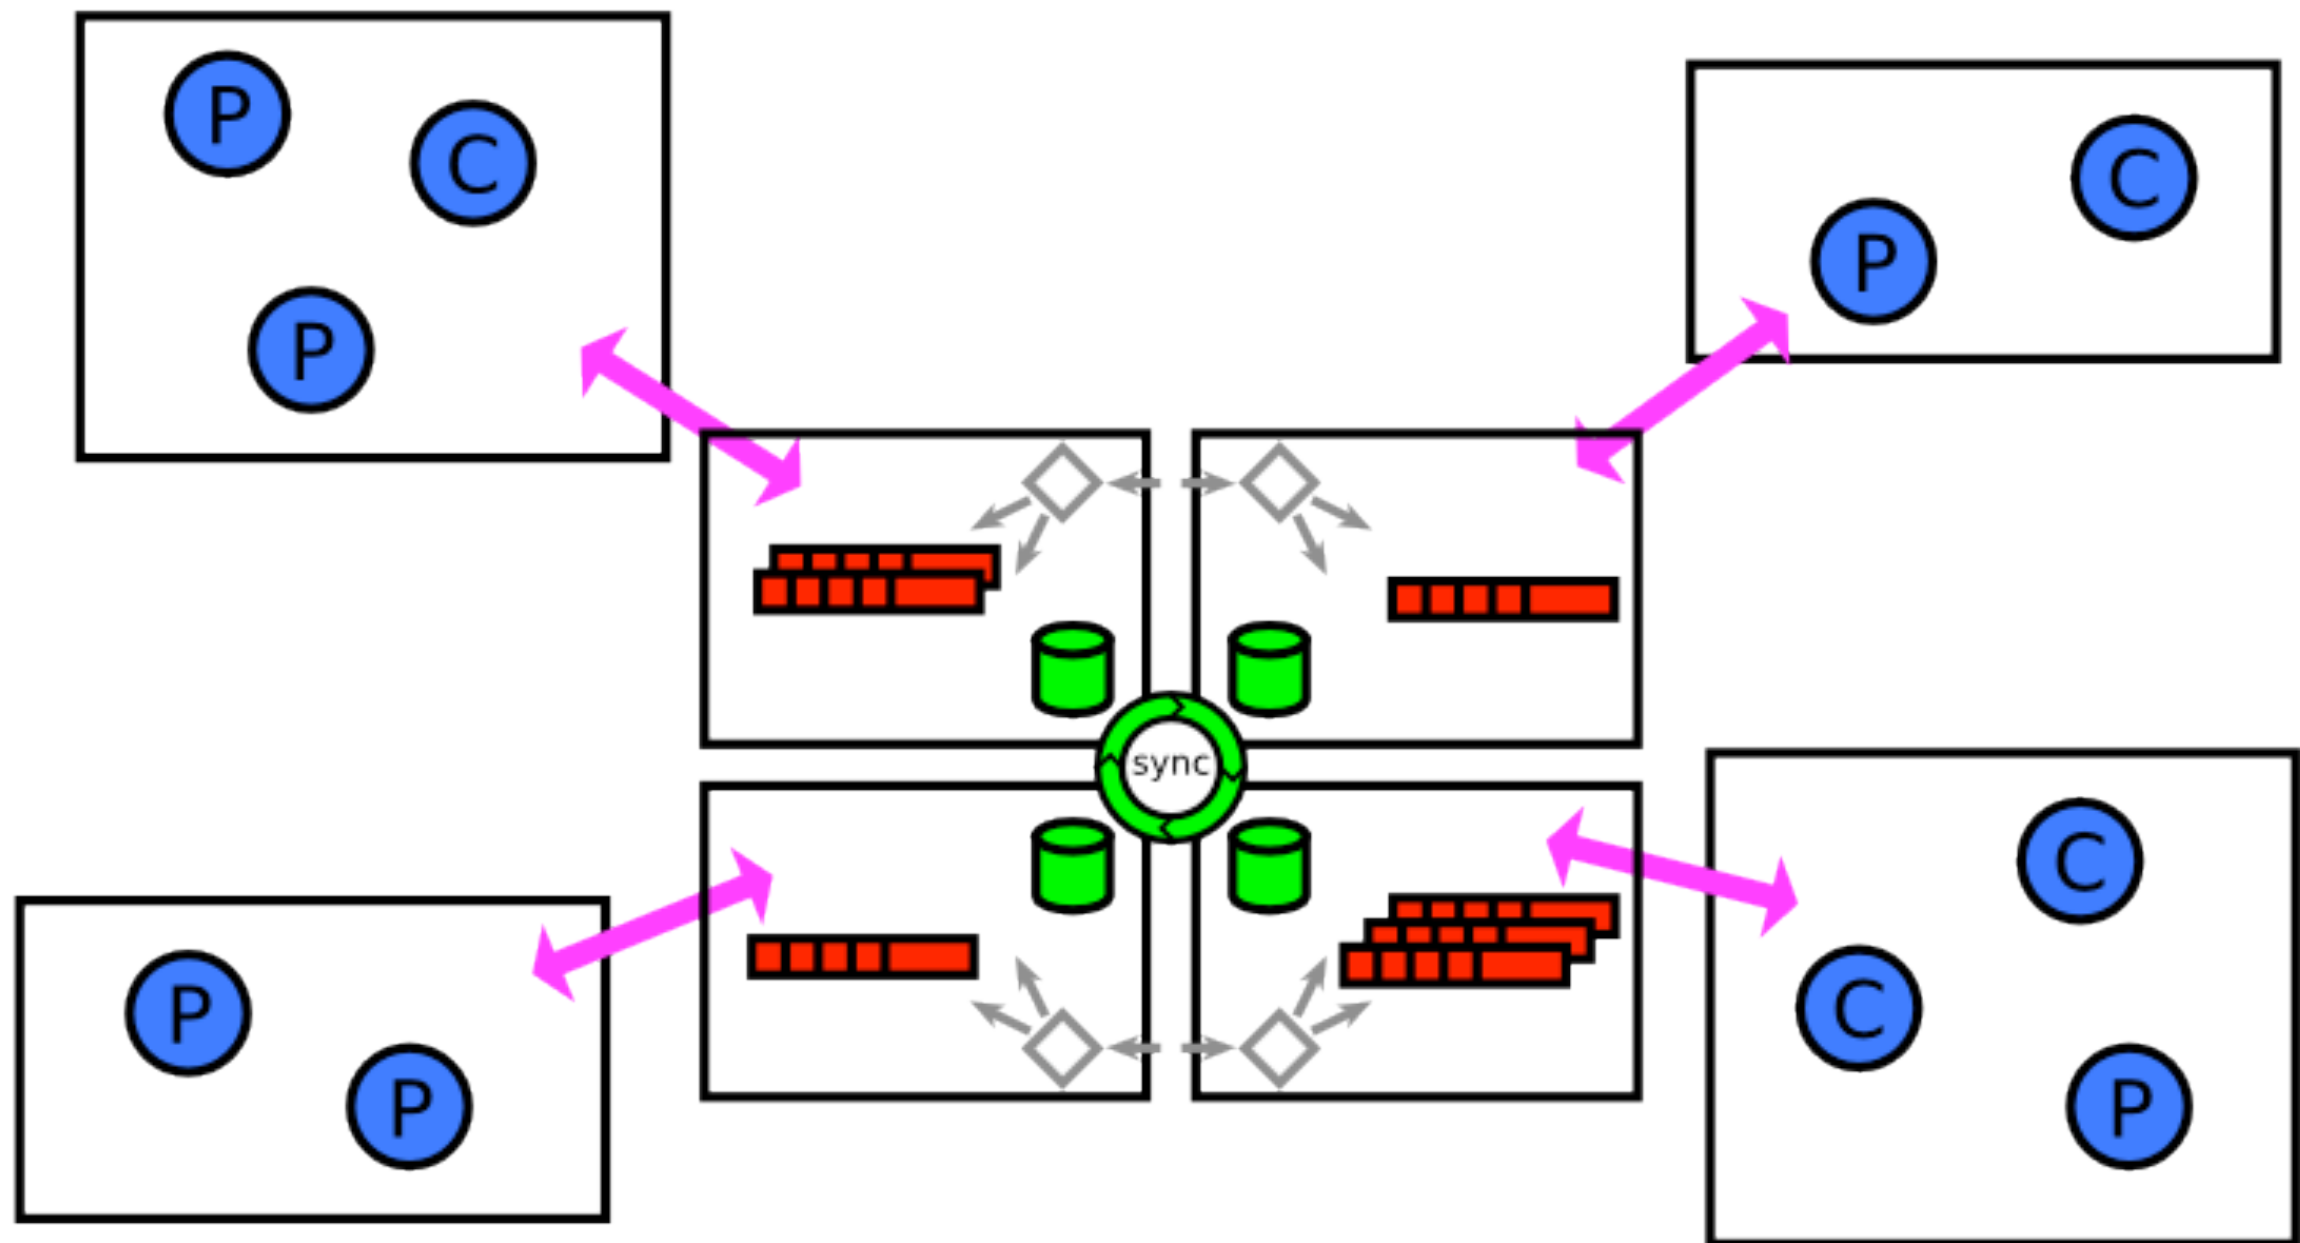


FIG. 1 — Centralized, Decentralized and Distributed Networks

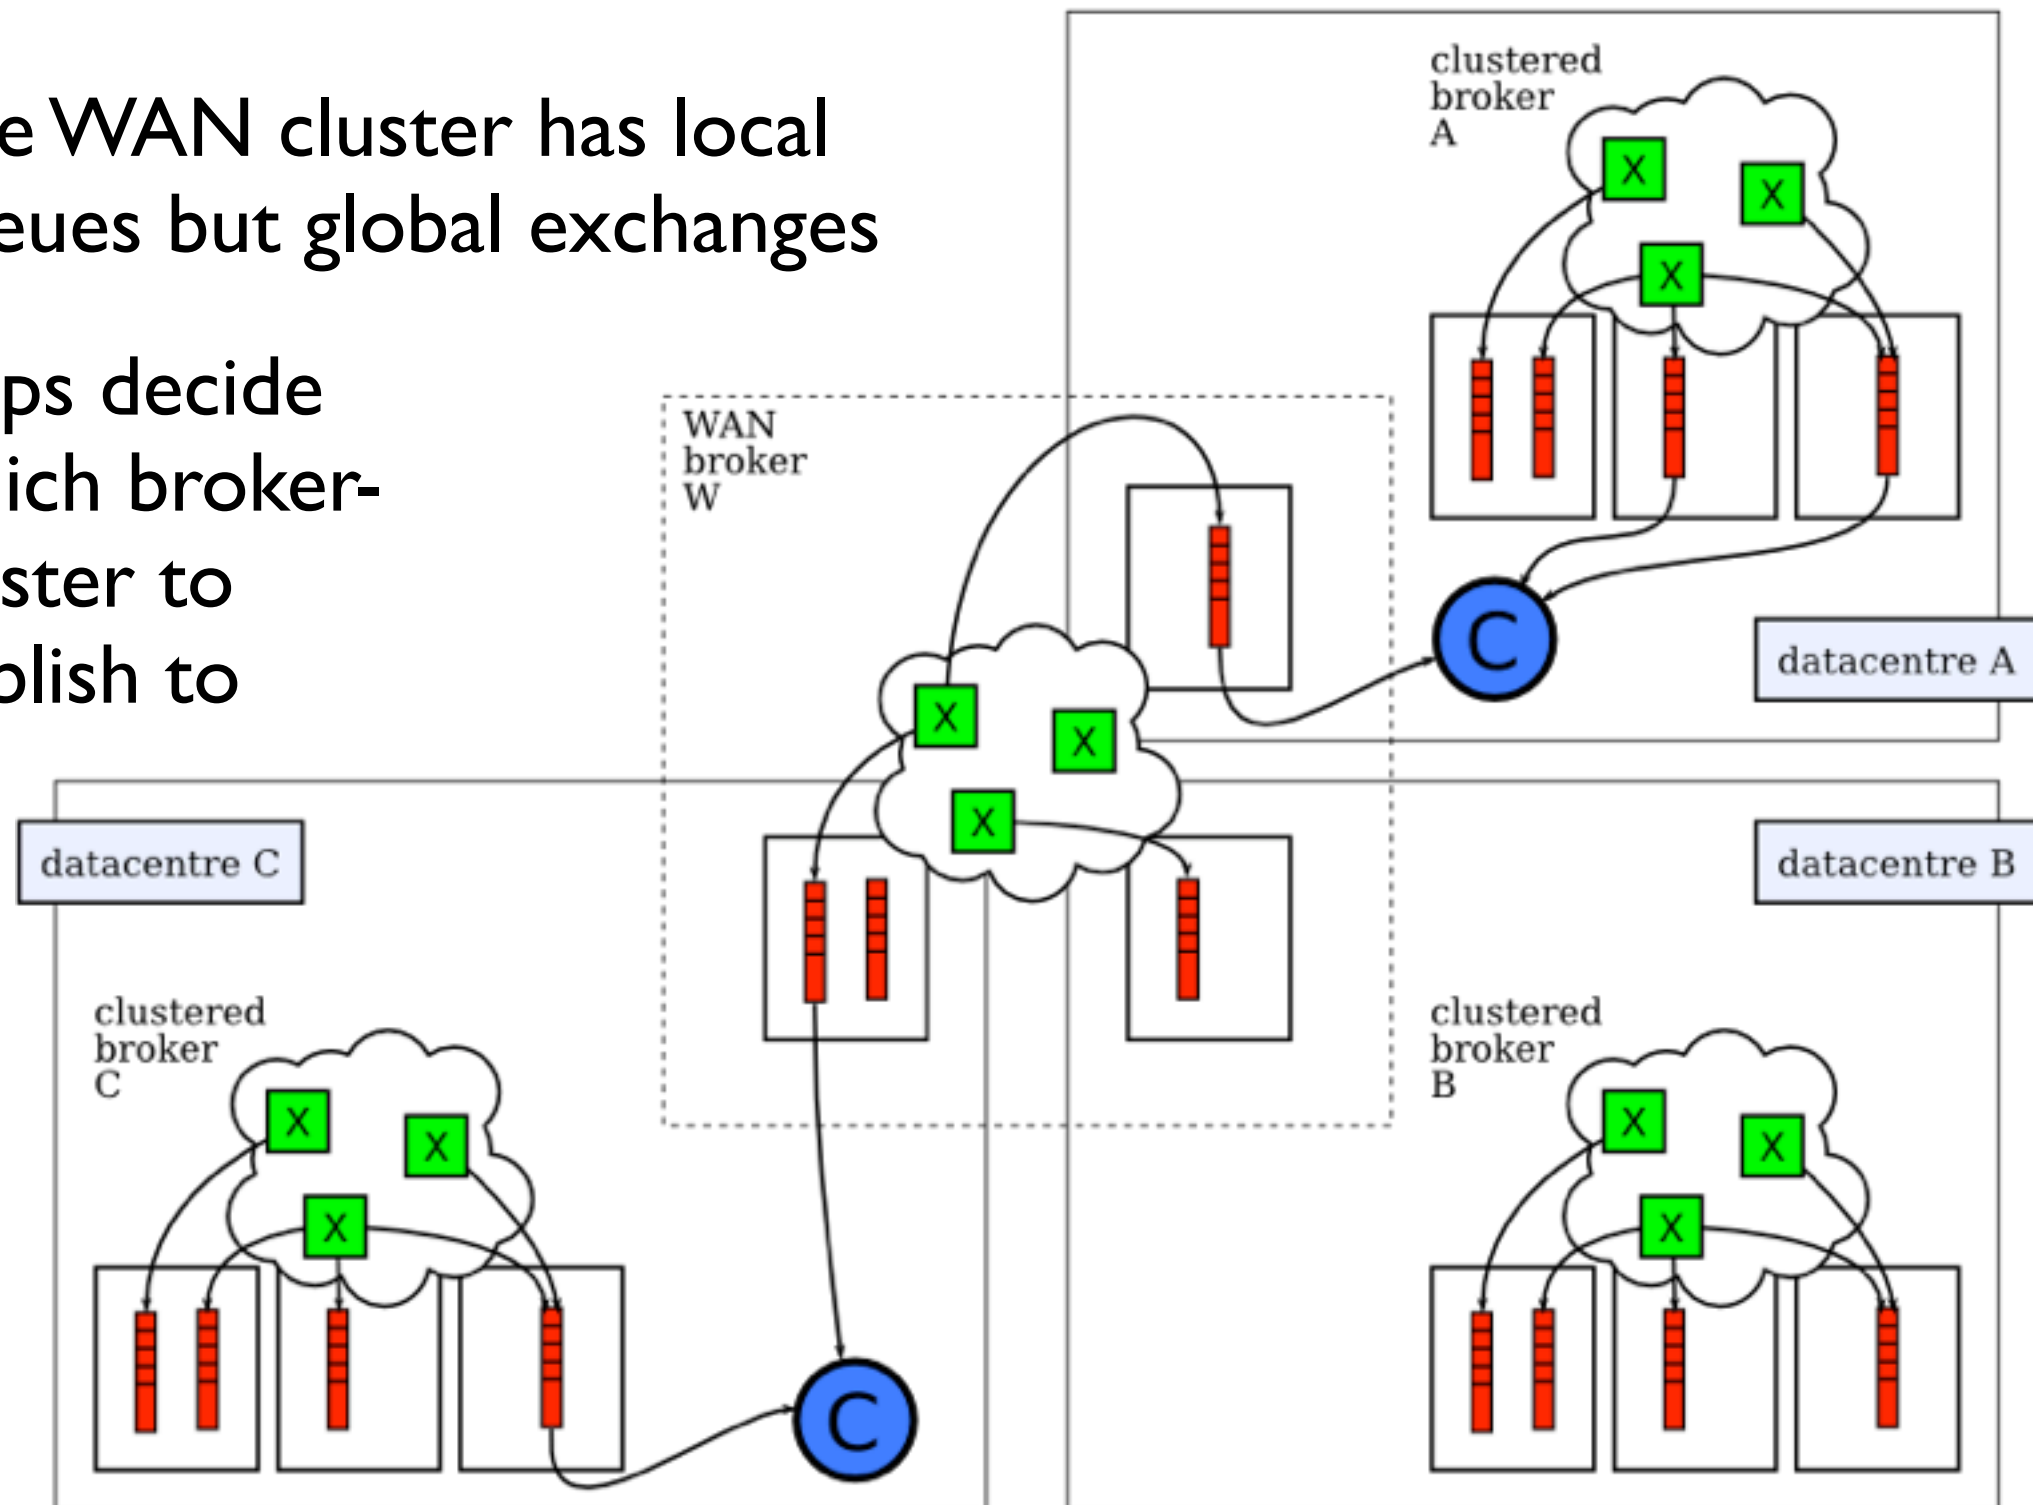
Clustering



WAN Cluster

The WAN cluster has local queues but global exchanges

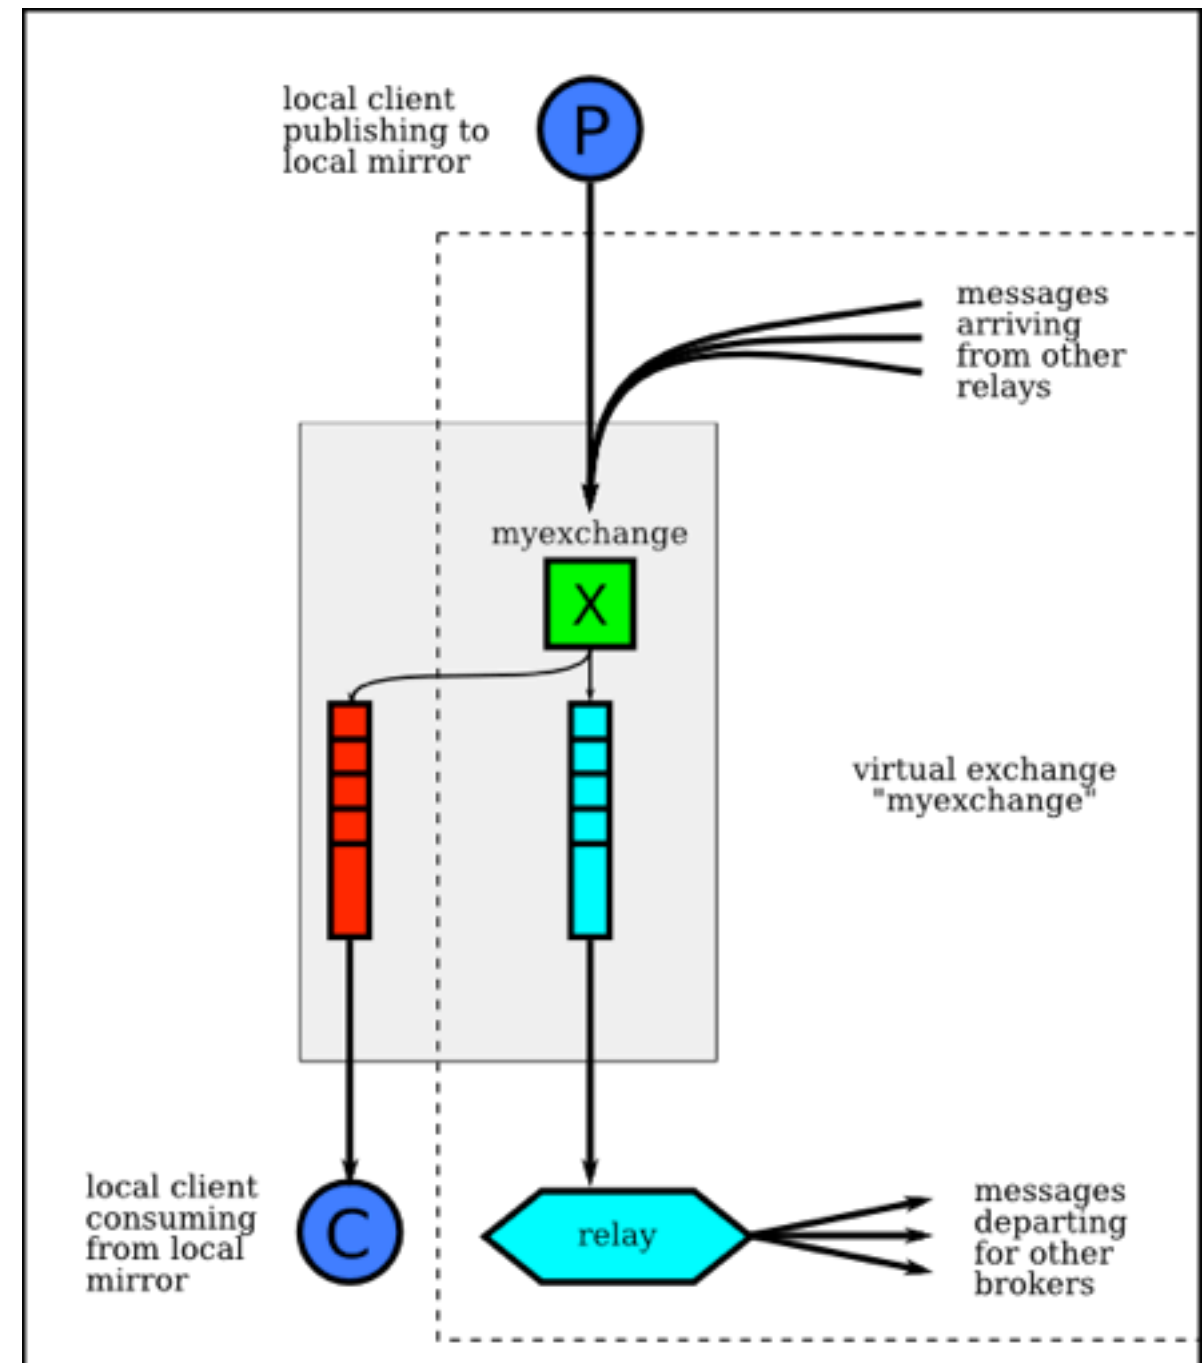
Apps decide which broker-cluster to publish to



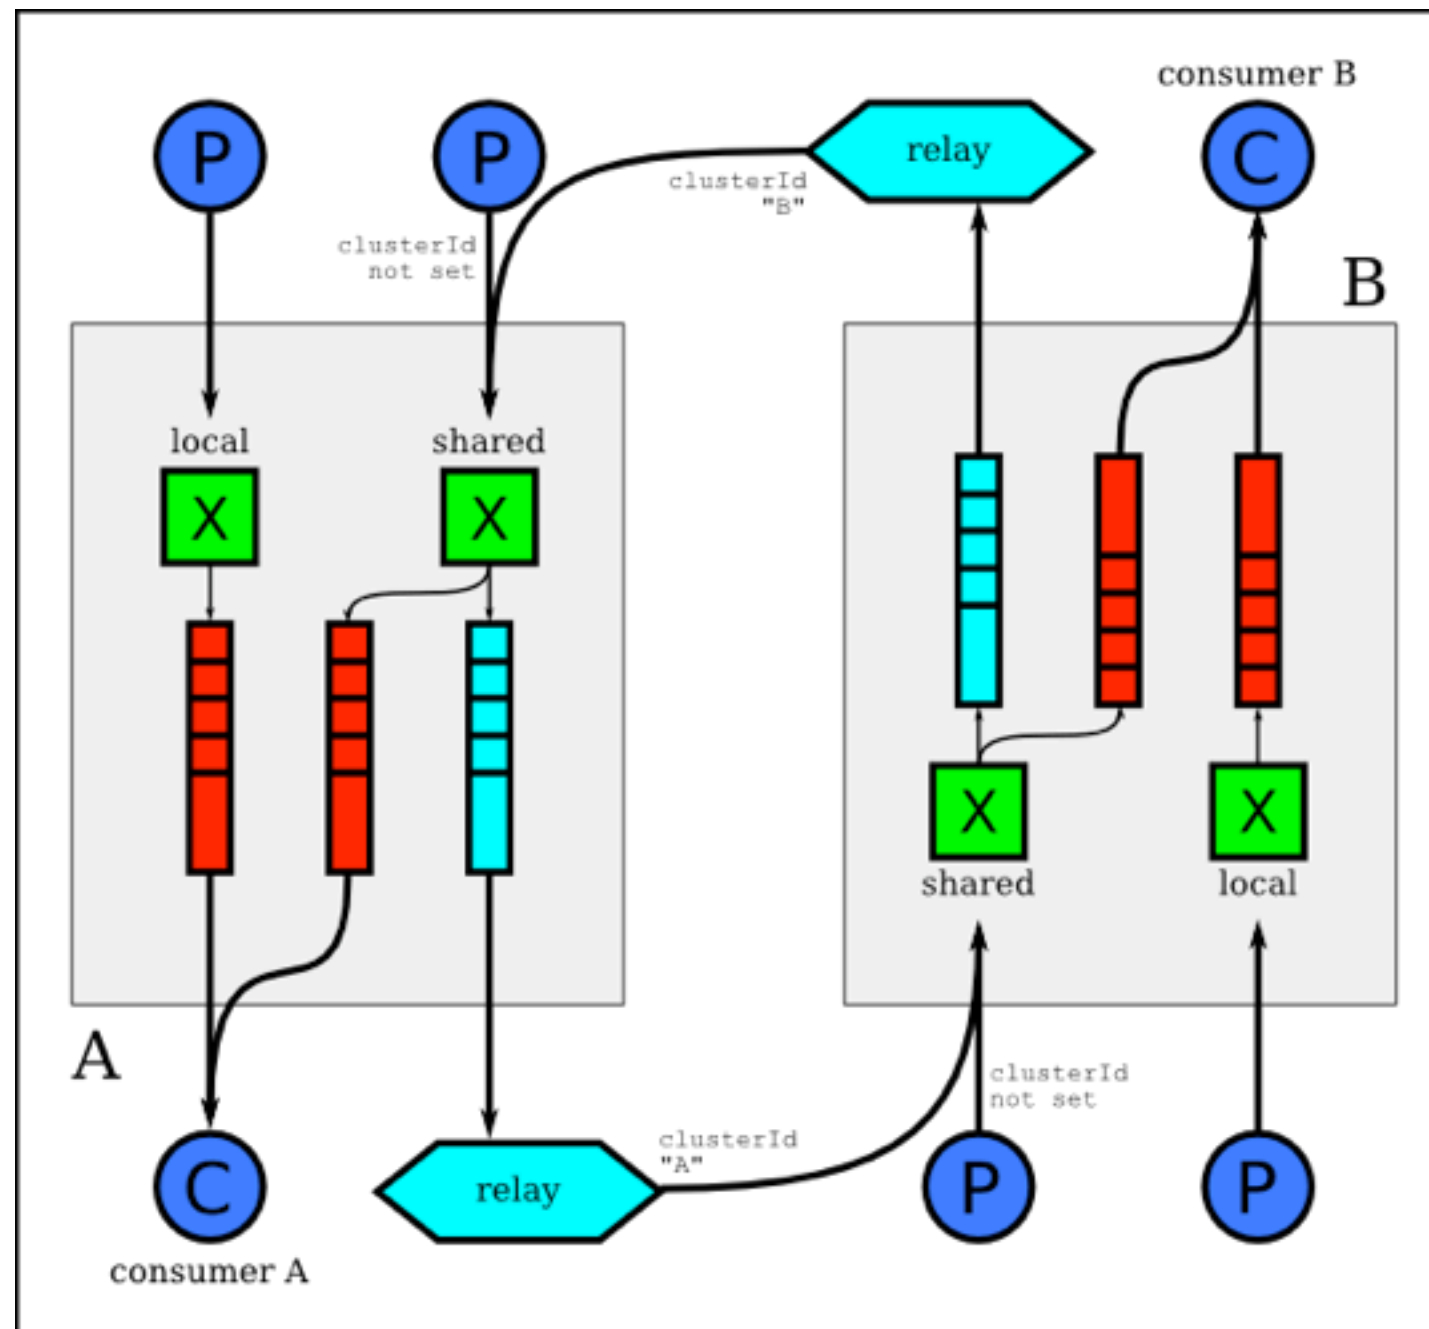
Synchronisation

When clustering might not be right:

- huge networks
- intermittent connectivity
- ruling bandwidth with an iron fist
- different administrative domains

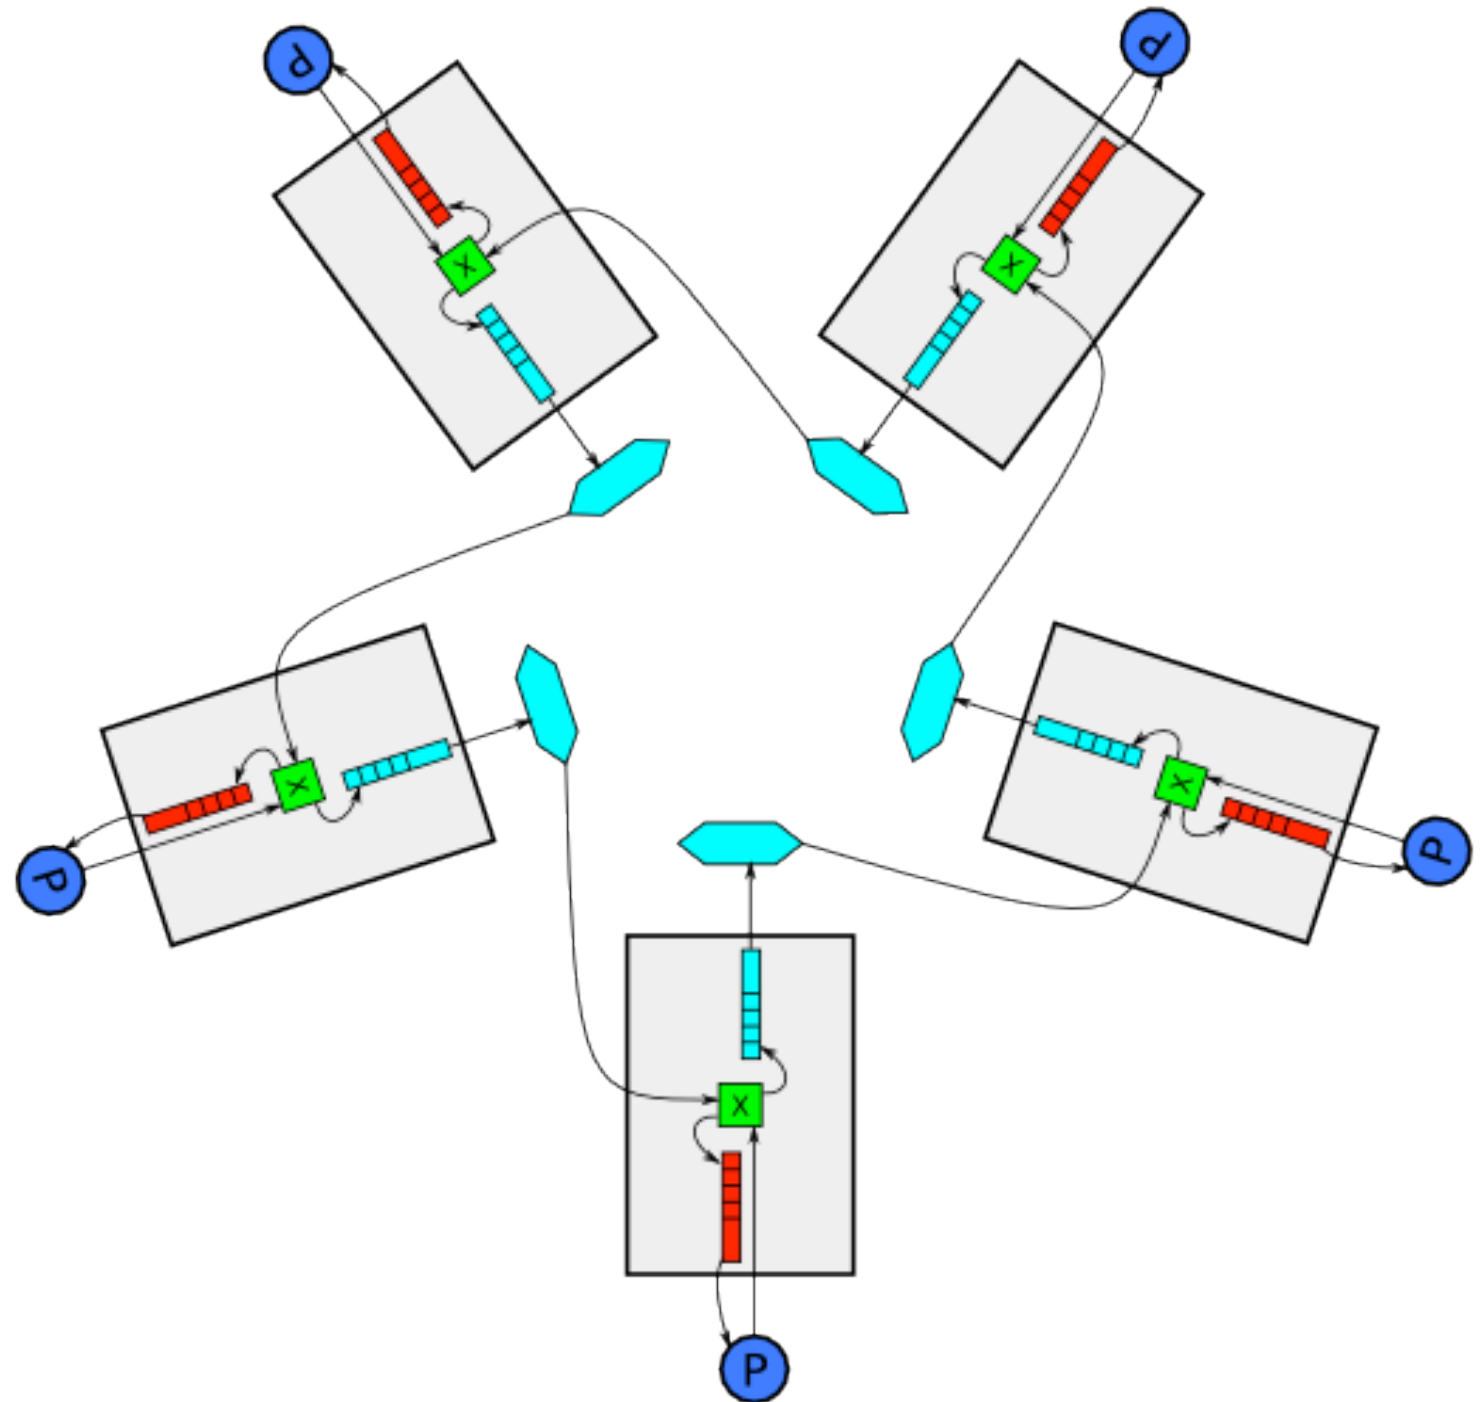


Synchronisation



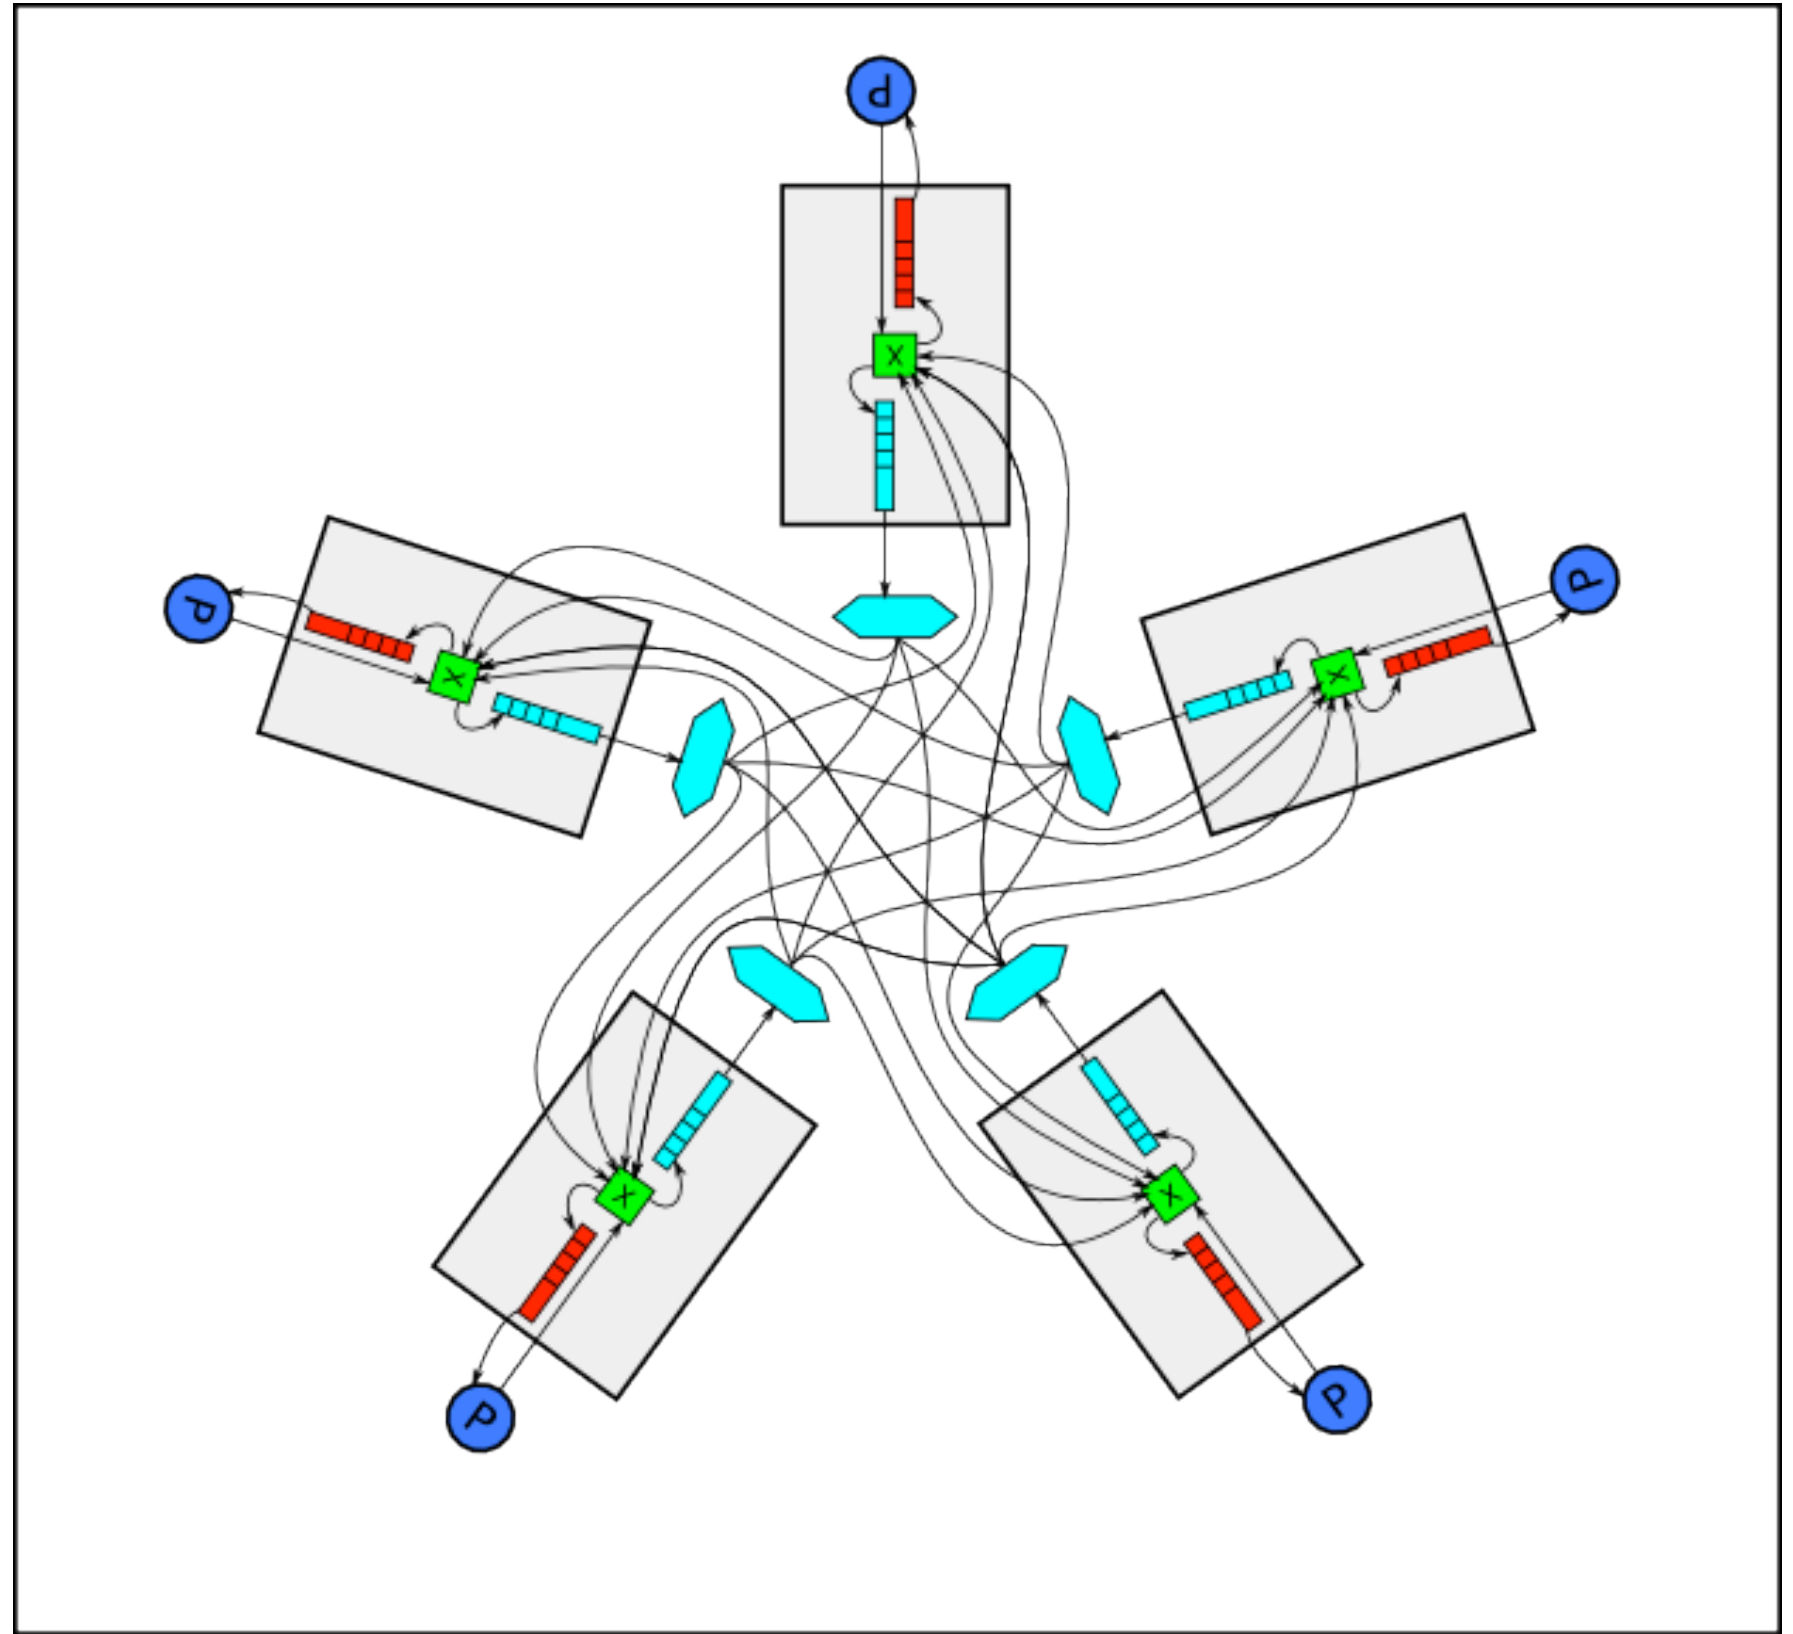
Ring

Pass it on to your neighbour if your neighbour's name isn't in the list yet



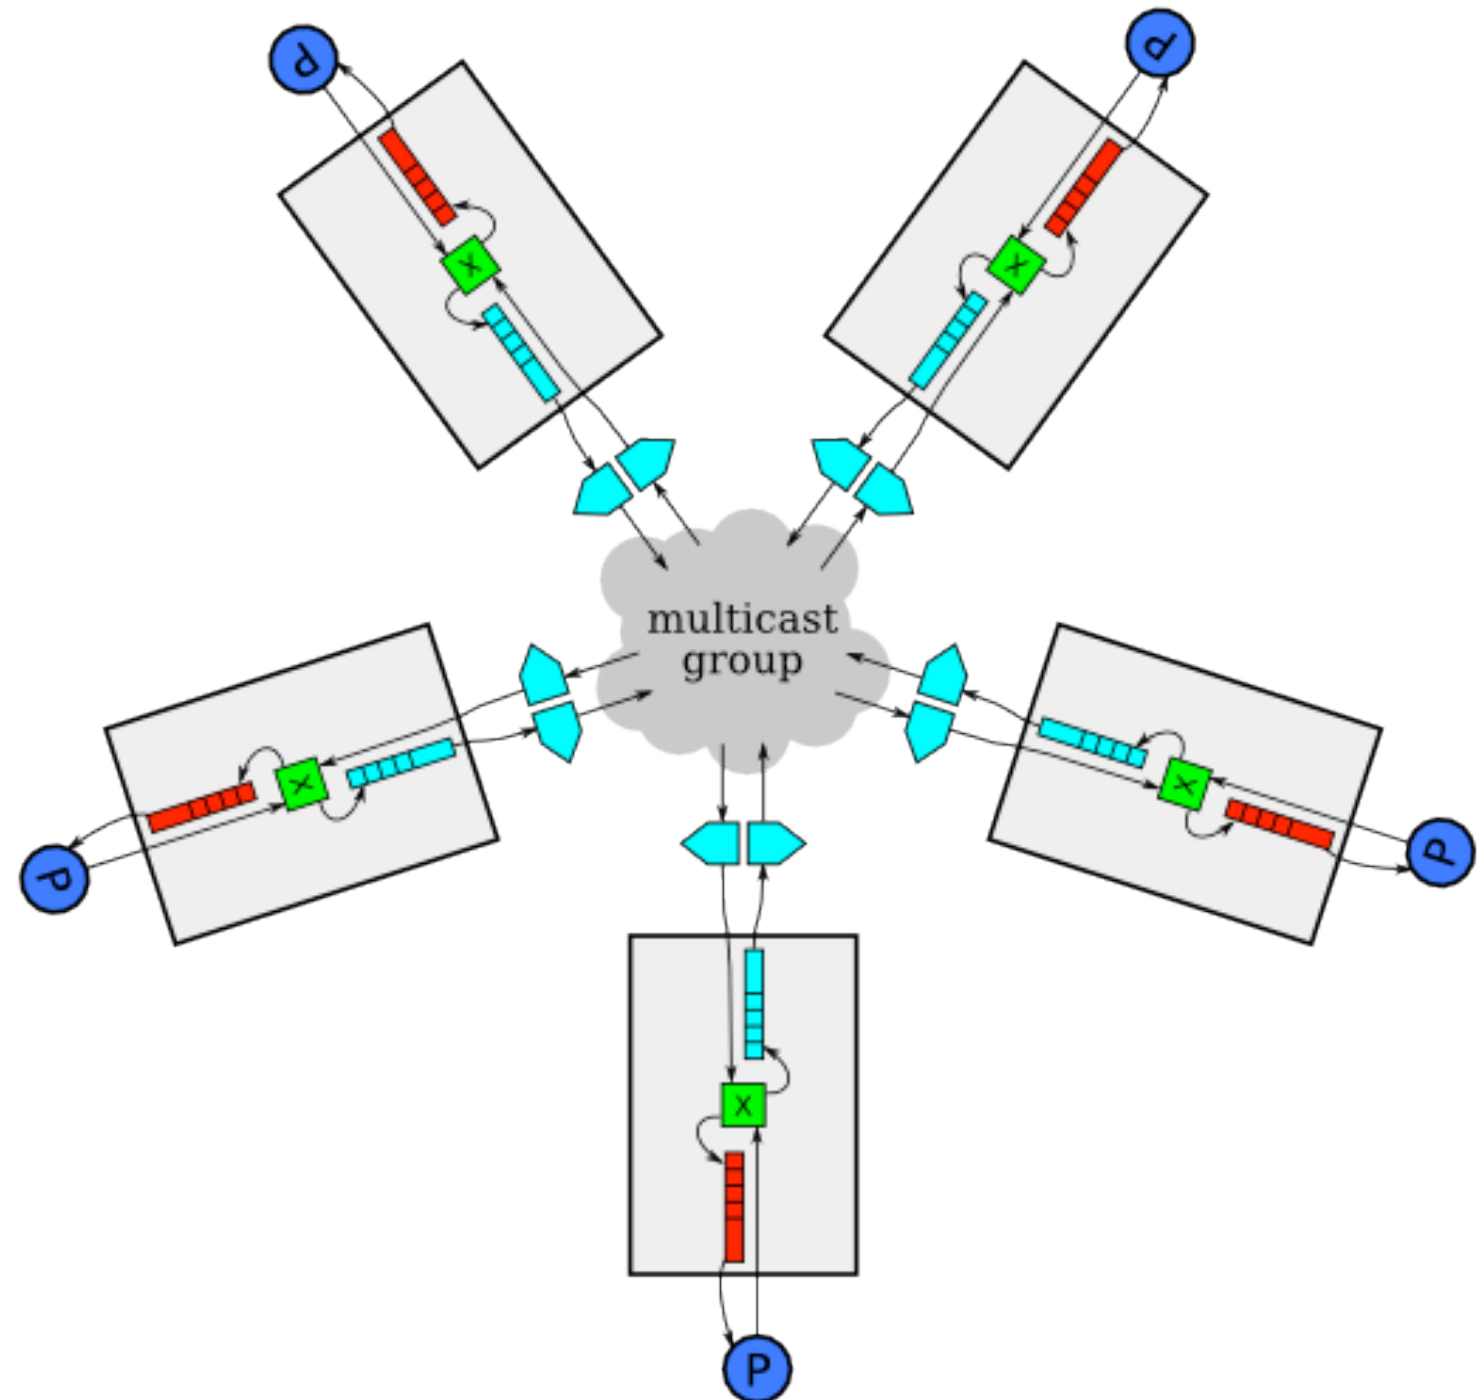
Complete Graph

Pass it on to your neighbour if it hasn't been labelled at all yet



Multicast

Pass it on to your neighbour if it hasn't been labelled at all yet



Durability, Persistence & Memory Usage

Terminology

- Durable: resource (exchange, queue, binding) that survives broker restart
- Persistent: message that survives broker restart
- Durability **and** persistence required for robust storage at the broker

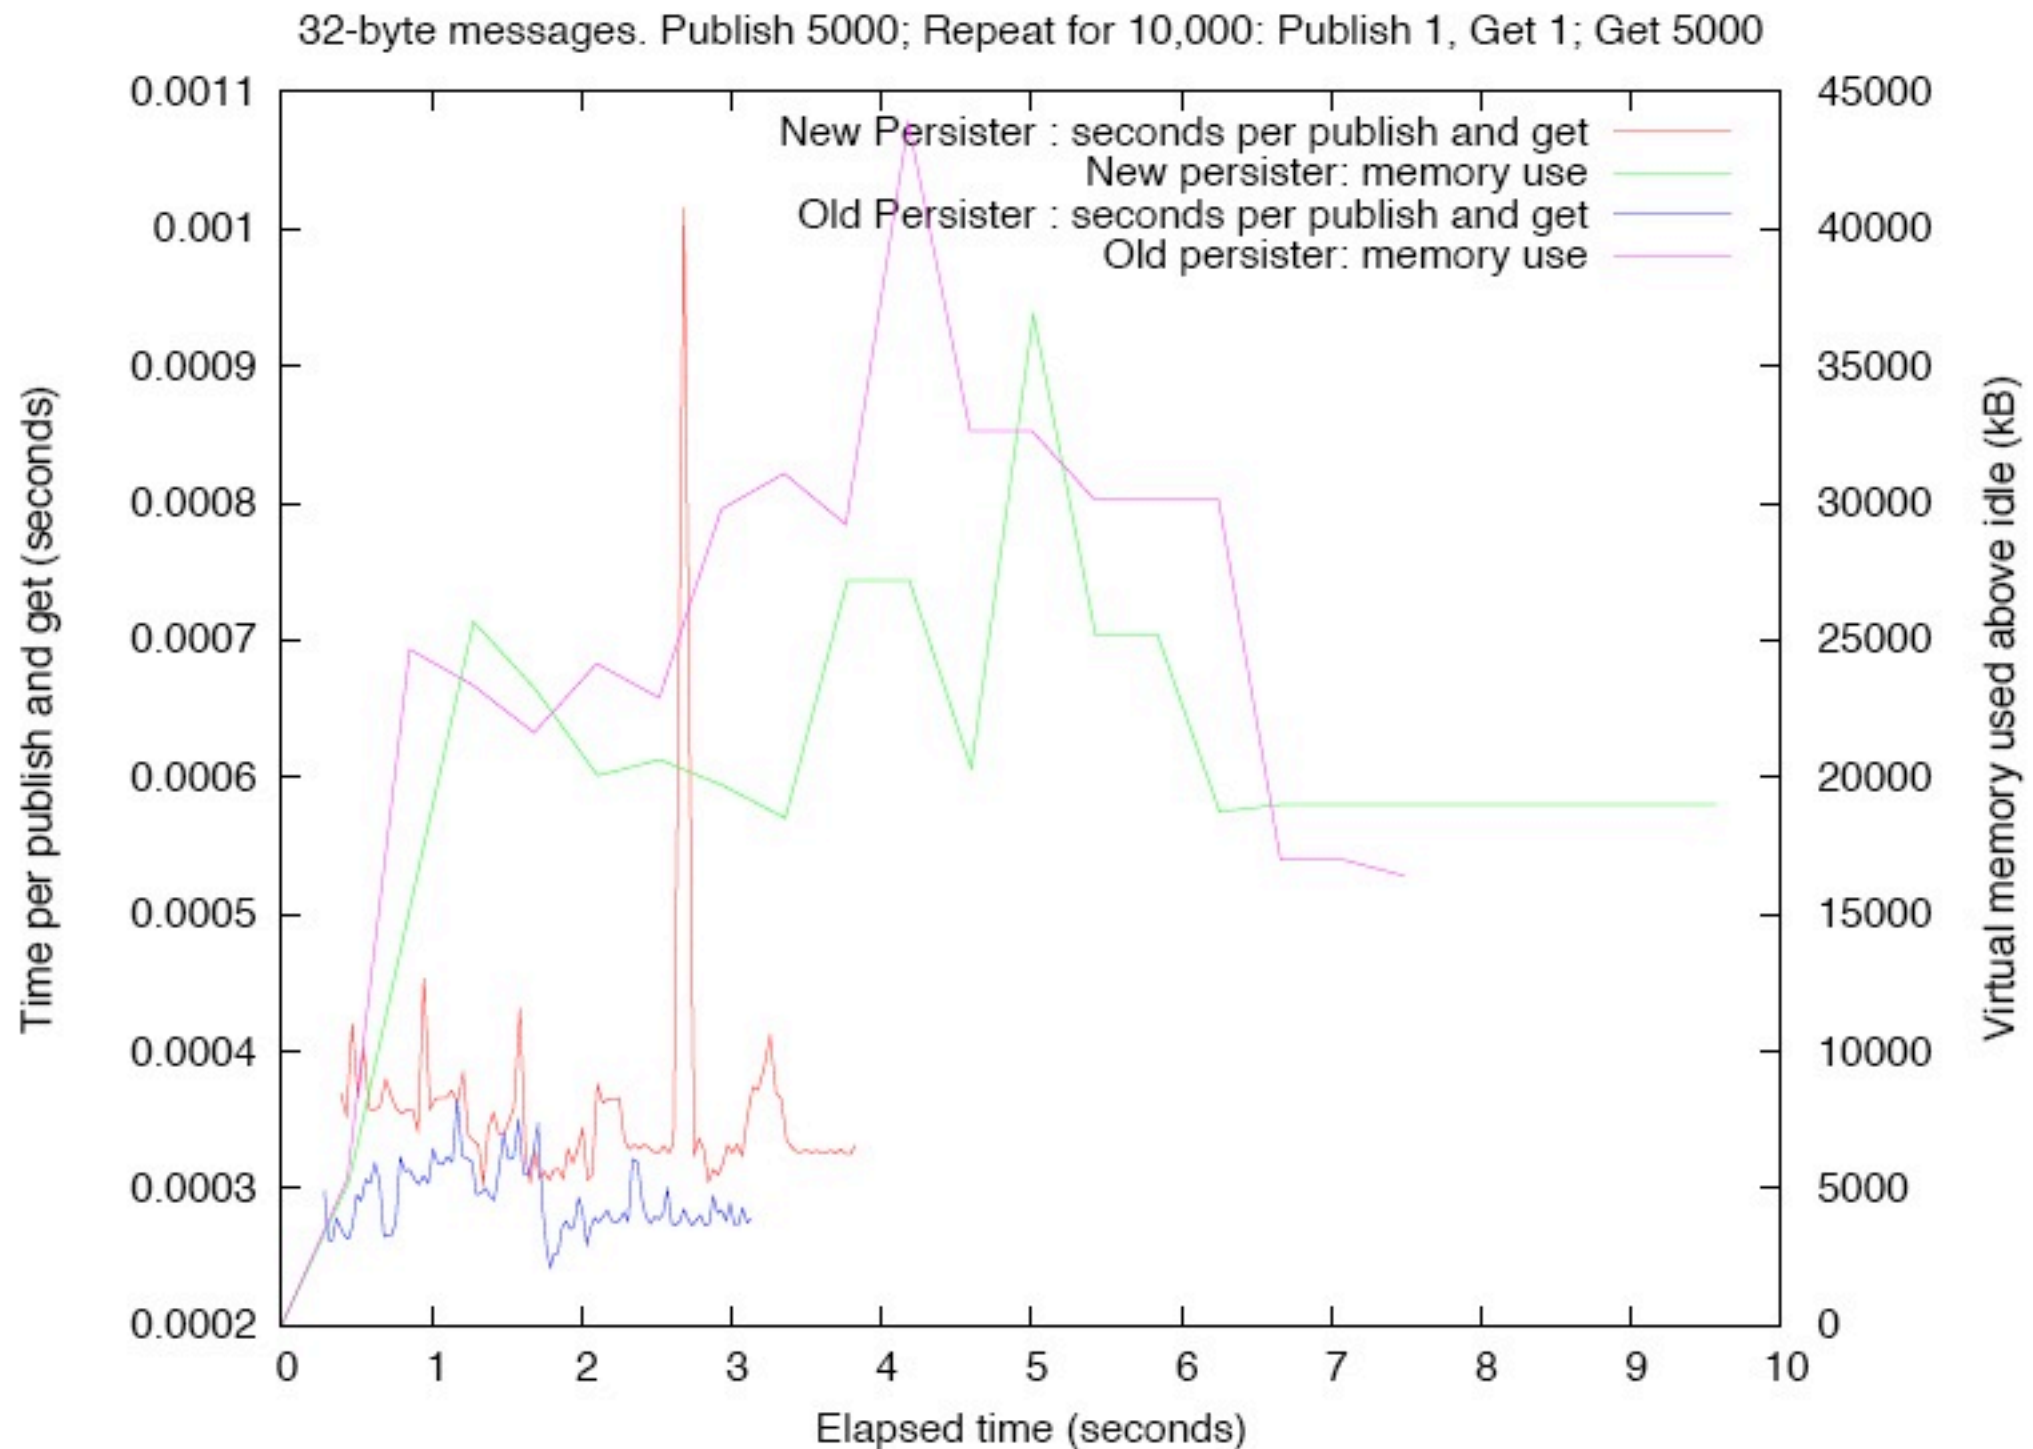
Persister: Old

- Old persister slows down dramatically once a backlog builds up
- Old persister uses disk as a backup only: still limited by RAM + swap
- Once swap runs out, it's all over

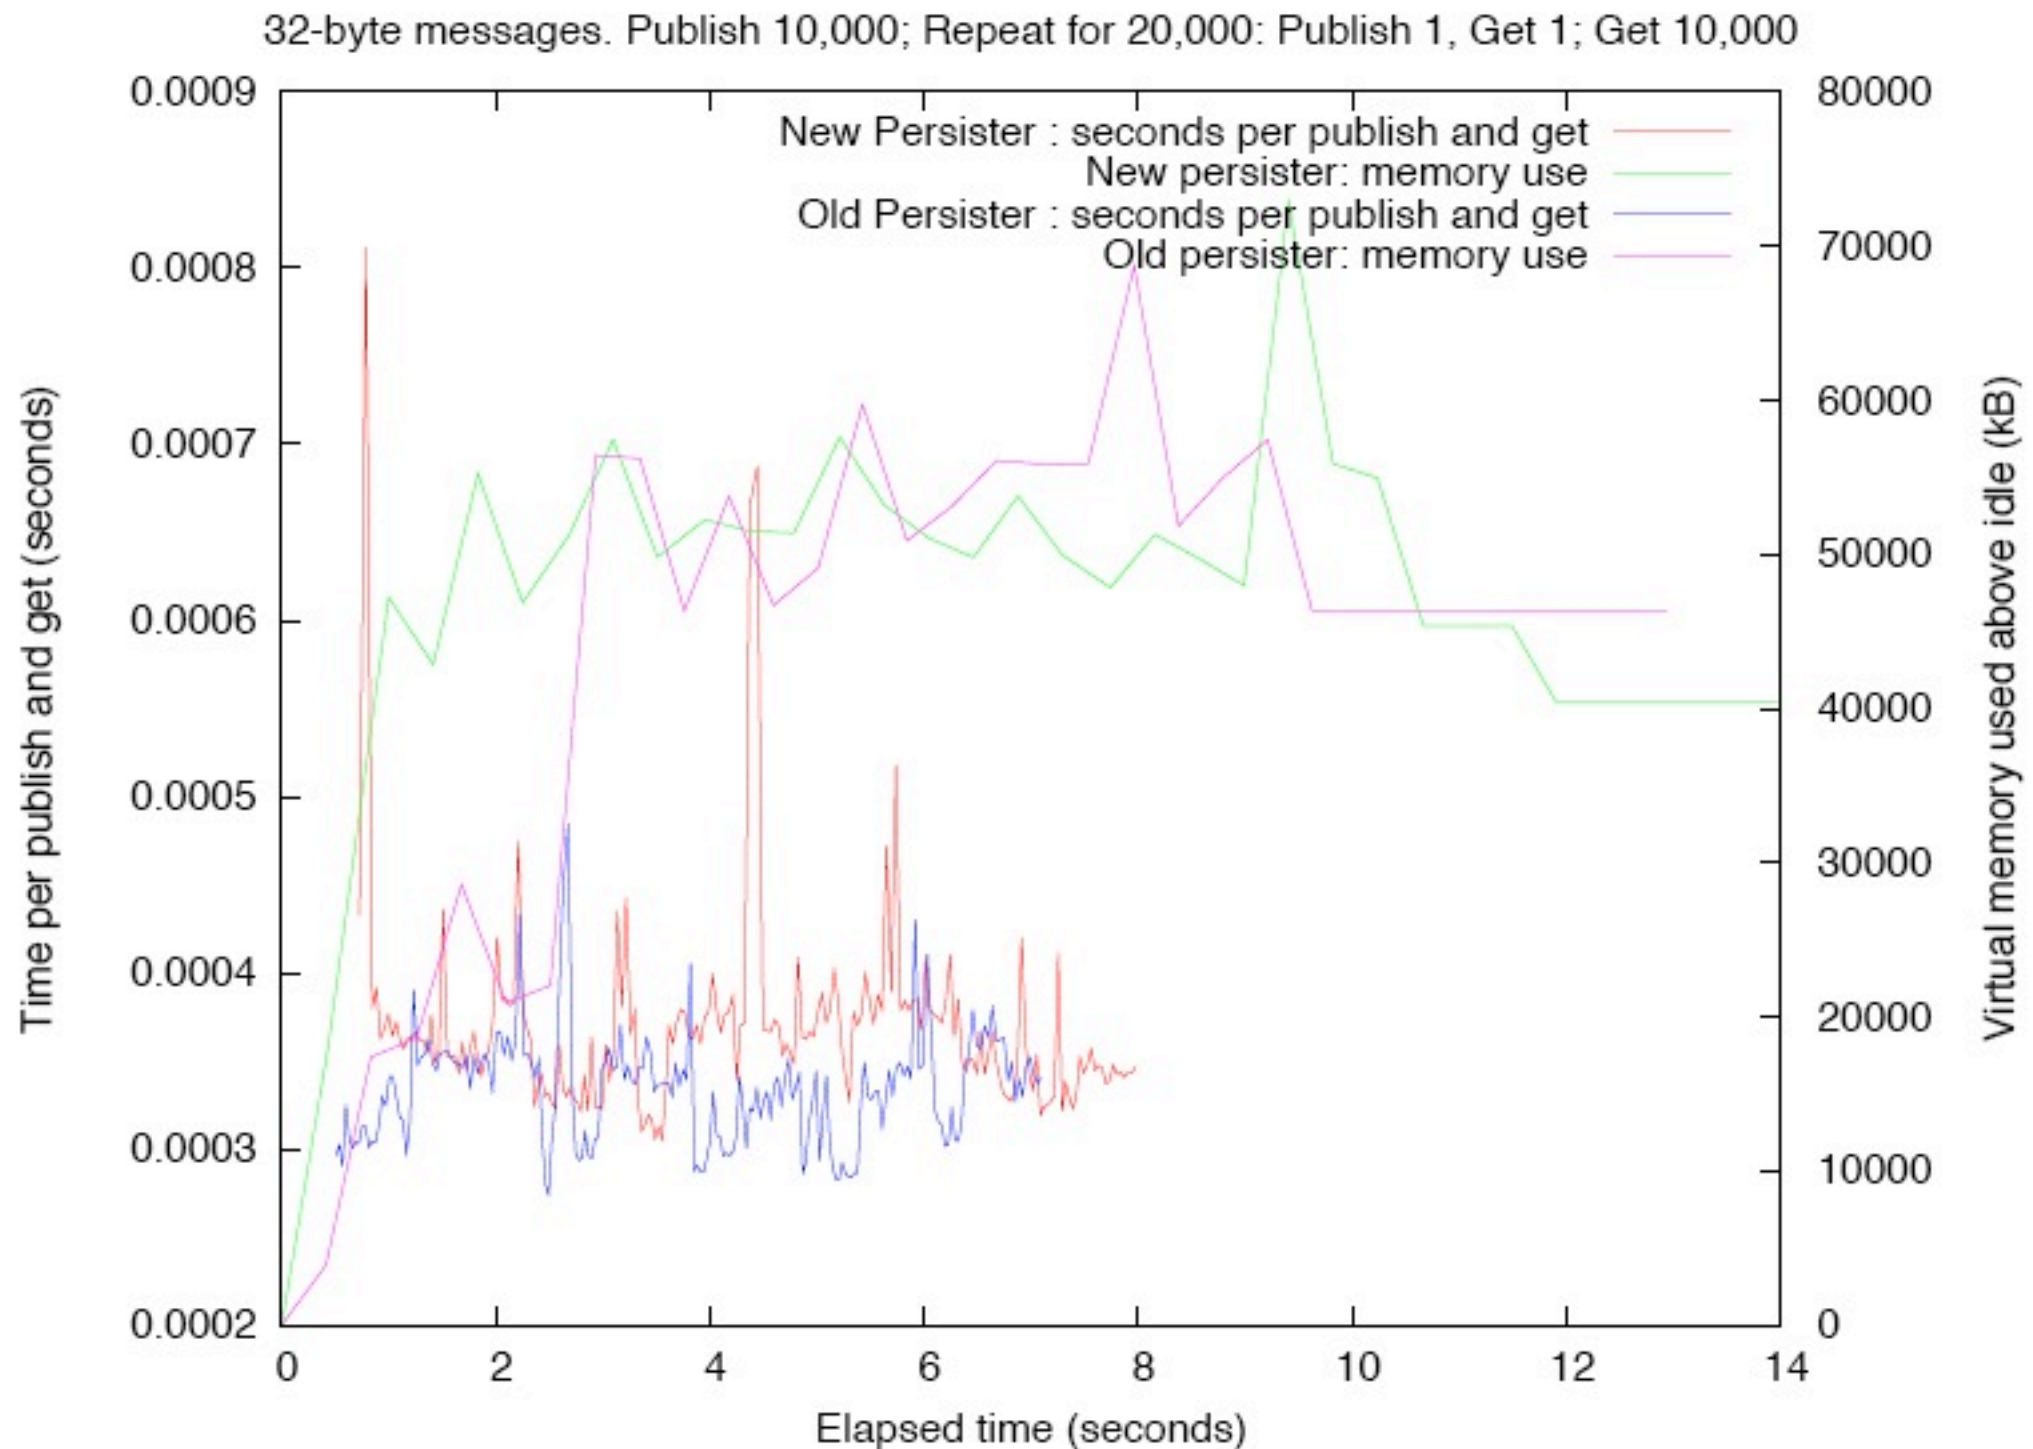
Persister: New

- New persister doesn't slow down with backlog
- New persister doesn't overflow RAM
(but does make good use of the RAM you have installed!)
- Automatic or manual decision when to switch to disk and then disk-only mode
- Carefully tuned to give good default behaviour

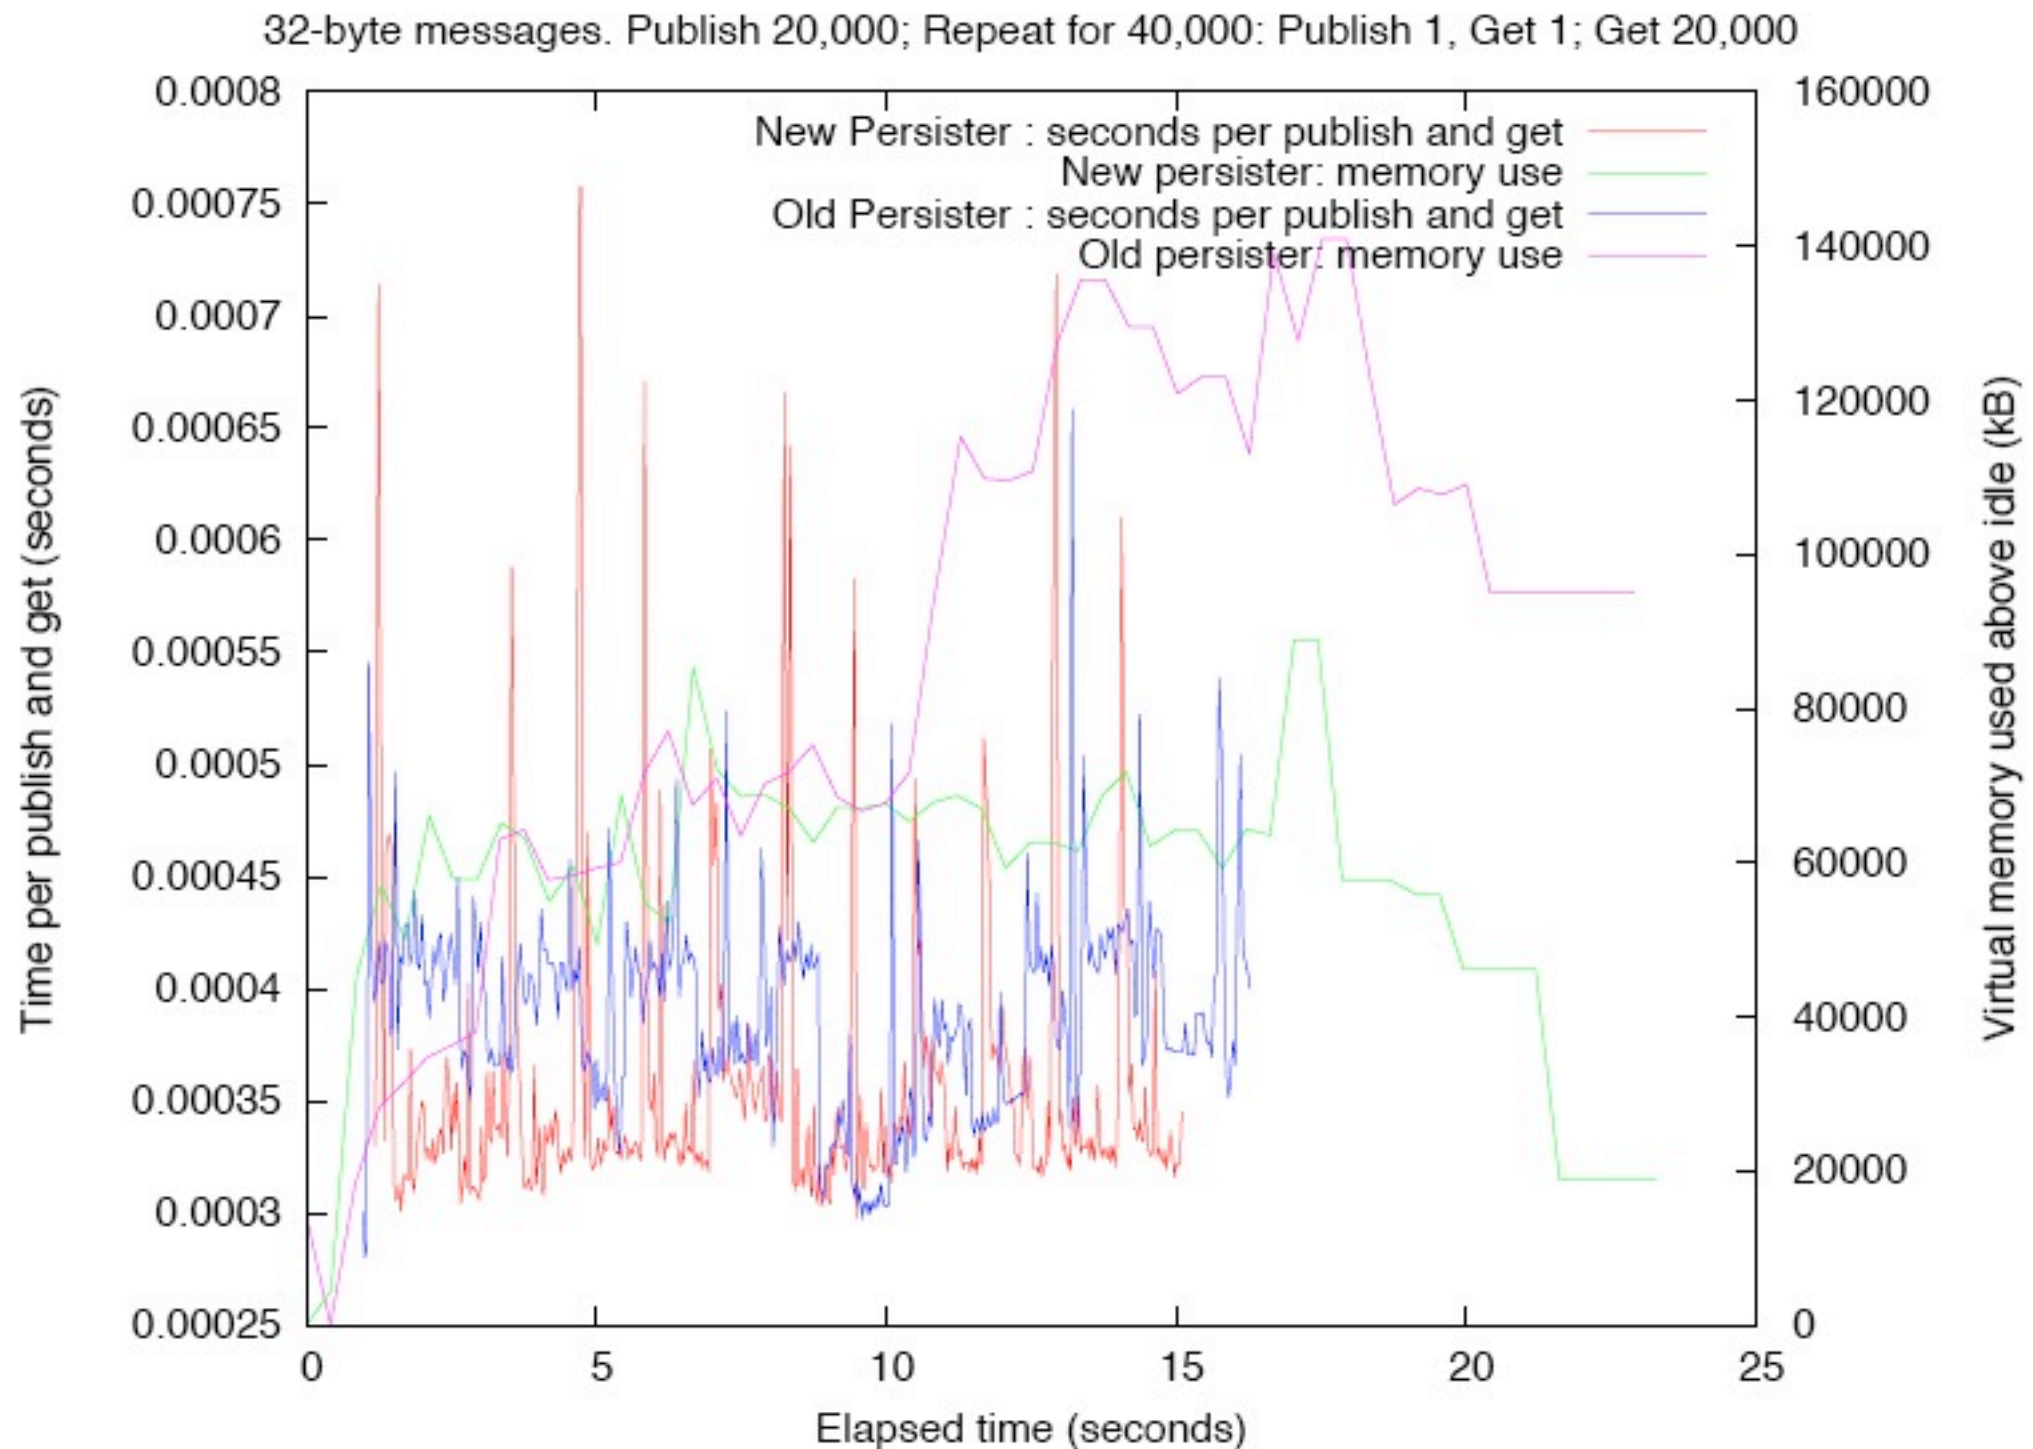
32B, Backlog 5,000



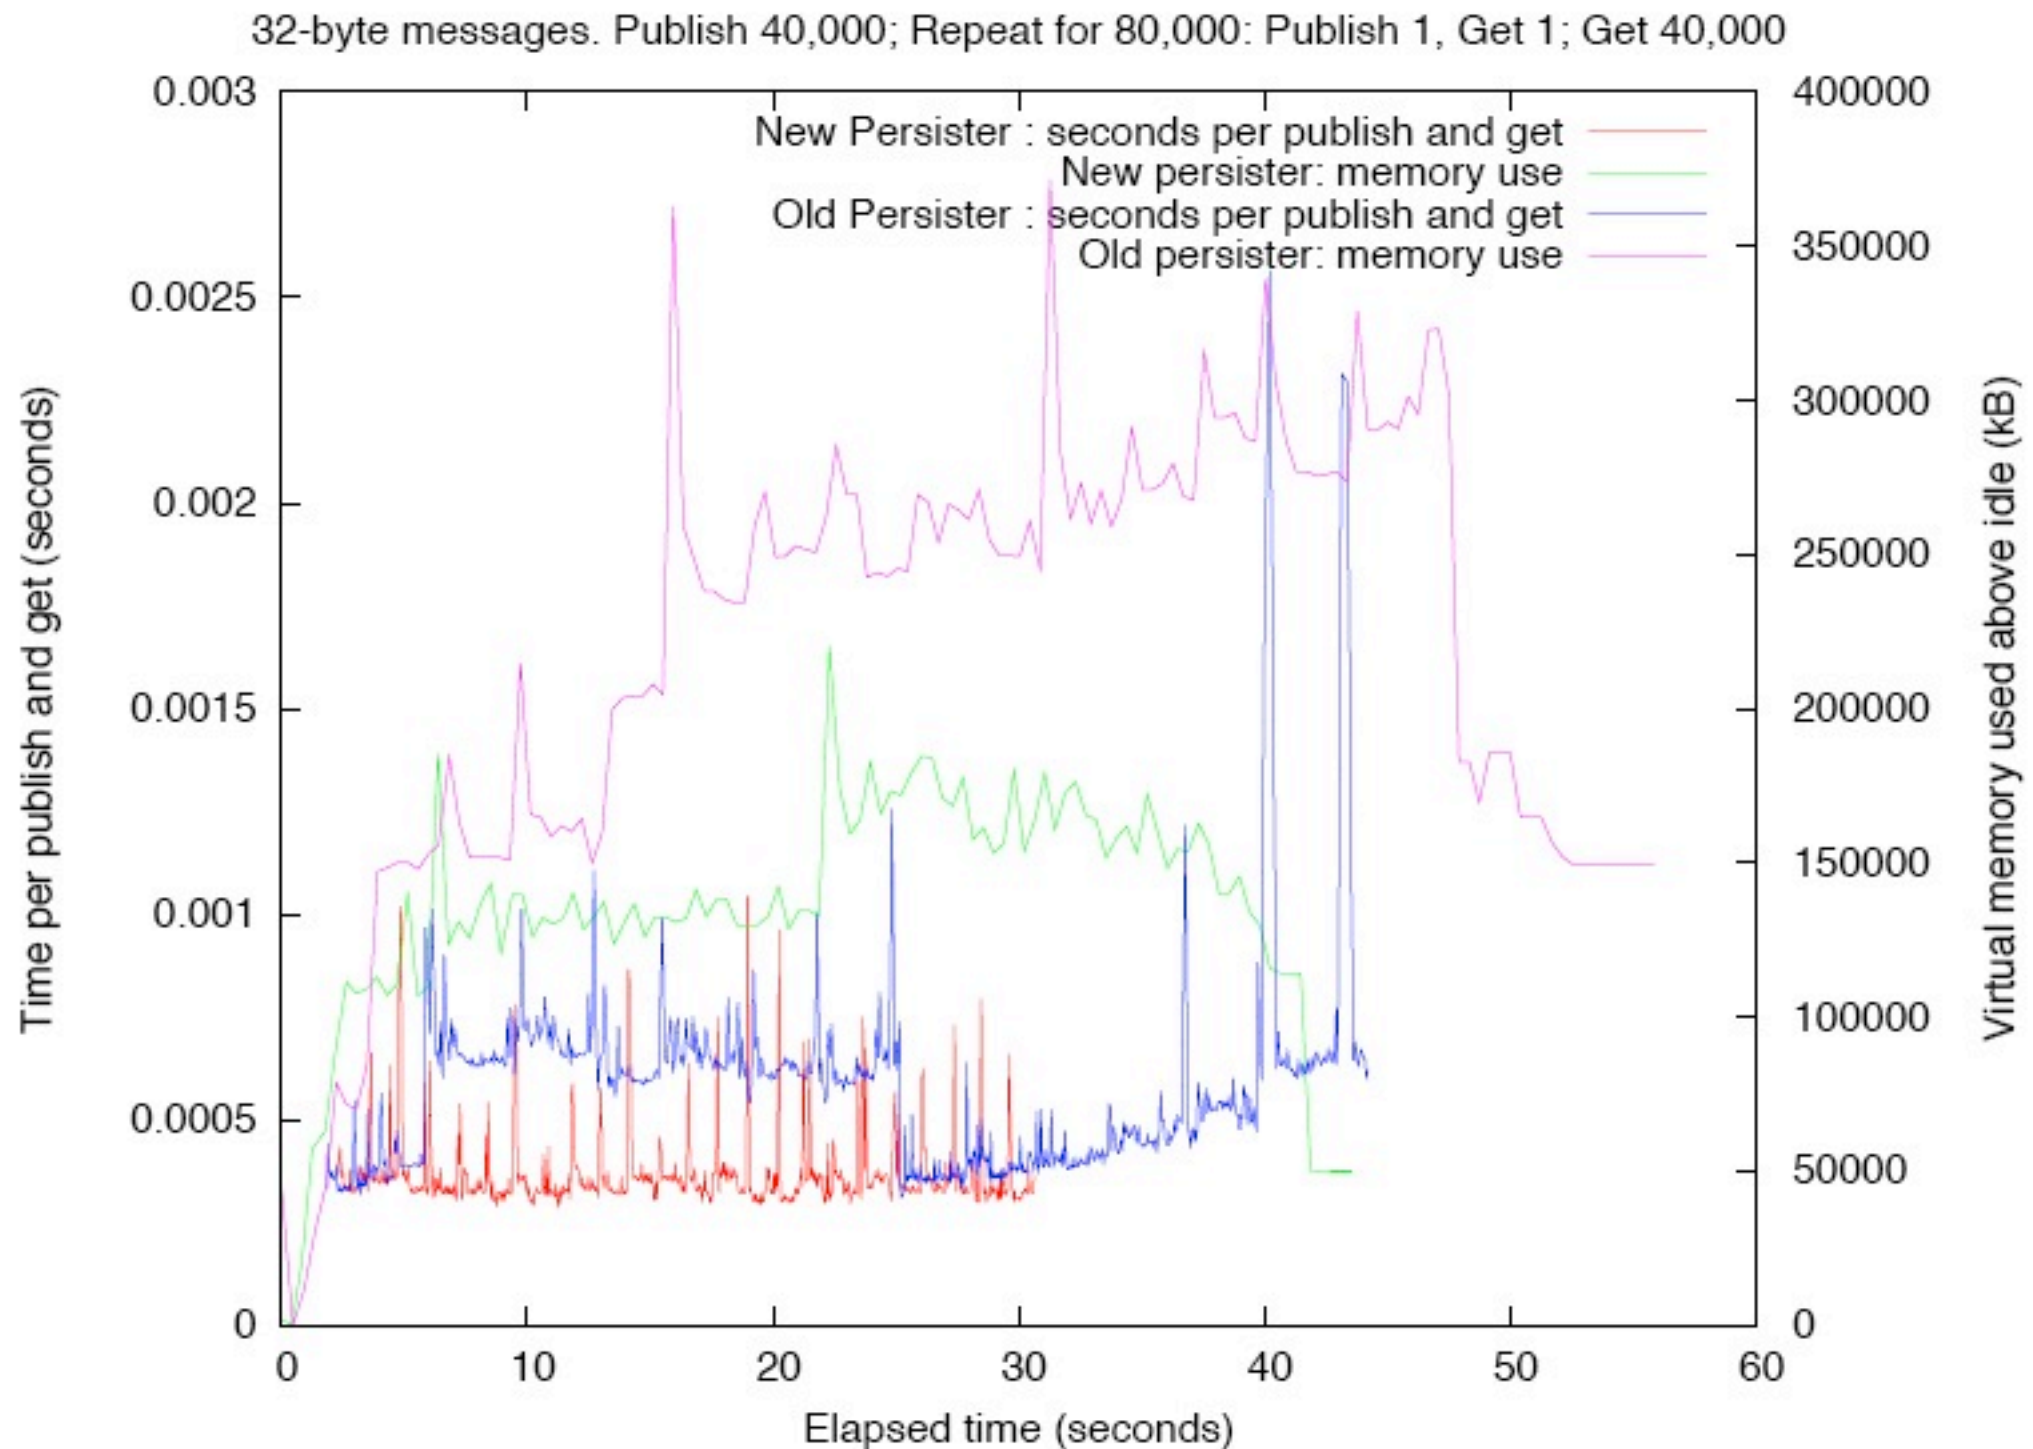
32B, Backlog 10,000



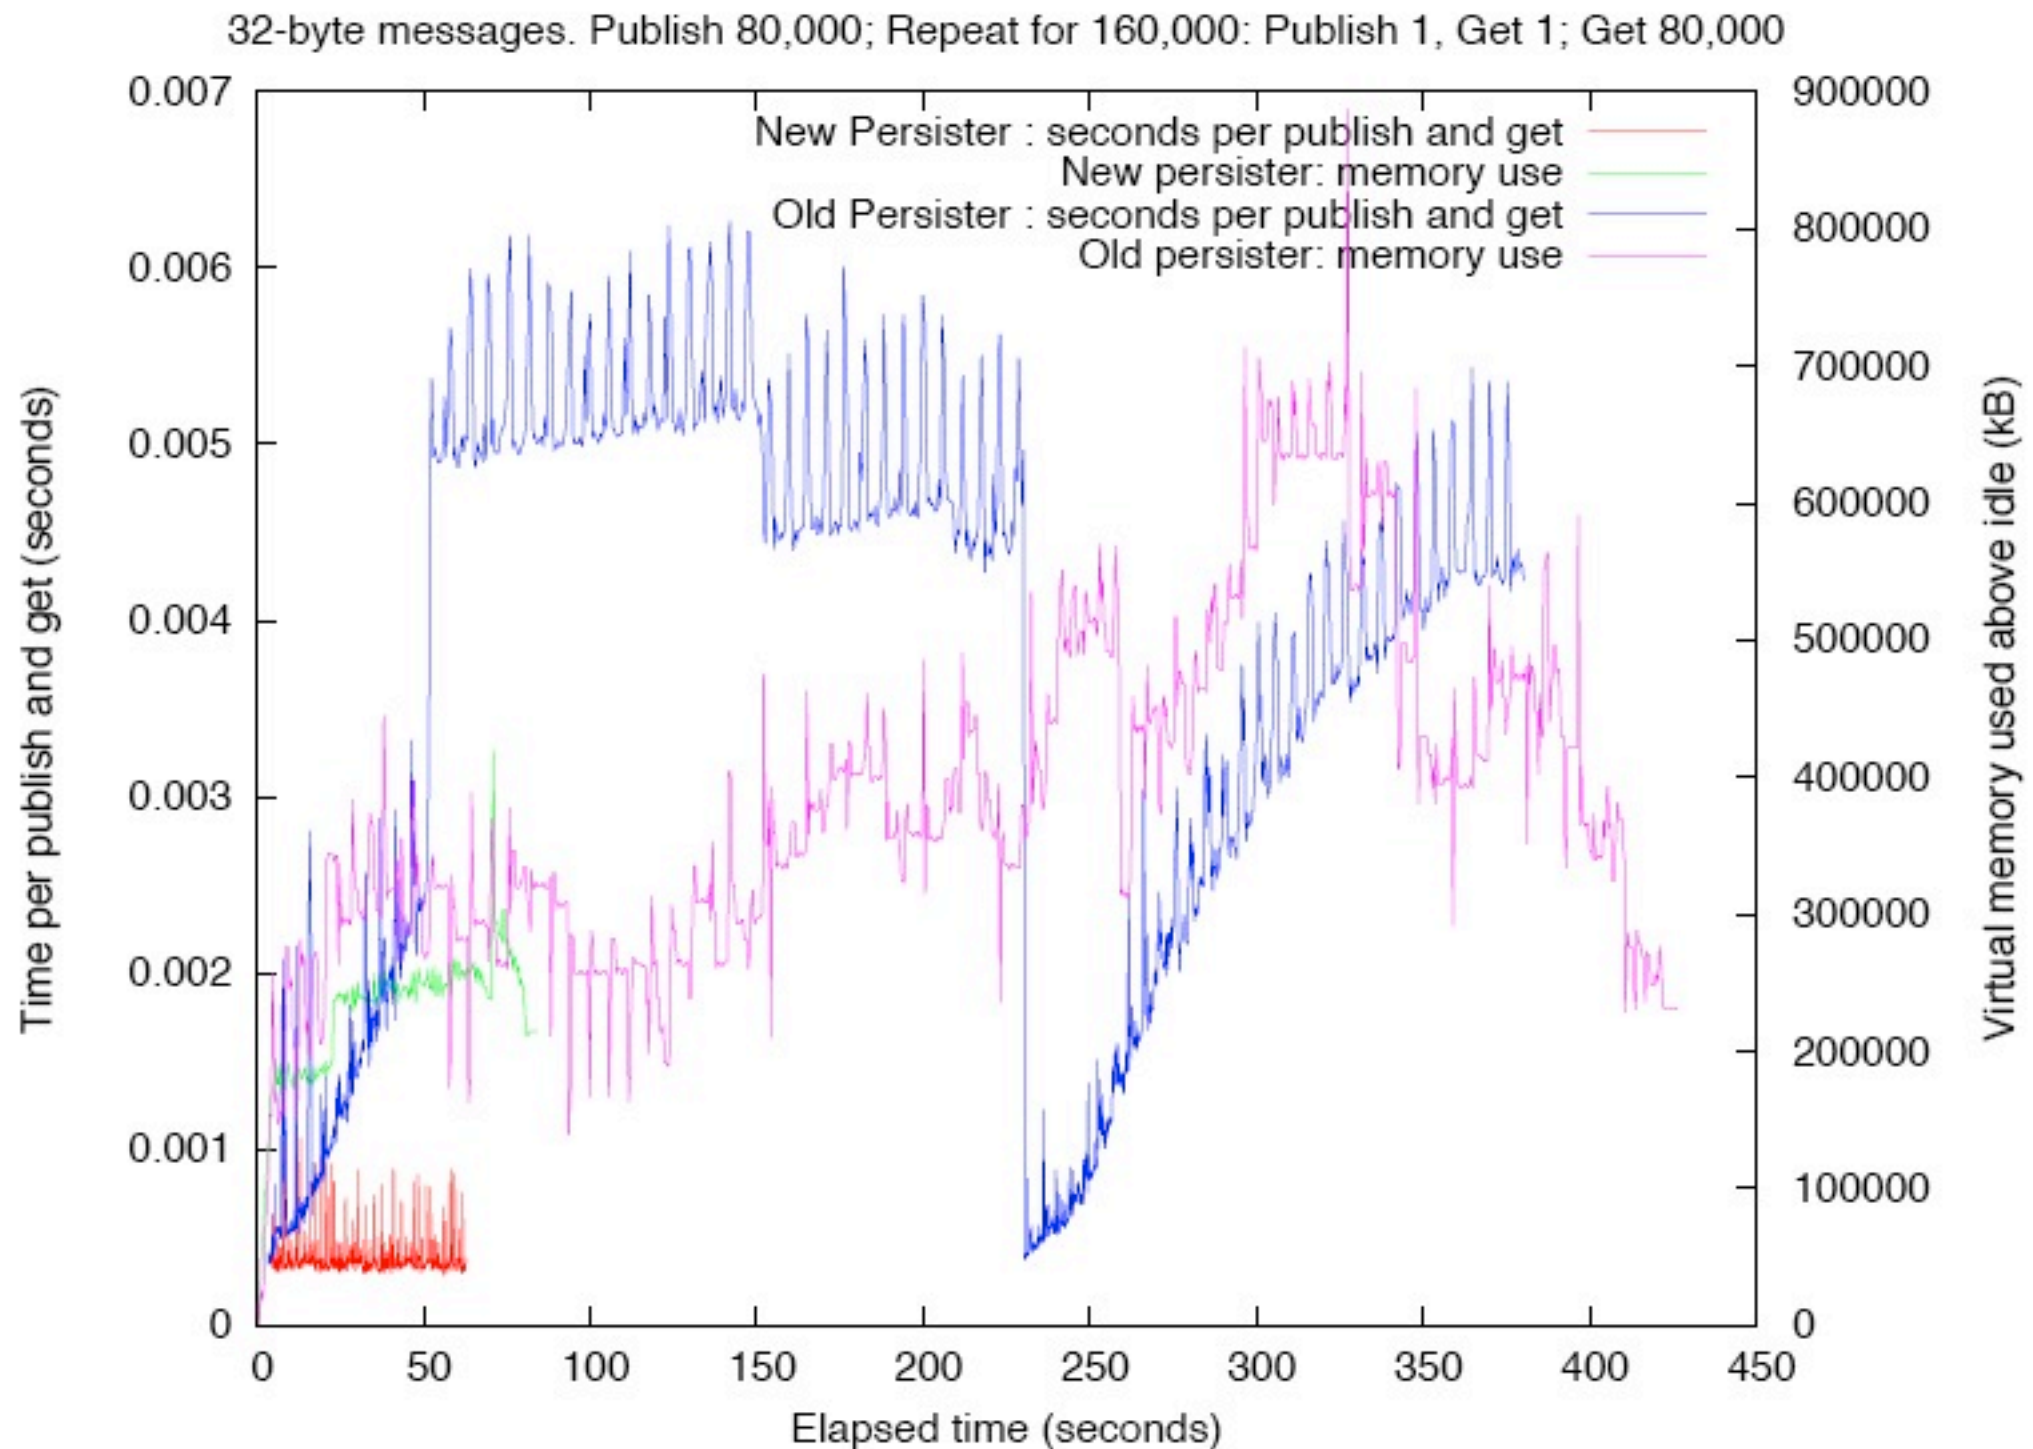
32B, Backlog 20,000



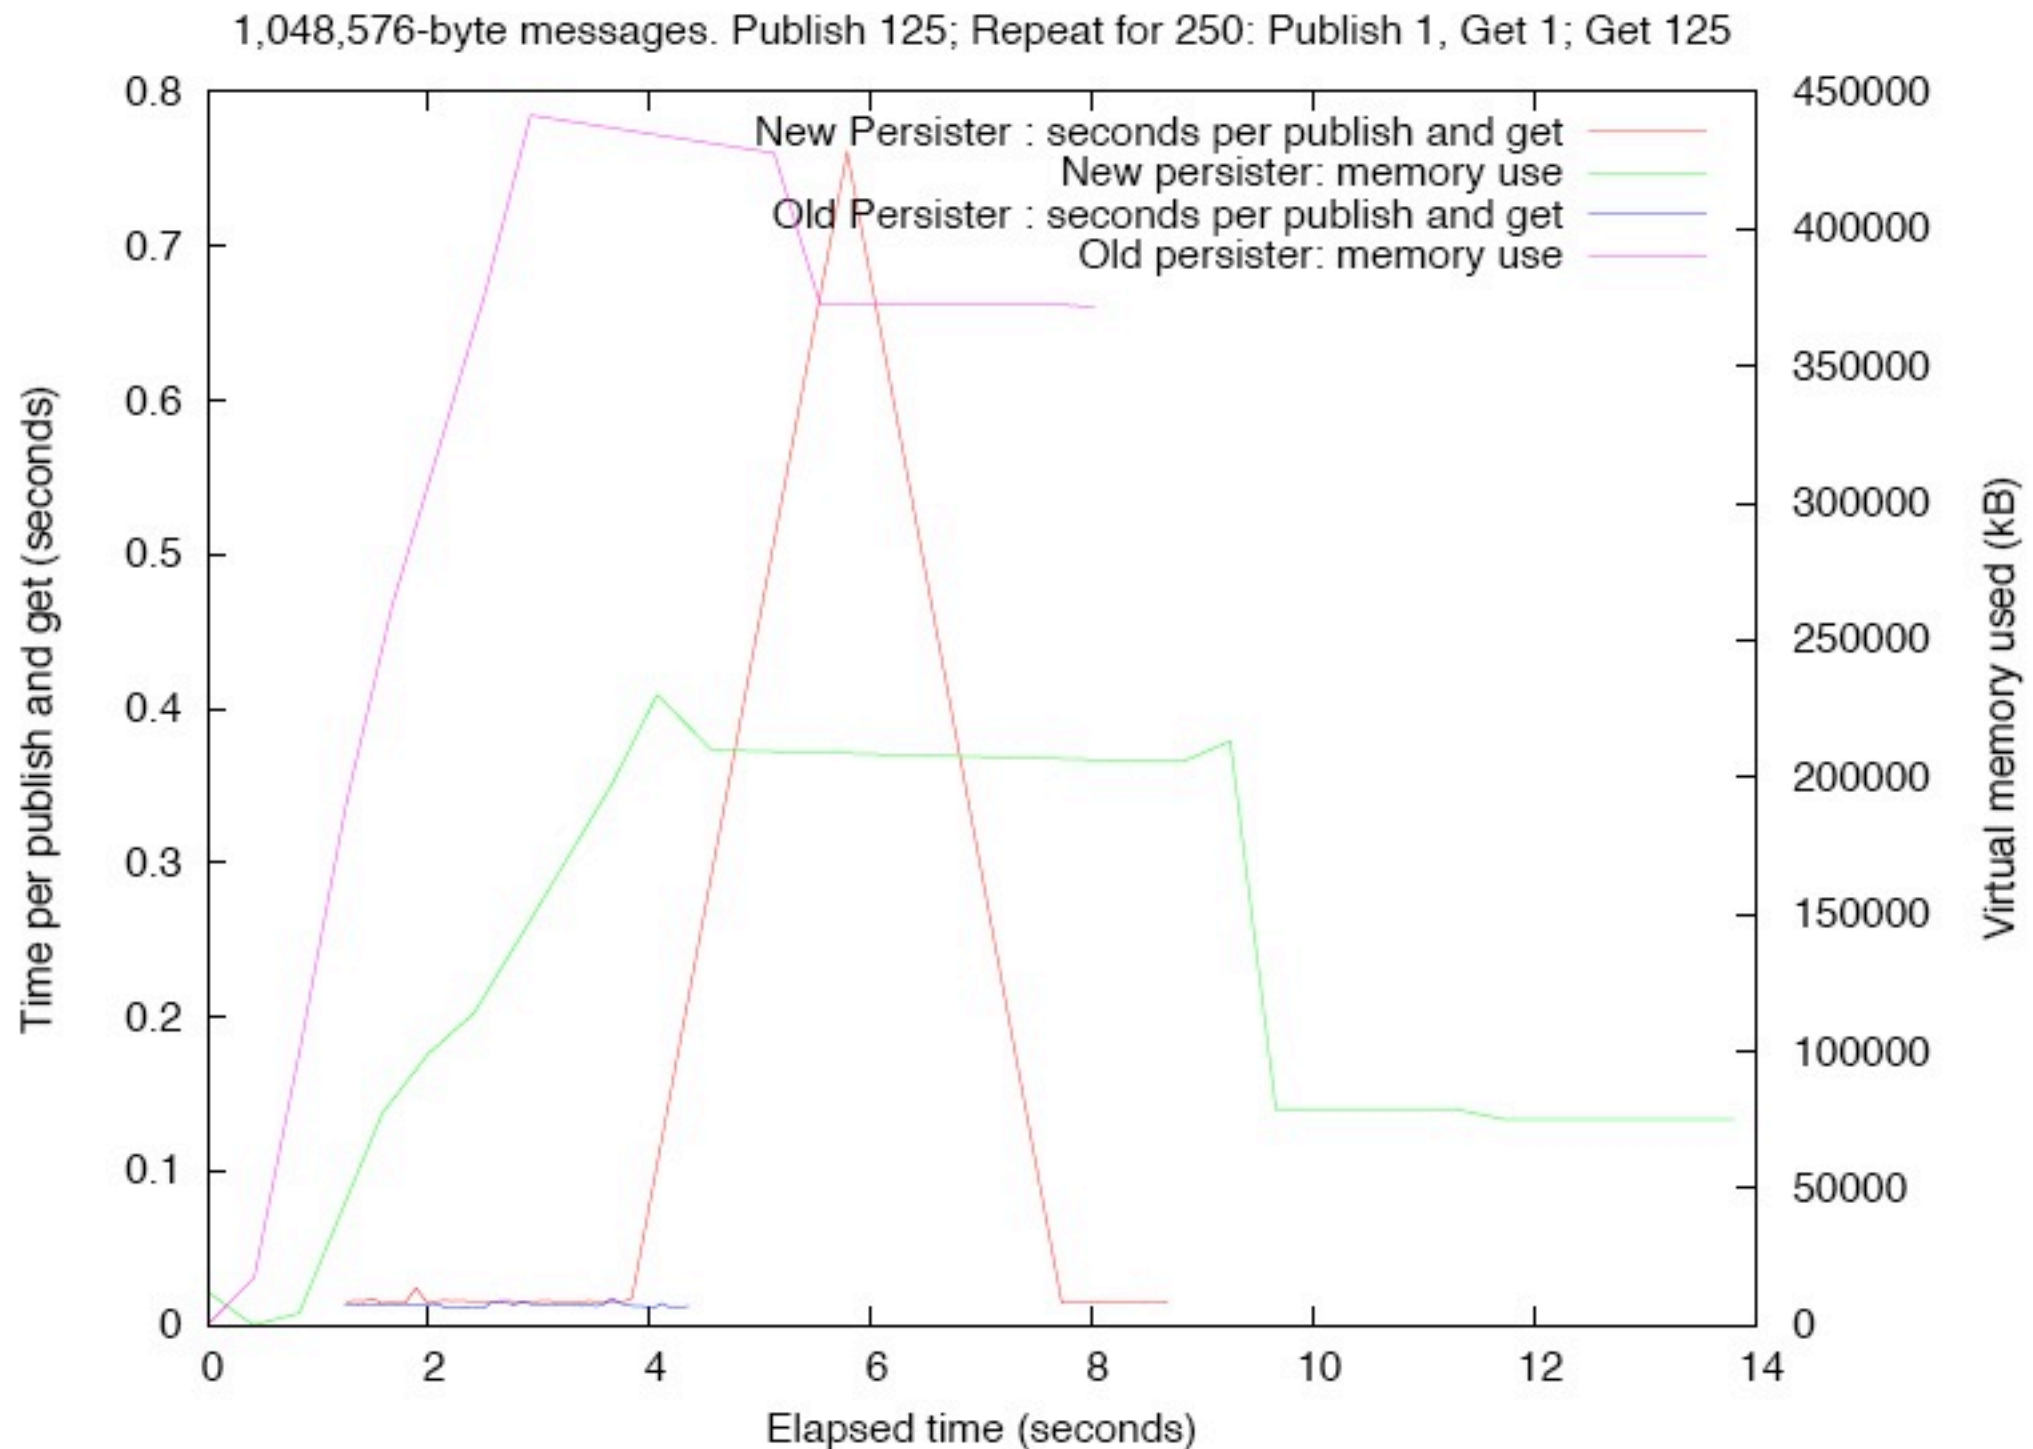
32B, Backlog 40,000



32B, Backlog 80,000

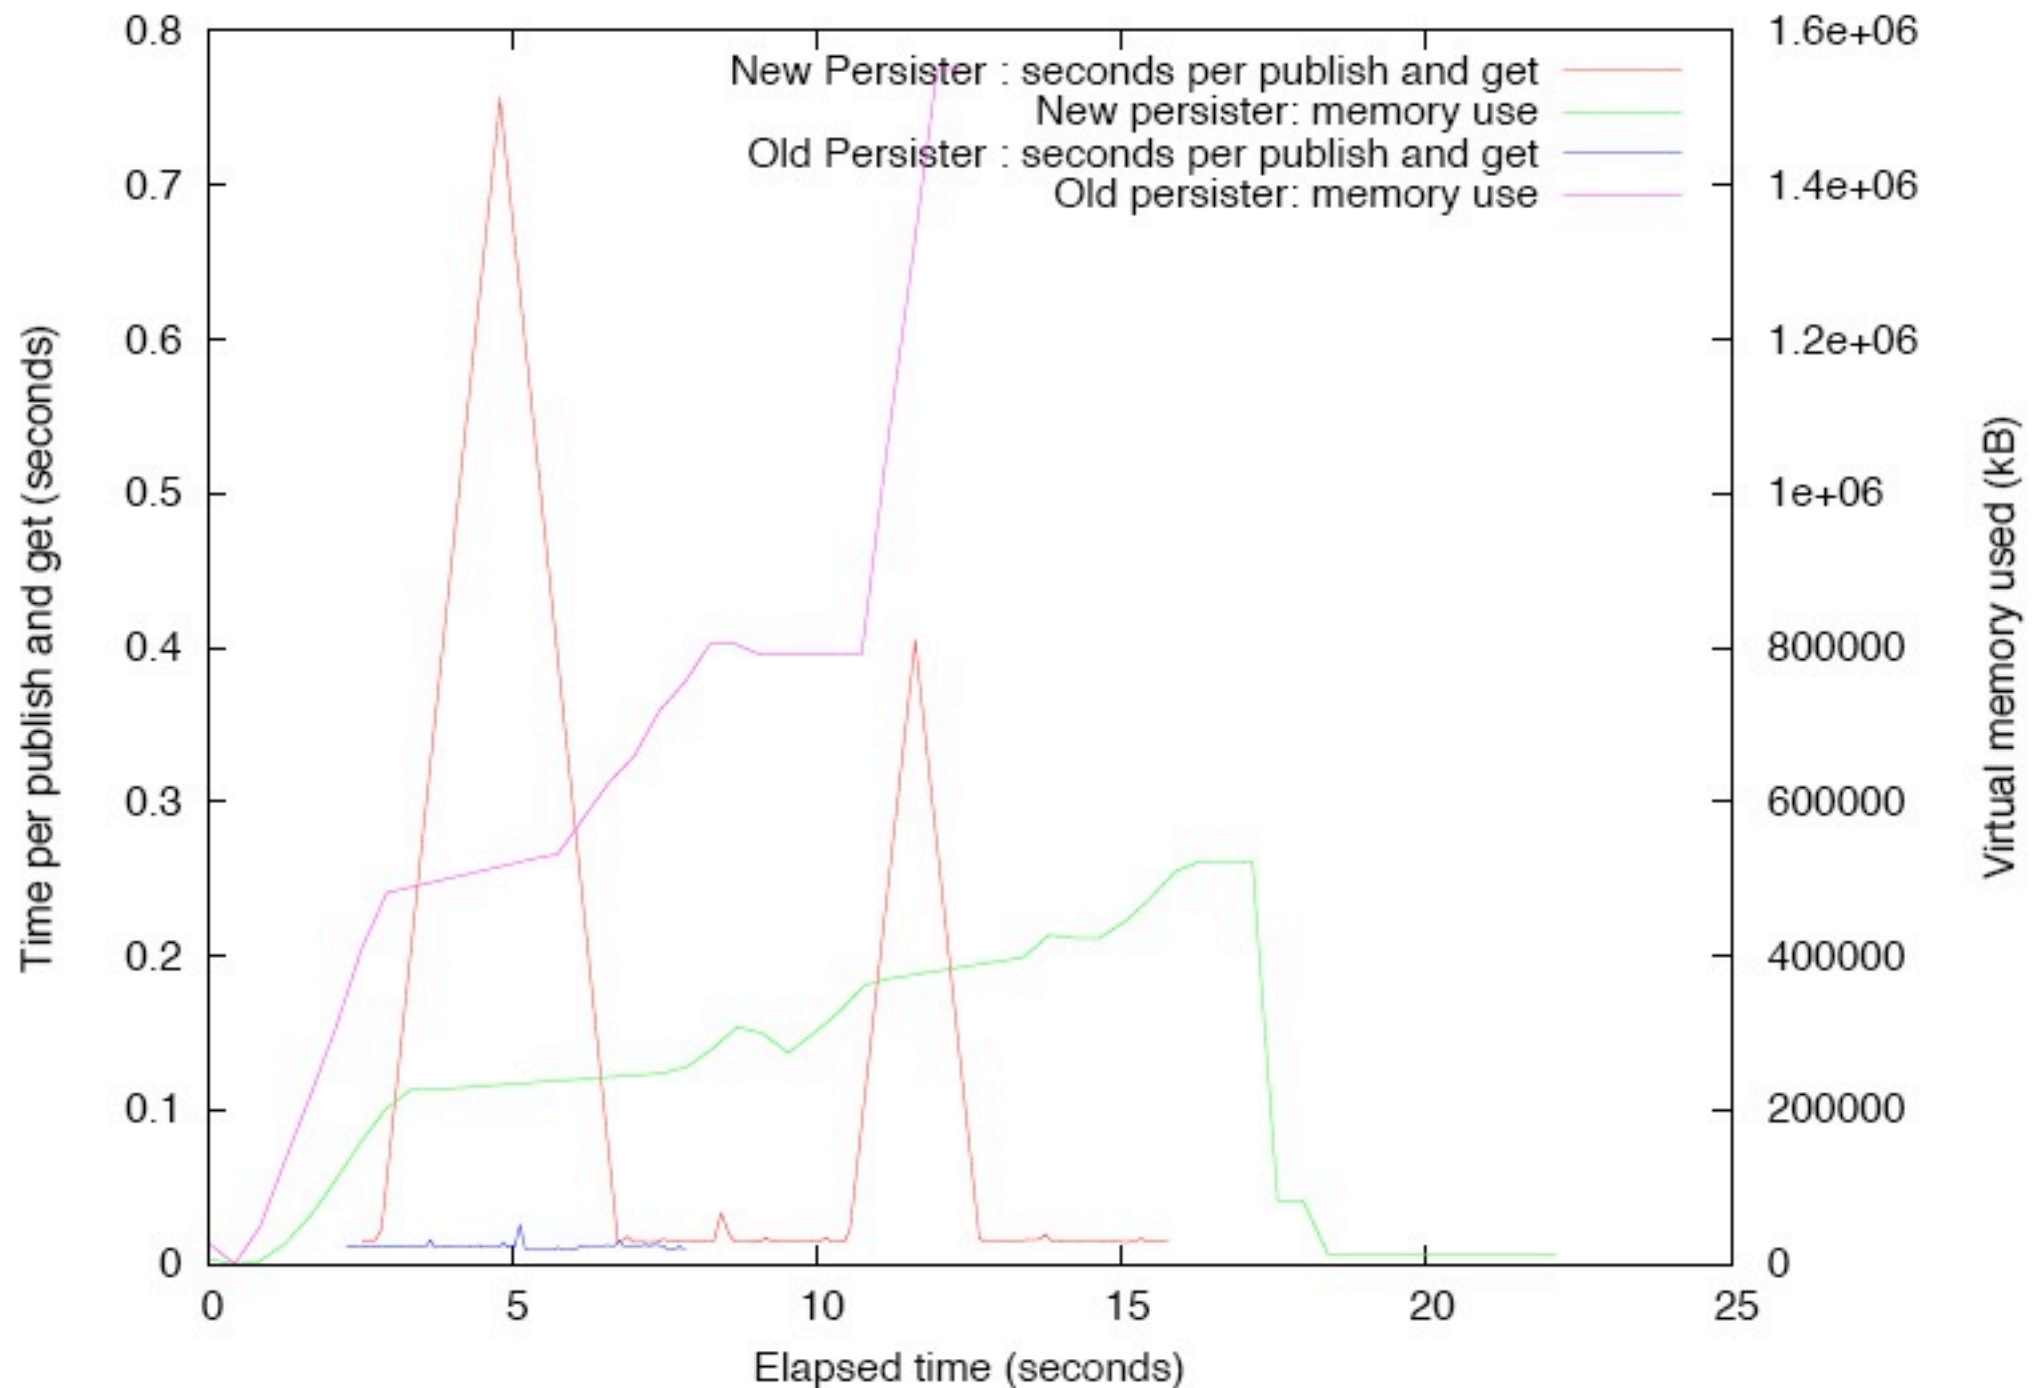


1 MB, Backlog 125



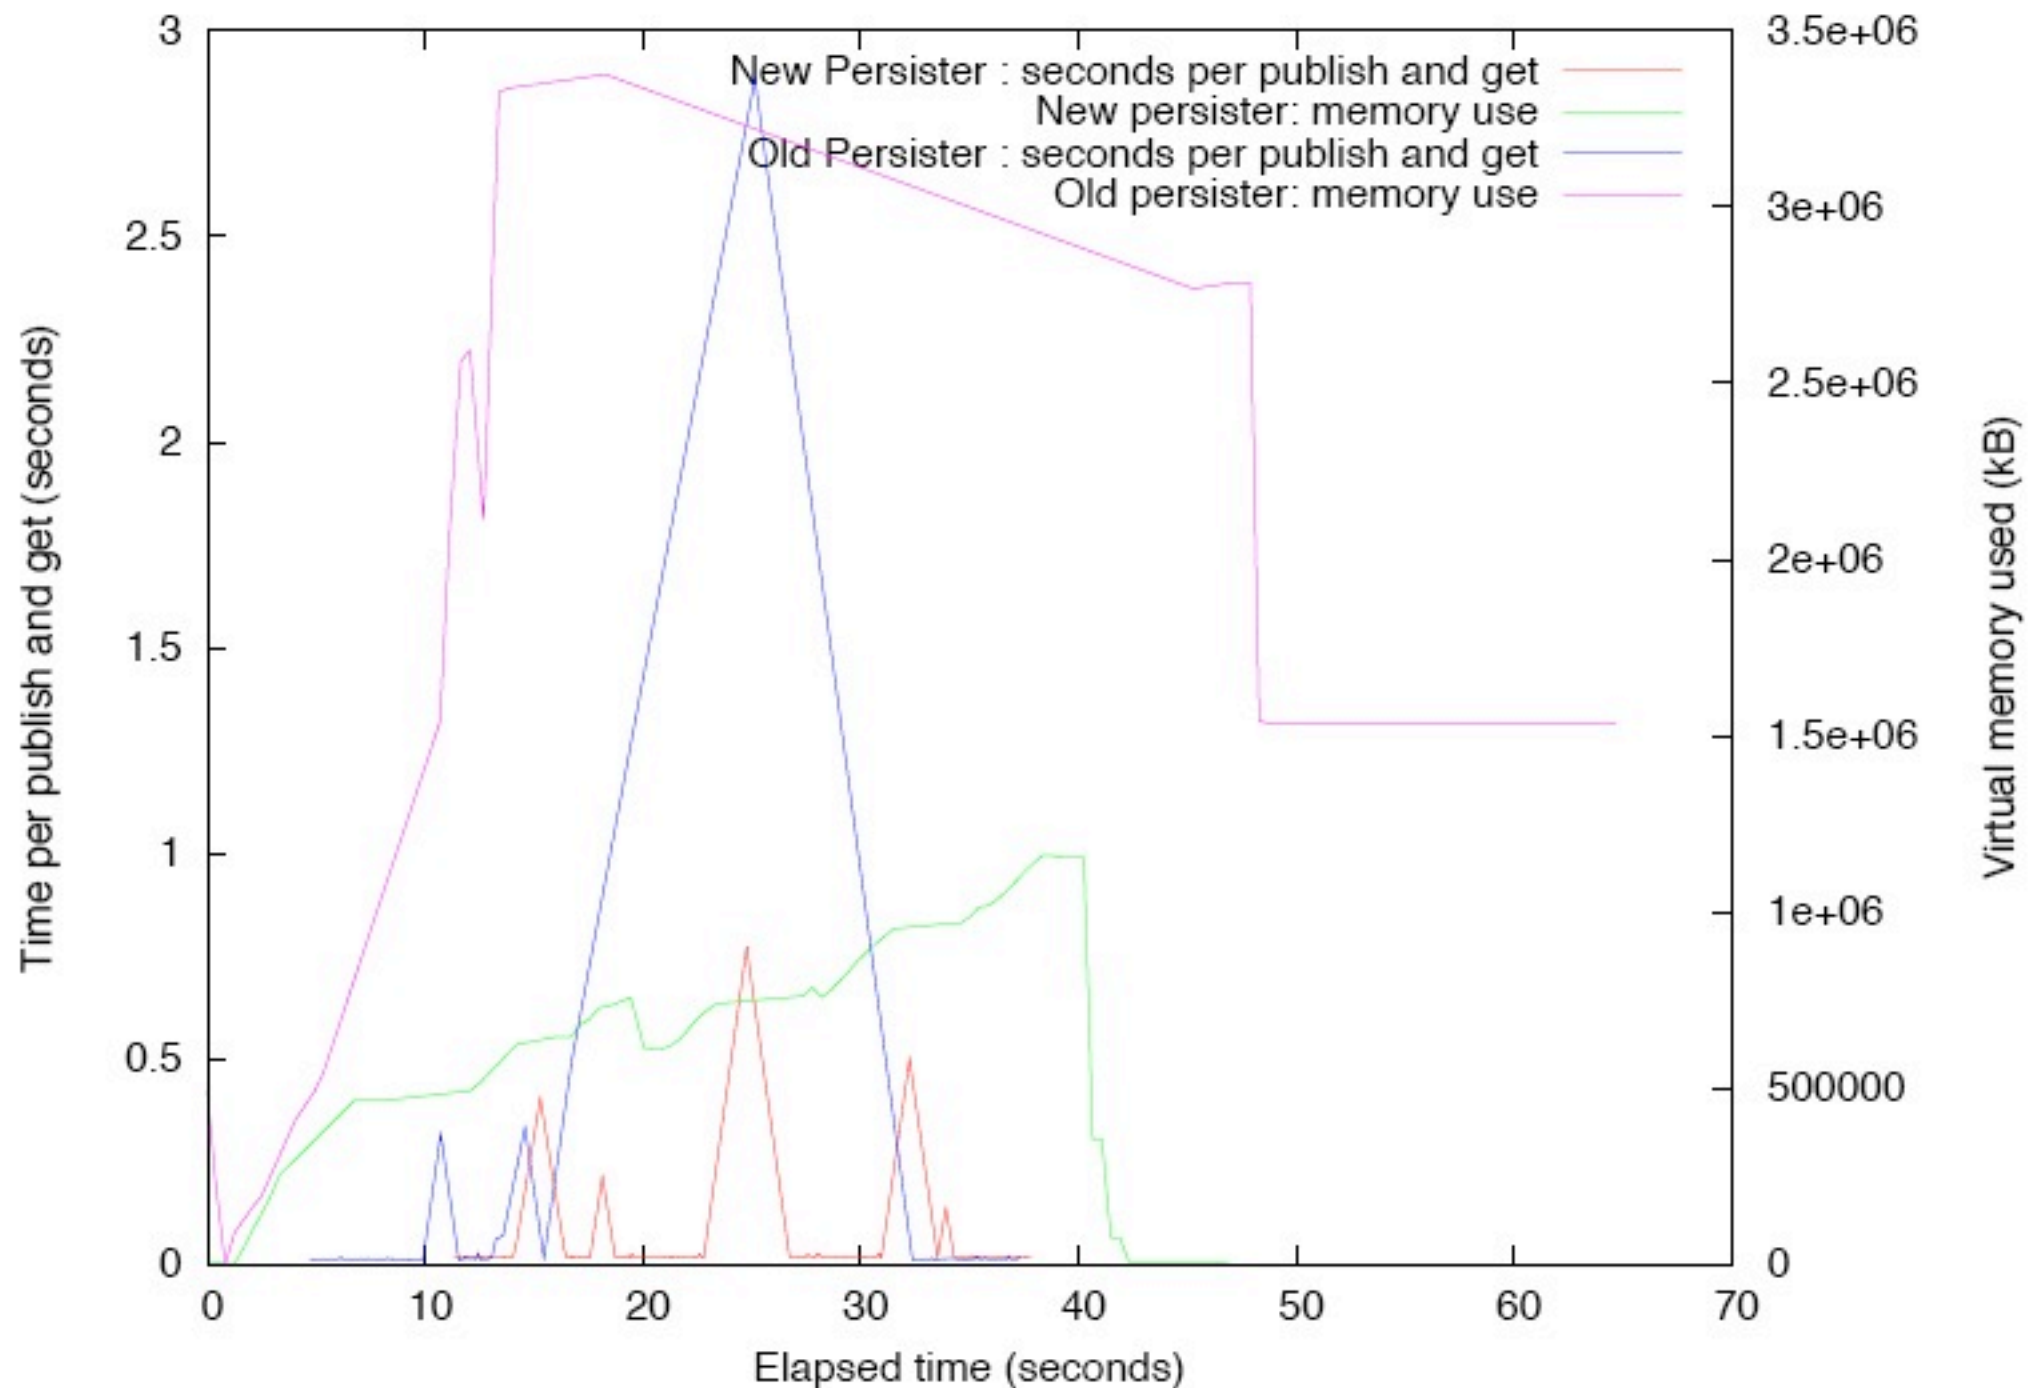
1MB, Backlog 250

1,048,576-byte messages. Publish 250; Repeat for 500: Publish 1, Get 1; Get 250



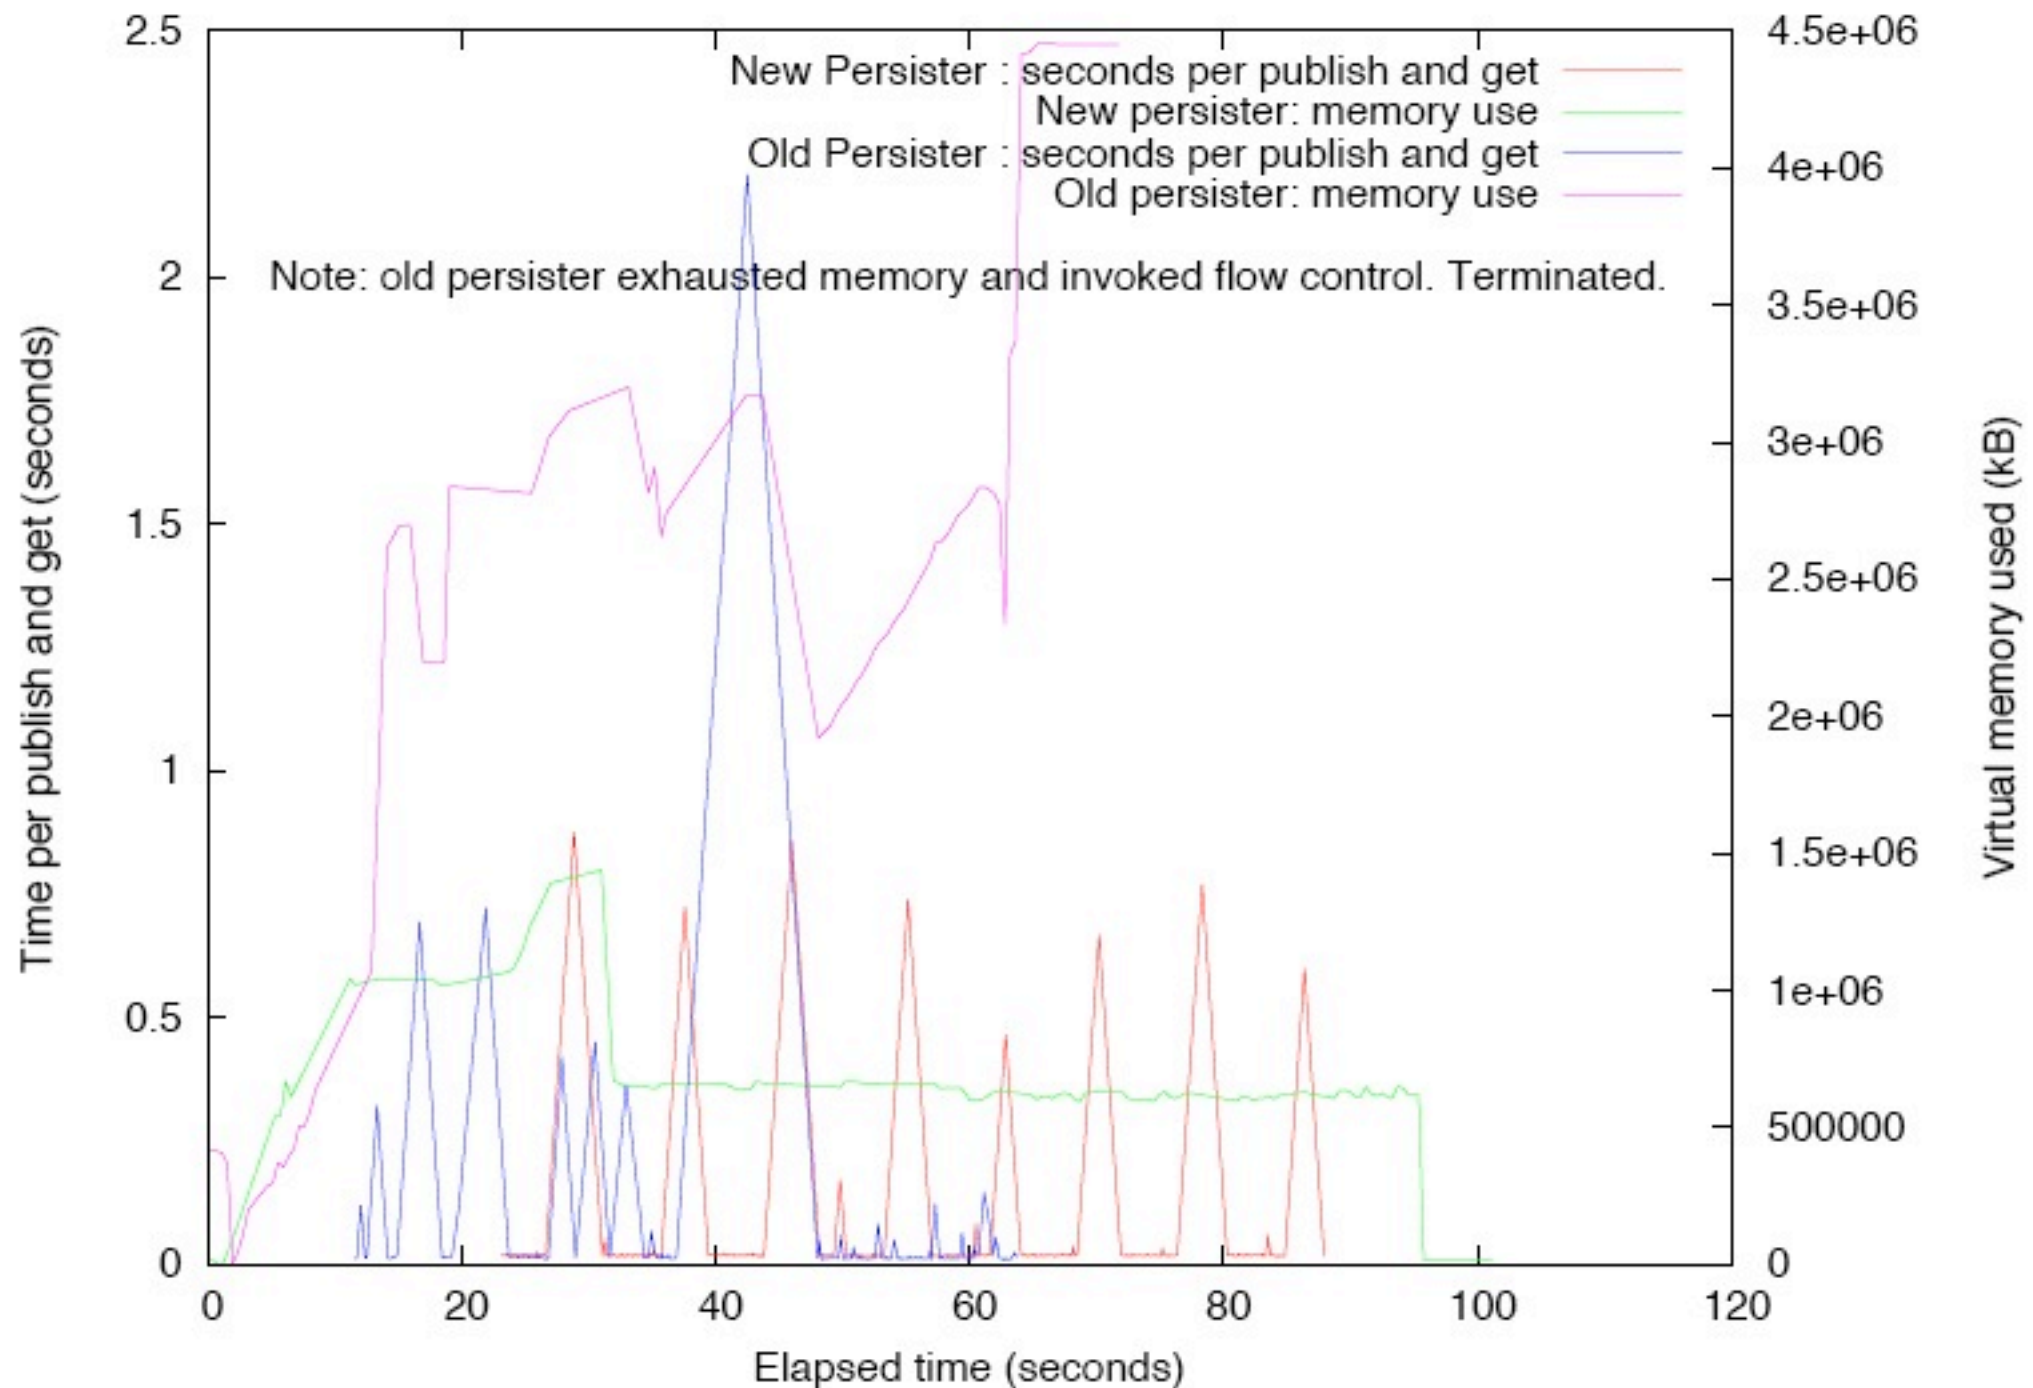
1MB, Backlog 500

1,048,576-byte messages. Publish 500; Repeat for 1000: Publish 1, Get 1; Get 500

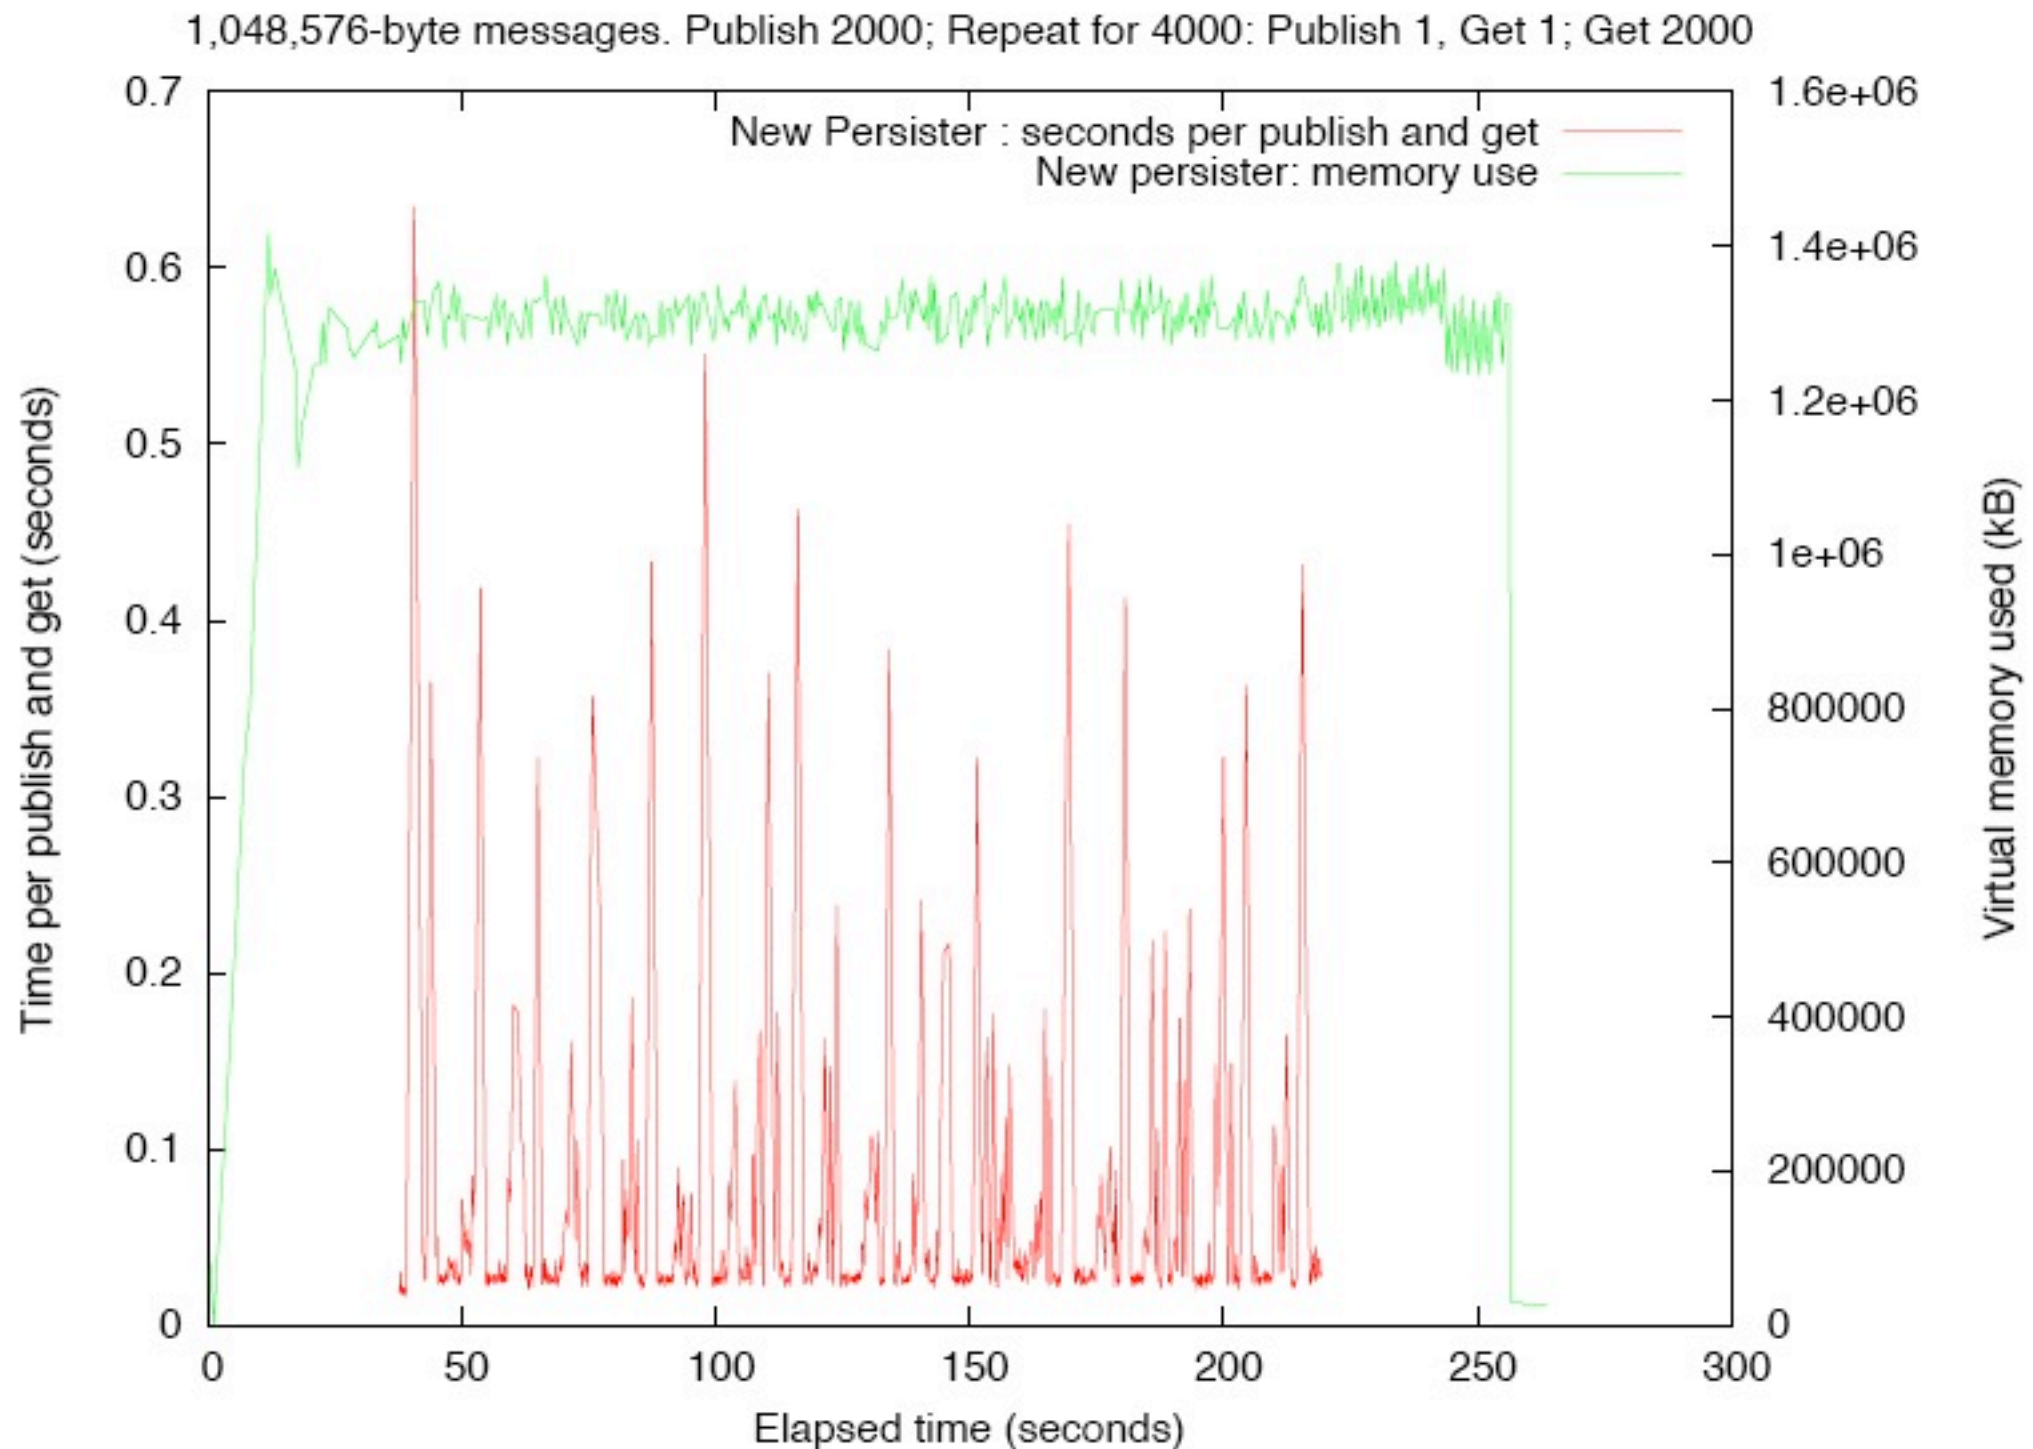


1MB, Backlog 1,000

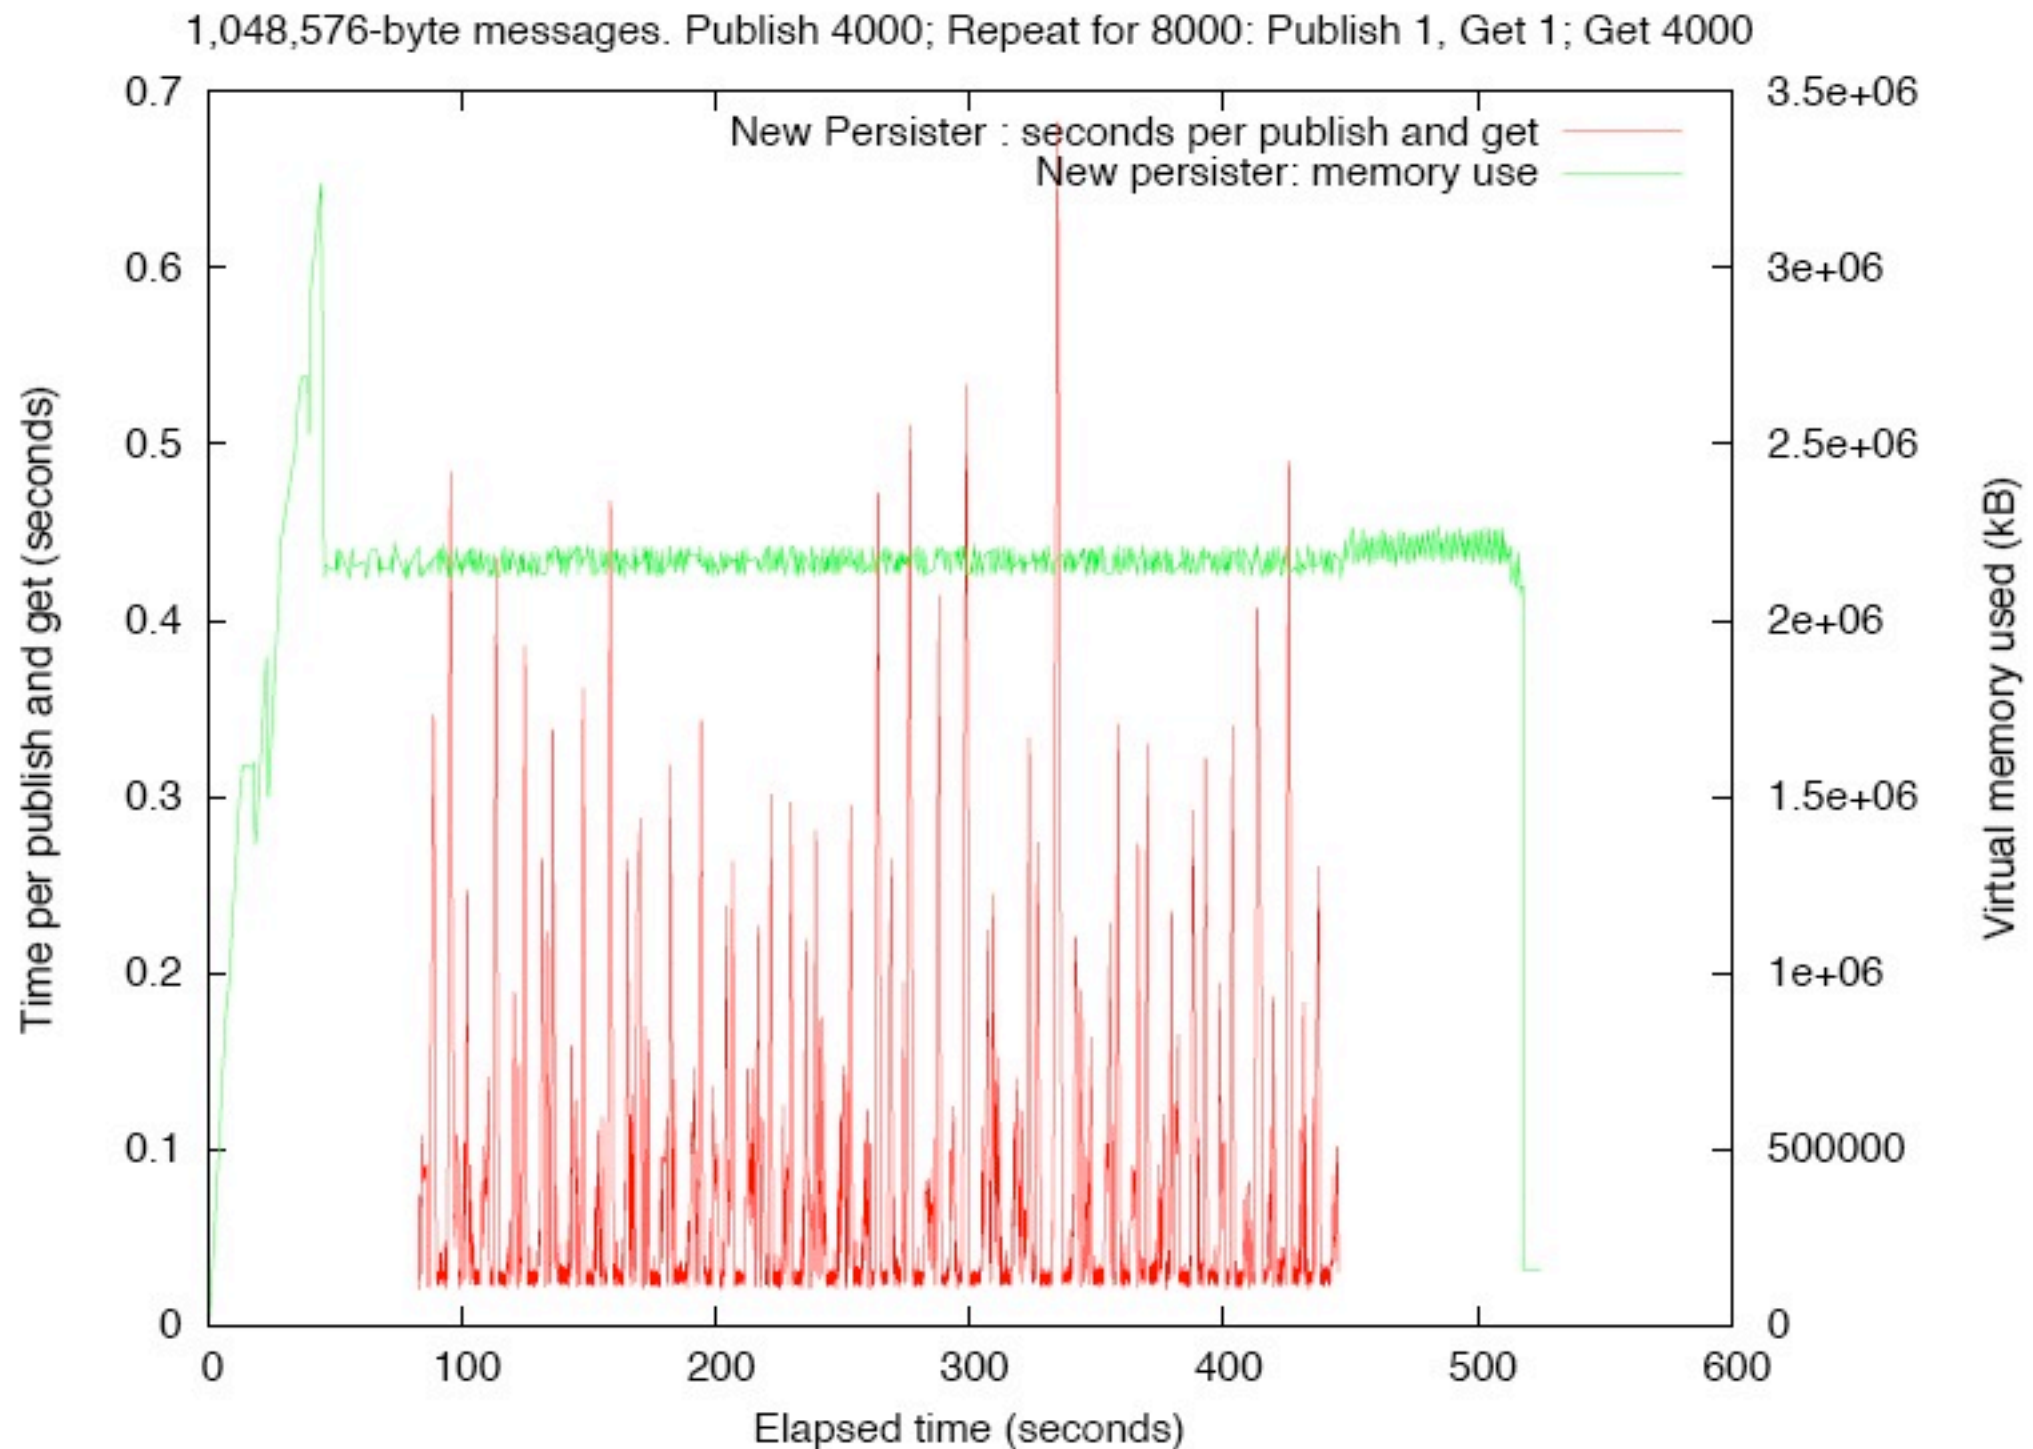
1,048,576-byte messages. Publish 1000; Repeat for 2000: Publish 1, Get 1; Get 1000



1 MB, Backlog 2,000



1 MB, Backlog 4,000



Persister: Summary

- Fast (slower than old persister, but still fast)
- Scales to fill your disks without filling RAM
- In QA at the moment; will land for 1.7

Security

Resource ACLs

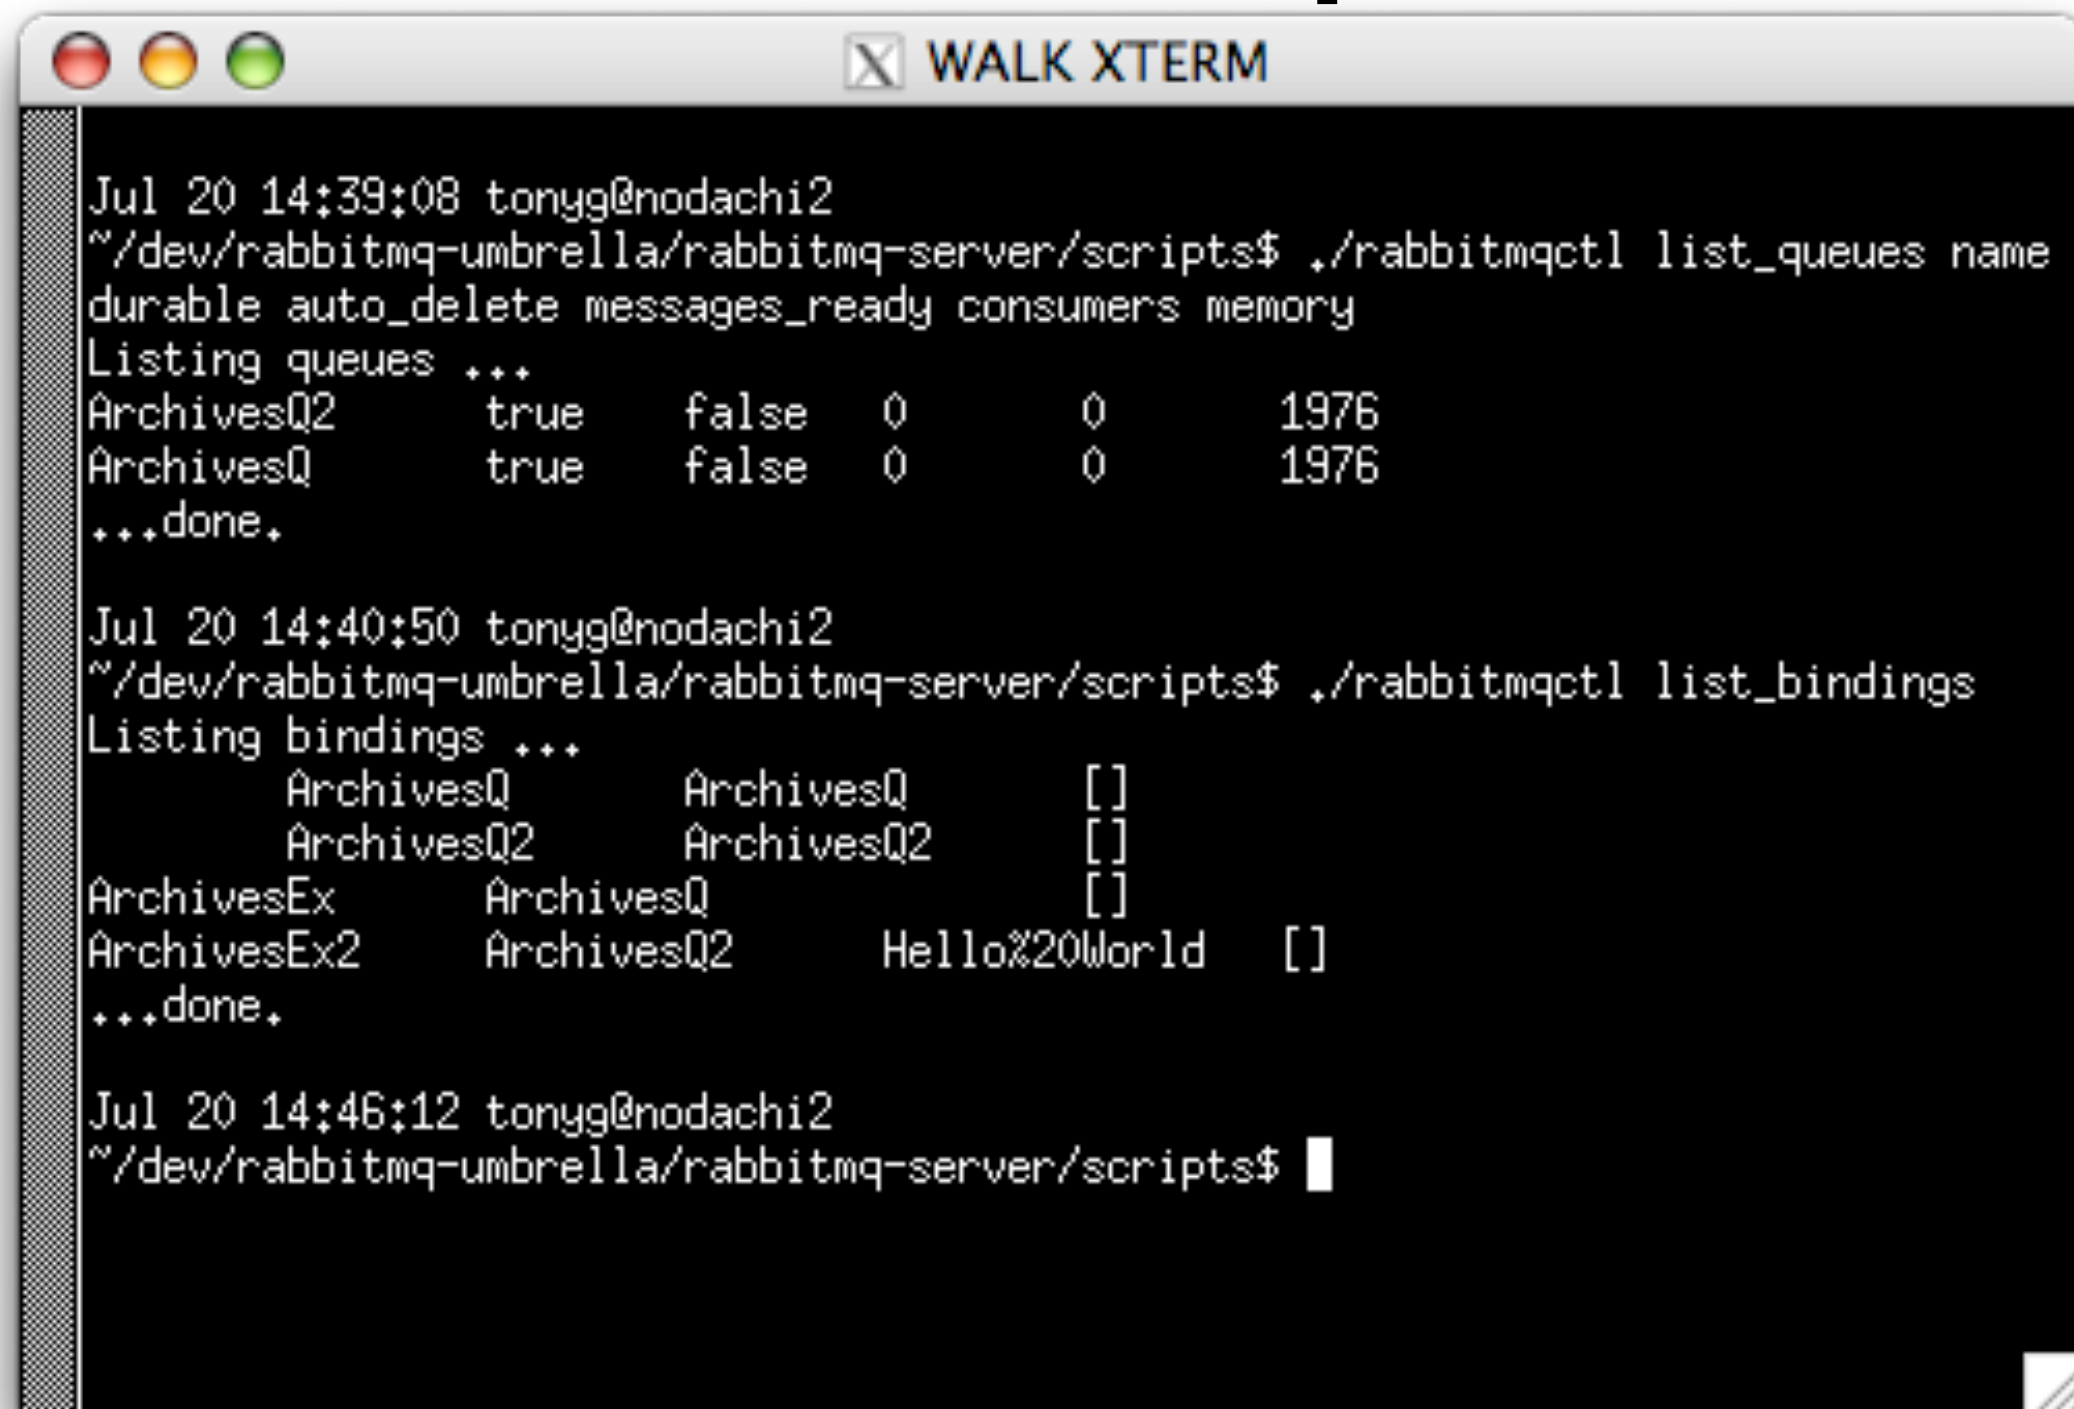
- Users are granted access to each vhost
- Each user has three regular expressions:
 - for Administrative actions (create, delete)
 - for Reading (bind from exch, consume from queue)
 - for Writing (publish to exch, bind to queue)

Encryption, PKI

- Securing the Transport
 - stunnel4
 - native SSL support in final QA
- Securing the Messages
 - key directory + per-message encryption and signing

Operational Tools

rabbitmqctl

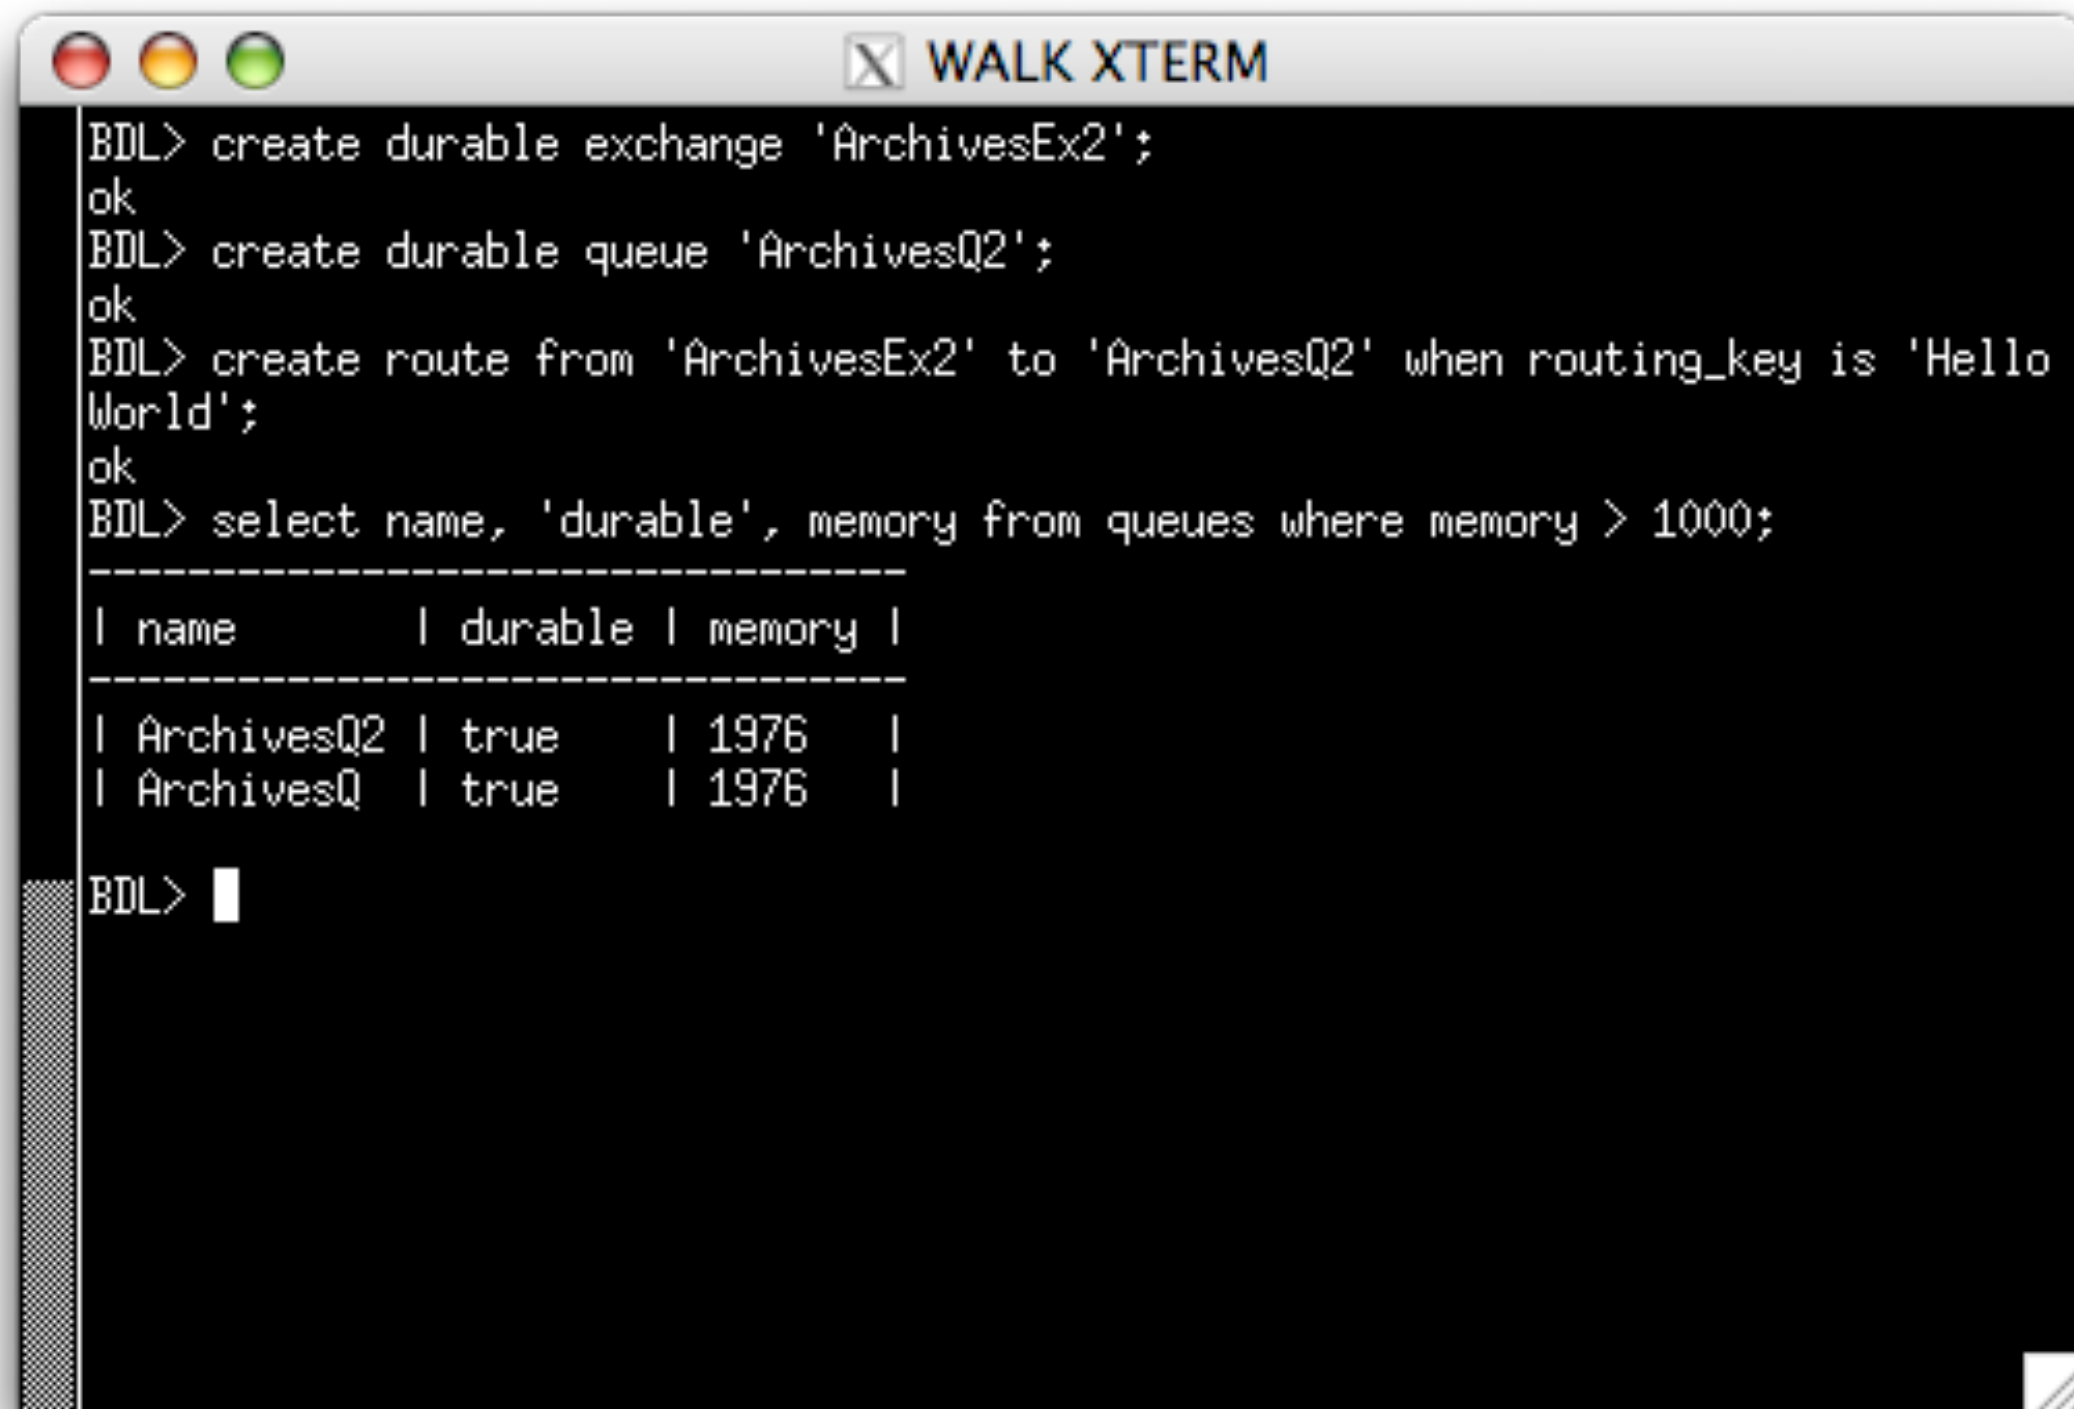


```
Jul 20 14:39:08 tonyg@nodachi2
~/dev/rabbitmq-umbrella/rabbitmq-server/scripts$ ./rabbitmqctl list_queues name
durable auto_delete messages_ready consumers memory
Listing queues ...
ArchivesQ2      true      false    0        0        1976
ArchivesQ       true      false    0        0        1976
...done.

Jul 20 14:40:50 tonyg@nodachi2
~/dev/rabbitmq-umbrella/rabbitmq-server/scripts$ ./rabbitmqctl list_bindings
Listing bindings ...
      ArchivesQ      ArchivesQ      []
      ArchivesQ2     ArchivesQ2     []
ArchivesEx      ArchivesQ      []
ArchivesEx2     ArchivesQ2     Hello%20World  []
...done.

Jul 20 14:46:12 tonyg@nodachi2
~/dev/rabbitmq-umbrella/rabbitmq-server/scripts$
```

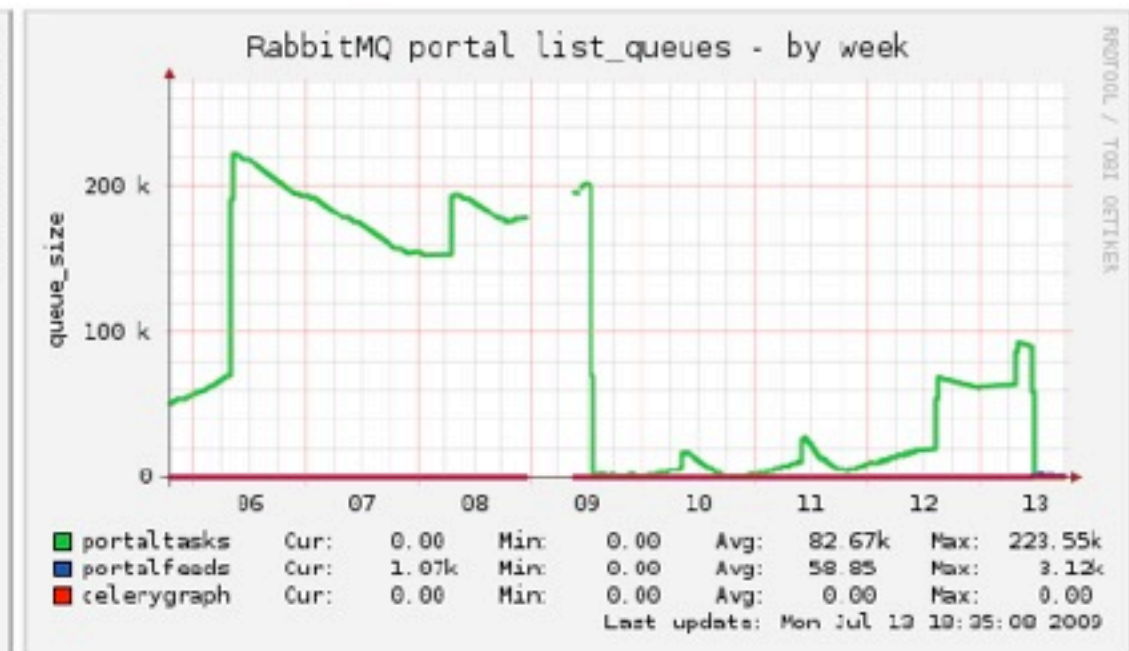
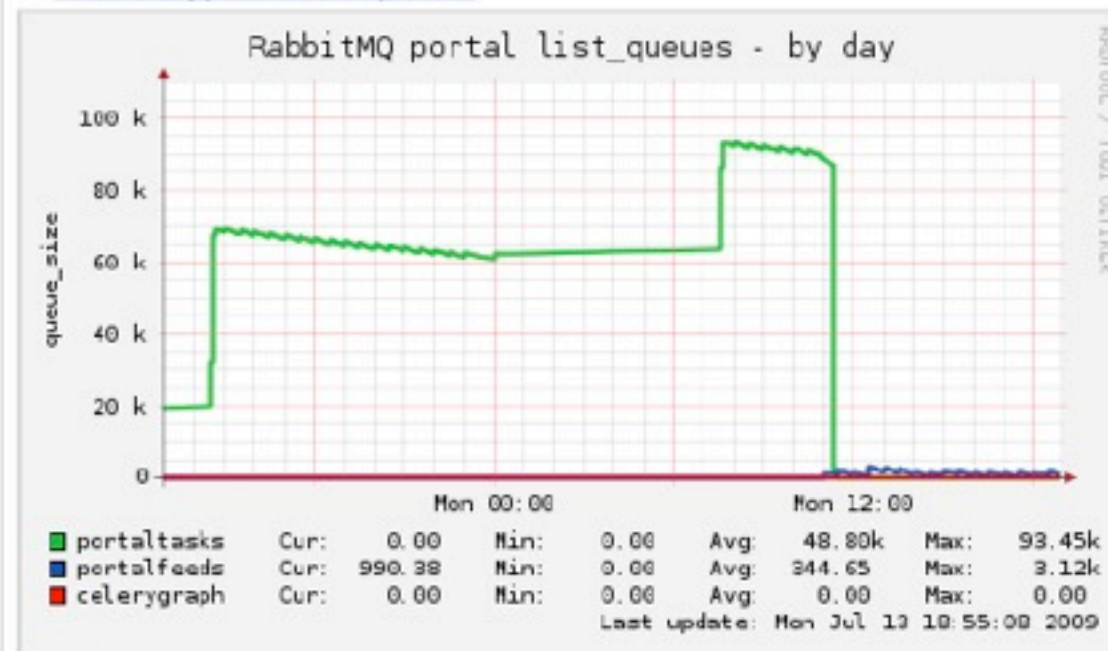
BDL



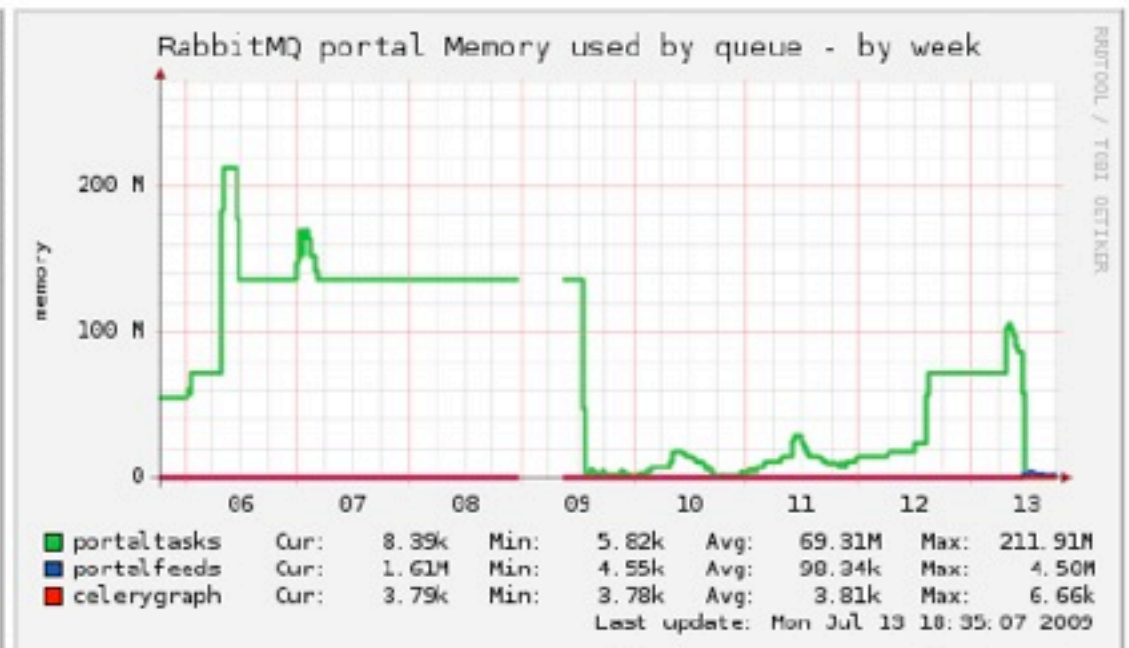
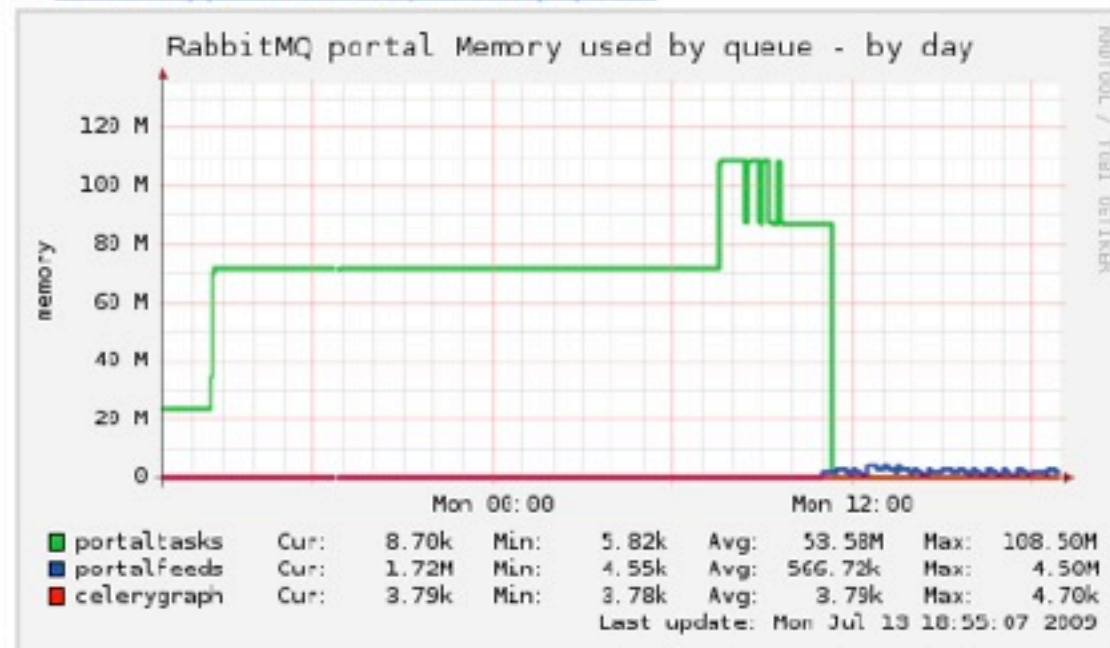
```
BDL> create durable exchange 'ArchivesEx2';
ok
BDL> create durable queue 'ArchivesQ2';
ok
BDL> create route from 'ArchivesEx2' to 'ArchivesQ2' when routing_key is 'Hello
World';
ok
BDL> select name, 'durable', memory from queues where memory > 1000;
-----
| name          | durable | memory |
-----
| ArchivesQ2    | true    | 1976    |
| ArchivesQ     | true    | 1976    |
BDL> 
```

Munin plugins

:: [RabbitMQ portal list_queues](#)



:: [RabbitMQ portal Memory used by queue](#)



RESTful APIs: Trixx

```
GET      http://localhost:8080/exchanges/:vhost
POST     http://localhost:8080/exchanges
GET      http://localhost:8080/queues/:vhost
POST     http://localhost:8080/queues
GET      http://localhost:8080/bindings/:vhost
GET      http://localhost:8080/vhosts
POST     http://localhost:8080/vhosts
GET      http://localhost:8080/connections
GET      http://localhost:8080/users/:user/permissions

GET      http://localhost:8080/users
PUT      http://localhost:8080/users

POST     http://localhost:8080/users

DELETE   http://localhost:8080/users/:user

GET      http://localhost:8080/rabbit/status
PUT      http://localhost:8080/rabbit/start
PUT      http://localhost:8080/rabbit/stop
PUT      http://localhost:8080/rabbit/reset

POST     http://localhost:8080/sessions/authenticate
```

RESTful APIs: Alice

/conn - Current connection information
/exchanges - Current exchanges information
/queues - Current queues
/users - Current users
/bindings - Current bindings
/control - Access to the RabbitMQ control
/permissions - Current permissions
/vhosts - Current vhosts

- - - - -
auser \$ curl -i -XPOST \
 -d '{"username": "ari", "password": "weak password"}' \
 http://localhost:9999/users

HTTP/1.1 200 OK
Server: MochiWeb/1.0 (Any of you quaid's got a smint?)
Date: Thu, 16 Jul 2009 00:10:35 GMT
Content-Type: text/json
Content-Length: 25

{"users": ["ari", "guest"]}

Web UIs: Trixx

The image displays two overlapping screenshots of the Trixx Web Console, a web-based interface for managing RabbitMQ services.

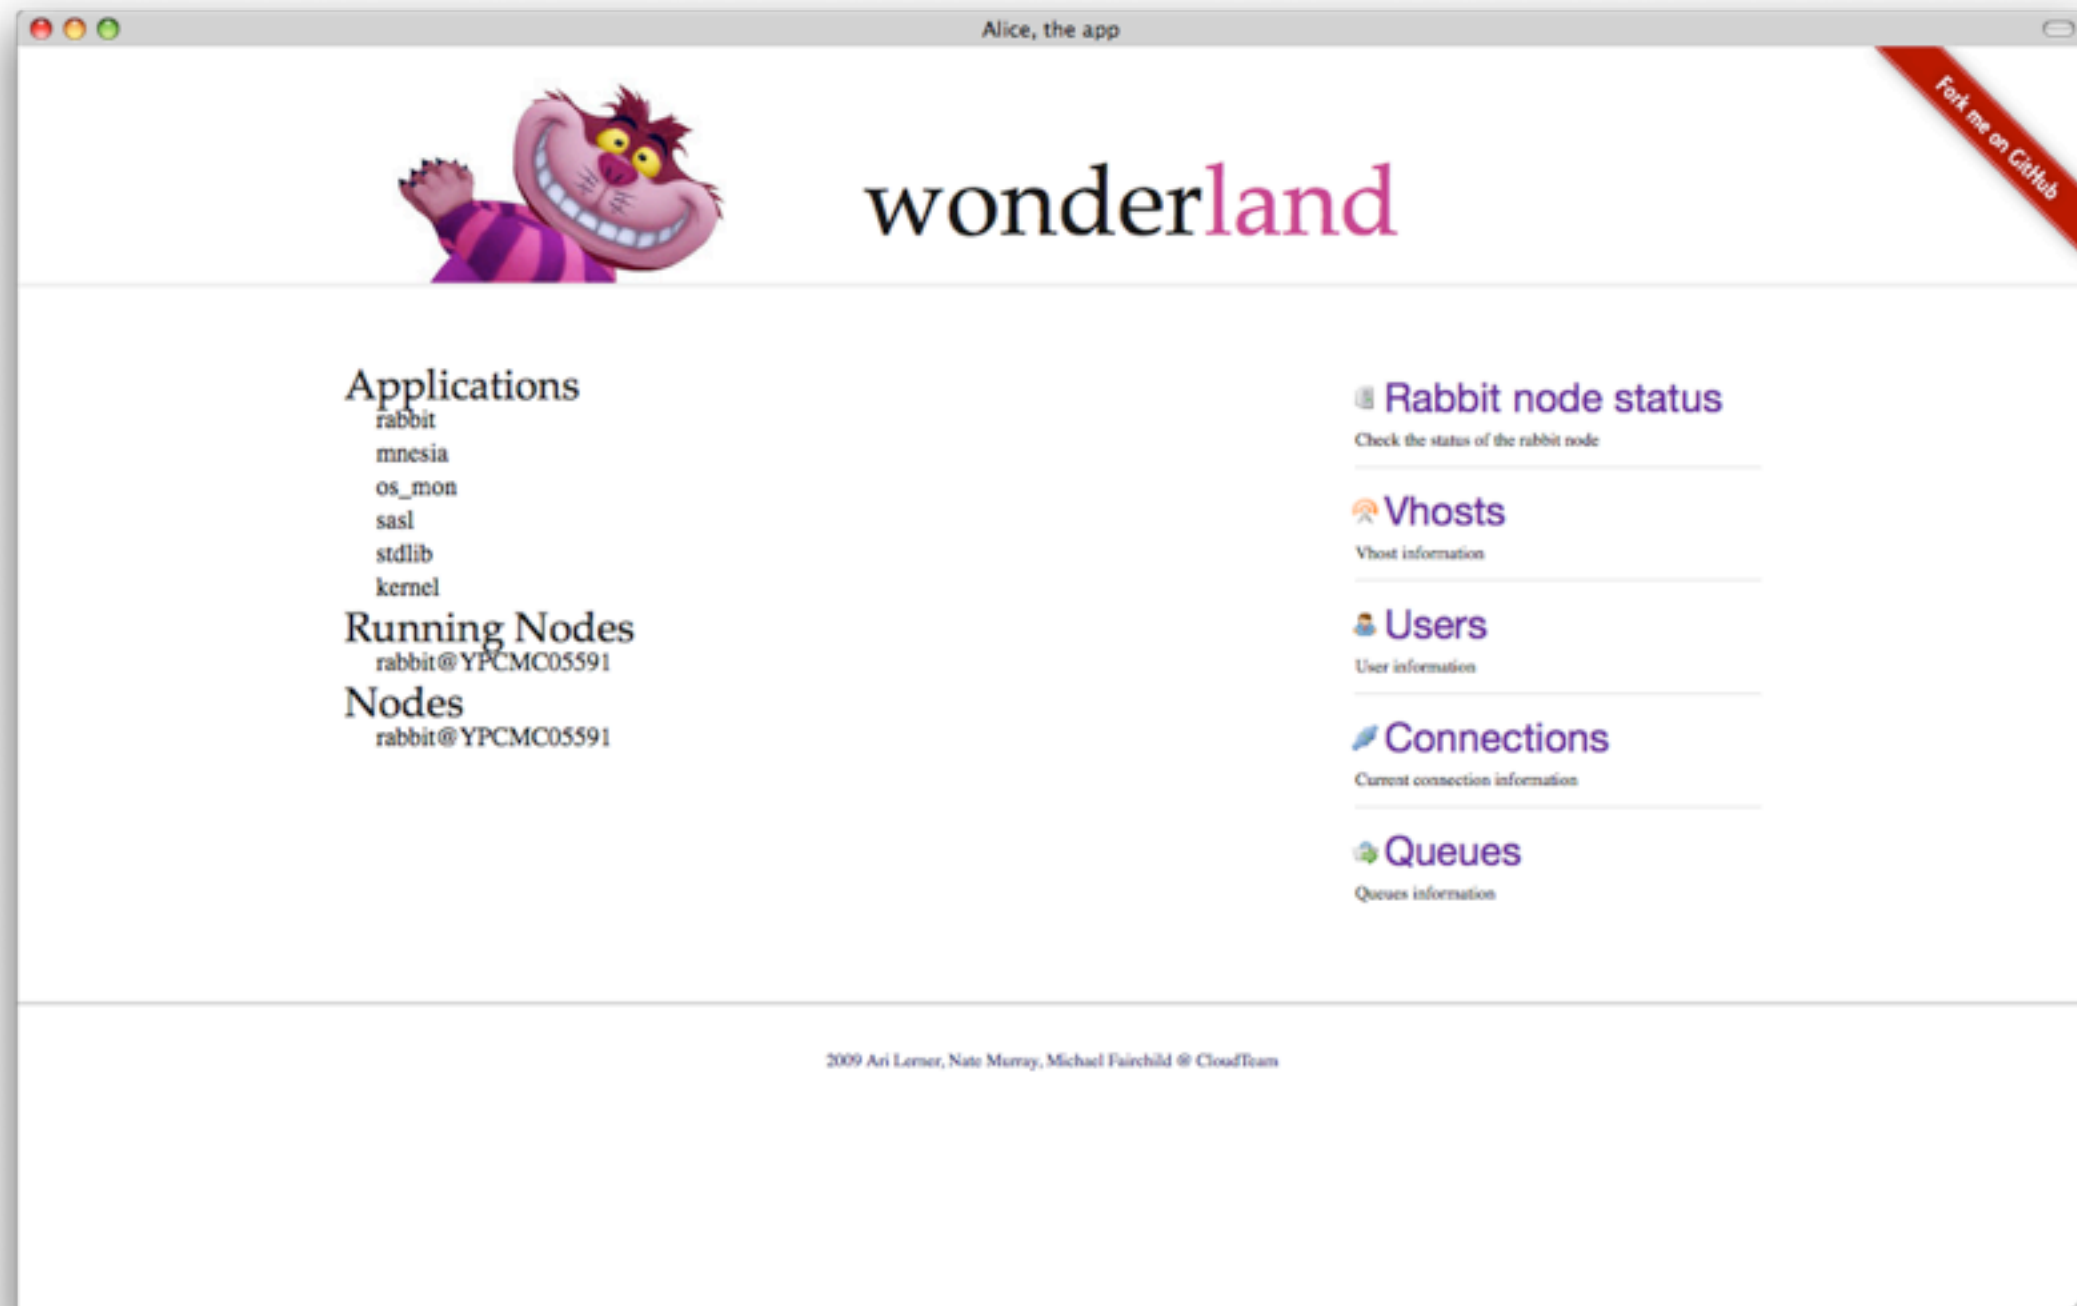
Left Screenshot: Trixx Dashboard

- Running Services:** rabbit, mnesia, os_mon, sasl, stdlib, kernel
- NODES:** running rabbit@nodachi2
- VHOST: /**
Select vhost
- Queues:**
QUEUE: ARCHIVESQ2 DURABLE
- Status:**
0 consumers 0 transactions 1.9 KB
- Messages:**
0 total 0 ready 0 unacknowledged

Right Screenshot: AMQ.FANOUT FANOUT

- Bindings:**
DIRECT
Queue: ArchivesQ (Routing Key: ArchivesQ)
Queue: ArchivesQ2 (Routing Key: ArchivesQ2)
- Users:**
Name Config Write Read
guest . * . *
- Connections:**
CONNECTION: 127.0.0.1:5672 FROM 127.0.0.1 BY GUEST RUNNING
1 channels 131,072 frame-max 0 timeout
225 recv oct 11 recv cnt
0 send pend 0 sent oct 5 sent cnt

Web UIs: Wonderland



Other ongoing projects

- Management
 - SNMP
 - JMX, HermesJMS
 - Presence, meta-events
 - AMQP Protocol extensions for management
- Gateways, federation
 - RabbitHub (HTTP)
 - Multicast
 - XMPP, JMS bridges
 - AMQP Protocol extensions for generalised gatewaying